



A blog dedicated to Software Developers, Product Managers and Marketing

[Software Analytics](#)[Product Management](#)[Company News](#)[Our Website](#)

Setting up Titan 1.0 with Apache HBase

2 February 2016 | [Arielle Bonnici](#) | Big Data, Software Analytics, Software

Development



During a recent project we researched a number of graph databases. One of these was Titan, currently regarded as one of the leading graph databases having native integration with the Apache TinkerPop graph stack. The plan was to create a test setup using an **HBase** storage backend which we had running on **Hadoop2**.

Although the official Titan Documentation is quite extensive and a good resource to start wrapping your head around how it works, the examples are quite generic and do not provide a clear picture on getting started with your choice of backend and putting the pieces together.

Getting it all up and running involved several steps and hiccups along the way which consumed a bunch of research time on blogs, forums and discussion groups. Therefore we decided to gather all the info and put together these basic steps in a single blog post in the hope that it will save you some time if you are also trying to get started with Titan on HBase.

This post will guide you to:

- Install Titan
- Install Elasticsearch
- Run the Gremlin console and create a new graph
- Run the Gremlin server for remote connections



Do you know how people use your software?

Find Out With Trackerbird

[TRY IT FOR FREE](#)

Articles by Category

Articles by Category

Select Category

Recent Posts

Using Open Source Libraries with Closed Source Software

Collecting Custom Data with Event Tracking

Setting up Titan 1.0 with Apache HBase

Related Links

Trackerbird API Reference

Trackerbird Pricing

Trackerbird Support Helpdesk

Archives

Pre-Requisite

Titan requires Java 8 and the `$JAVA_HOME` environment variable needs to be configured to point to the directory where the JRE or JDK is installed.

We also had HBase 0.98.12 running on Hadoop 2.7.1, all set up on an Ubuntu 14.04 machine.

Installing Titan

The first thing you'll need to do is download Titan from <https://github.com/thinkaurelius/titan/wiki/Downloads>. At the time of writing the latest version of Titan is 1.0 and you will need to download the **titan-1.0.0-hadoop1.zip** file.

Unzip the file you downloaded from a shell:

```
$ unzip titan-1.0.0-hadoop1.zip
```

Let me side-track for a minute here and say that initially we were led to believe that since we have Hadoop 2 in our setup, we needed to use the titan-1.0.0-hadoop2.zip file. When we tried to start the Gremlin console we encountered problems right away and got an error to resolve `'com.thinkaurelius.titan.hadoop.MapReduceIndexManagement'`.

*Invalid import definition:
'com.thinkaurelius.titan.hadoop.MapReduceIndexManagement';
reason: startup failed:
script14522627513712042298222.groovy: 1: unable to
resolve class
com.thinkaurelius.titan.hadoop.MapReduceIndexManagement*

Then if you try to open the graph by executing: `graph = TitanFactory.open("conf/titan-hbase-es.properties")`, you will get the error below:

*Could not instantiate implementation:
com.thinkaurelius.titan.diskstorage.hbase.HBaseStoreManager*

This issue occurs because the `$TITAN_HOME/lib/titan-hadoop-`

February 2016

January 2016

October 2015

September 2015

April 2015

January 2015

December 2014

October 2014

July 2014

May 2014

April 2014

January 2014

December 2013

October 2013

July 2013

April 2013

February 2013

January 2013

December 2012

November 2012

September 2012

Sign up for our mailing list

First Name :

Last Name :

Email Address :

1.0.0.jar is missing when using titan-1.0.0-hadoop2.zip.

Installing Elasticsearch

Titan comes with native support for composite indexes, however if you want to make use of mixed indexes you will need to configure an indexing backend. Titan currently supports Lucene, Solr and Elasticsearch. We chose to install Elasticsearch for the test setup which was pretty straight forward and “just worked”. Note that we didn’t bother with security as at this stage we just needed to get things going, but that is something you will need to look into if you decide to take it further. Have a look [here](#).

To install Elasticsearch:

Go to the \$TITAN_HOME/bin folder and search for the elasticsearch jar. In the case of Titan 1.0 it comes with Elasticsearch 1.5.1 therefore you should find elasticsearch-1.5.1.jar. This is the Elasticsearch version you should install.

Open a shell and download the corresponding Elasticsearch version:

```
$ wget  
https://download.elastic.co/elasticsearch/elasticsearch  
1.5.1.deb
```

Install it by running the dpkg command:

```
$ sudo dpkg -i elasticsearch-1.5.1.deb
```

Elasticsearch will be installed in /usr/share/elasticsearch and its configuration files will be located in /etc/elasticsearch.

Run the following command so that Elasticsearch starts and stops automatically:

```
$ sudo update-rc.d elasticsearch defaults
```

Open the /etc/elasticsearch/Elasticsearch.yaml file and set a name for the node and cluster (remember to remove the # at the beginning of the lines):

```
cluster.name: mycluster
```

```
node.name: mynode
```

For testing purposes, you can set the number of shards and replicas as follows:

```
index.number_of_shards: 1
index.number_of_replicas: 0
```

Save the changes and start Elasticsearch:

```
$ sudo service elasticsearch start
```

Running the Gremlin Console

You are now ready to connect to Titan, create a graph and start playing with it. Using the Gremlin console is a good starting point to do that. This will create a new table in HBase called 'titan'.

Open `$TITAN_HOME/conf/titan-hbase-es.properties`, set the storage hostname / IP of your HBase server(s) and save the file:

```
storage.hostname: 127.0.0.1
```

We have left the default 127.0.0.1 since we have everything set up locally.


Open a shell and change the working directory to the Titan home folder:

```
$ cd titan-1.0.0-hadoop1
```

Start the Gremlin console:

```
$ sudo ./bin/gremlin.sh
```

If you omit the `sudo`, you'll encounter an access denied error when Gremlin is loading:



```
Exception in thread "main" java.io.FileNotFoundException:
/usr/local/titan-1.0.0-hadoop1/ext/plugins.txt (Permissions
denied)
```

Open the graph and specify the Titan / HBase / Elasticsearch


properties file:

```
graph = TitanFactory.open('conf/titan-hbase-  
es.properties')
```

If that works you will see:

```
==>standardtitangraph[hbase:[127.0.0.1]]
```

The path to the properties file you specify is relative to the current working directory. If you didn't change the working directory to the Titan home folder before running the Gremlin console, you will encounter an error here:



```
Backend shorthand unknown: /conf/titan-hbase-  
es.properties
```

Or if you try to specify the complete path to the properties file:



```
Could not open global configuration
```

At this stage you can start adding vertices and edges to your graph. If you're new to the Gremlin language, you will find that this presentation on [The Gremlin Graph Traversal Language](#) will help to get you started, as well as the [TinkerPop3 Documentation](#).

Running the Gremlin Server for Remote Connections

When you connect to the Titan graph from the Gremlin console, the graph is only accessible through that session and once the Gremlin console is stopped Titan will also stop, since its main thread of execution was part of the Gremlin shell. To run Titan as a long running process which you'll be able to connect to remotely, you will need to configure and run the Gremlin server.

Copy \$TITAN_HOME/conf/titan-hbase-es.properties to the \$TITAN_HOME/conf/gremlin-server directory and add the following:

```
gremlin.graph=com.thinkaurelius.titan.core.TitanFactory  
storage.hbase.tablename = titan
```

Open the \$TITAN_HOME/conf/gremlin-server/gremlin-server.yaml

and change:

```
host: localhost
```

to:

```
host: 0.0.0.0
```

Find the graph entry and change:

```
graph: conf/gremlin-server/titan-berkelyje-  
server.properties
```

to:

```
graph: conf/gremlin-server/titan-hbase-es.properties
```

Open a shell and change the working directory to the Titan home folder:

```
$ cd titan-1.0.0-hadoop1
```

Start the Gremlin server:

```
$ sudo ./bin/gremlin-server.sh
```

You can test the connection by opening another shell and starting the Gremlin console (see above), then run the following command:

```
:remote connect tinkerepop.server conf/remote.yaml
```

If the connection is successful you will see:

```
==>Connected - localhost/127.0.0.1:8182
```

You will then need to enter `:>` before your Gremlin commands, so they will be sent to the currently active remote connection:

```
:> g = graph.traversal()
```


We hope this blog post has provided some useful information to make your Titan implementation less painful.

G+1

Like

8

Comments are closed.



EXPLORE TRACKERBIRD

- Overview
- How It Works
- Why Trackerbird
- Business Benefits

ABOUT US

- About Trackerbird
- Supporting Startups

CONTACT

- Contact Trackerbird
- Partners & Affiliates

FOLLOW US:

- rss
- facebook
- twitter
- linkedin



Copyright © 2016 - Trackerbird Software Analytics