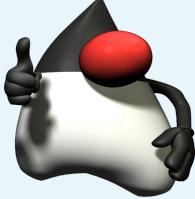


# The Gremlin Graph Traversal Machine and Language

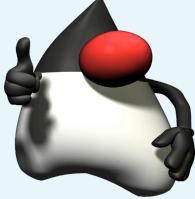


Dr. Marko A. Rodriguez  
Director of Engineering at DataStax, Inc.  
Project Management Committee, Apache TinkerPop





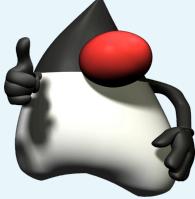
Java is a virtual machine.  
Java is a programming language.



Java is a virtual machine.  
Java is a programming language.



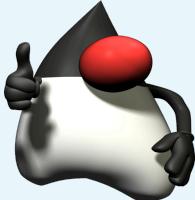
Gremlin is a traversal machine.  
Gremlin is a traversal language.



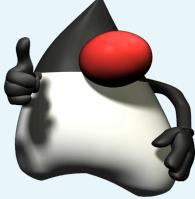
Java is a virtual machine.  
Java is a programming language.



Gremlin is a traversal machine.  
Gremlin is a traversal language.



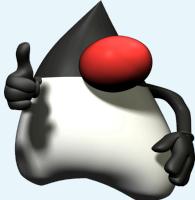
Java is operating system agnostic.  
MacOSX, Linux, Windows, etc.



Java is a virtual machine.  
Java is a programming language.



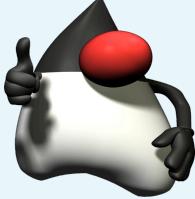
Gremlin is a traversal machine.  
Gremlin is a traversal language.



Java is operating system agnostic.  
MacOSX, Linux, Windows, etc.



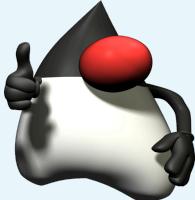
Gremlin is graph system agnostic.  
Titan, Neo4j, Stardog, Giraph, Spark, Hadoop, etc.



Java is a virtual machine.  
Java is a programming language.



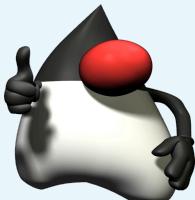
Gremlin is a traversal machine.  
Gremlin is a traversal language.



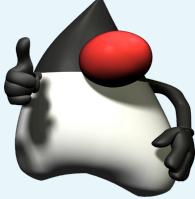
Java is operating system agnostic.  
MacOSX, Linux, Windows, etc.



Gremlin is graph system agnostic.  
Titan, Neo4j, Stardog, Giraph, Spark, Hadoop, etc.



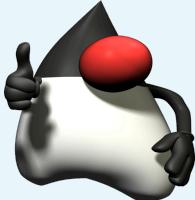
Other languages compile to the Java virtual machine.  
Groovy, Scala, Clojure, JavaScript Nashorn, etc.



Java is a virtual machine.  
Java is a programming language.



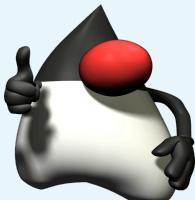
Gremlin is a traversal machine.  
Gremlin is a traversal language.



Java is operating system agnostic.  
MacOSX, Linux, Windows, etc.



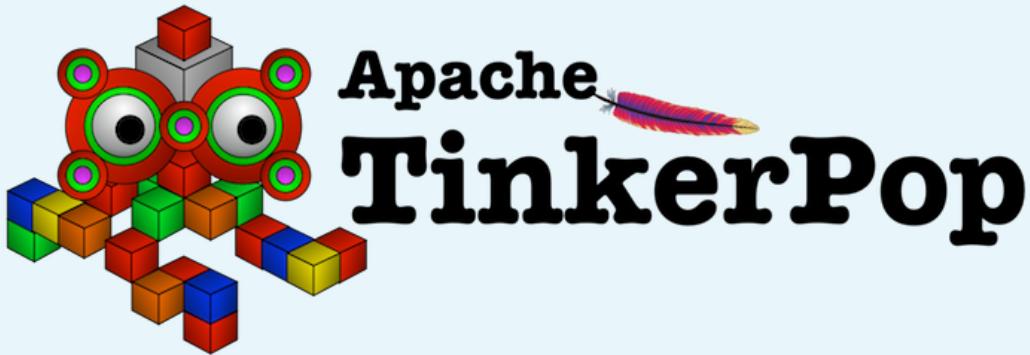
Gremlin is graph system agnostic.  
Titan, Neo4j, Stardog, Giraph, Spark, Hadoop, etc.



Other languages compile to the Java virtual machine.  
Groovy, Scala, Clojure, JavaScript Nashorn, etc.

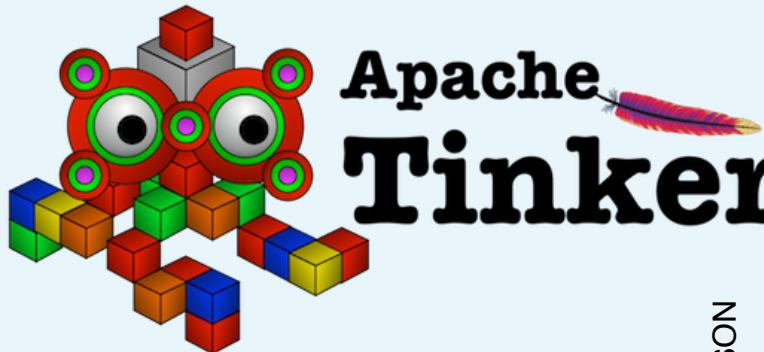


Other languages compile to the Gremlin traversal machine.  
Gremlin-Groovy, Gremlin-Scala, SPARQL, etc.

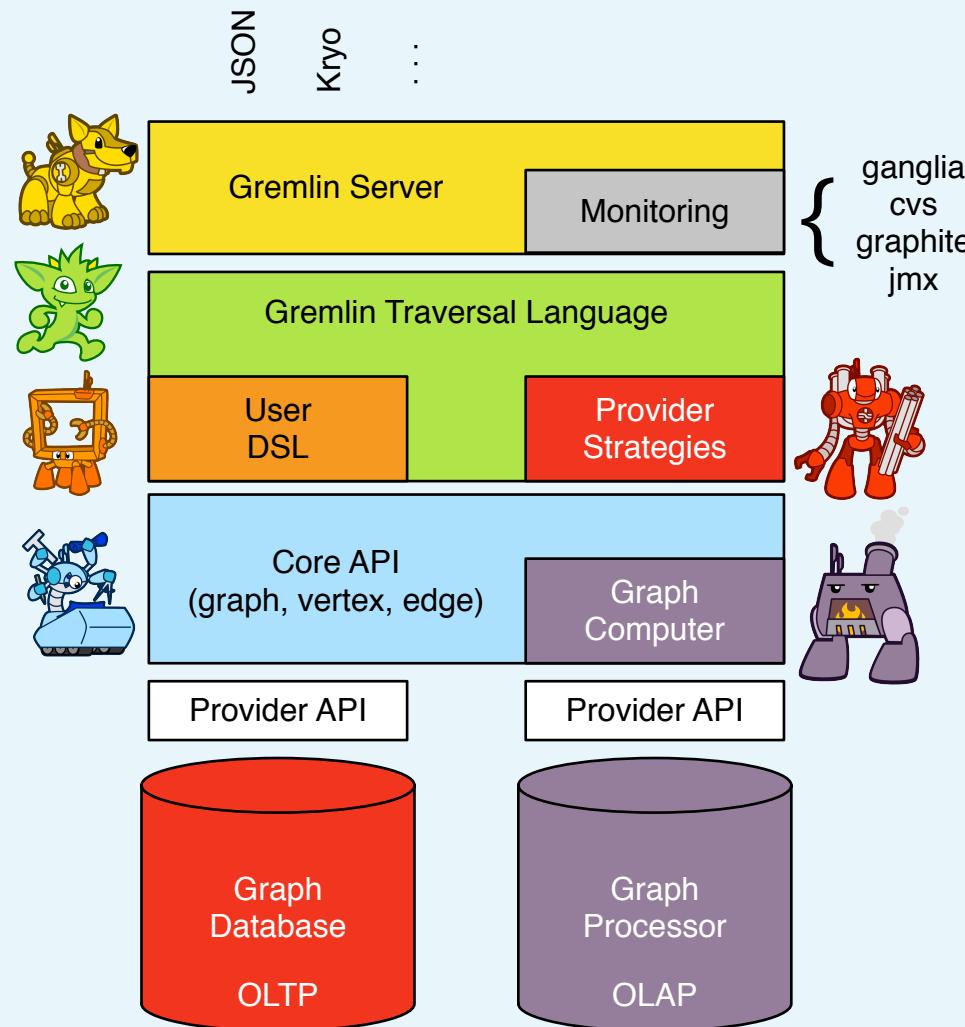


- TinkerPop is an open source graph computing framework.
- TinkerPop was started in 2009 by Marko A. Rodriguez, Josh Shinavier, and Peter Neubauer.
- TinkerPop joined the Apache Software Foundation January 2015 (currently in incubation).
- TinkerPop designs and develops the Gremlin graph traversal machine and language.
- TinkerPop is supported by nearly every graph system provider  
(both commercial and open source).
- TinkerPop<sup>3</sup> has been in development for 2 years and was officially released July 2015.

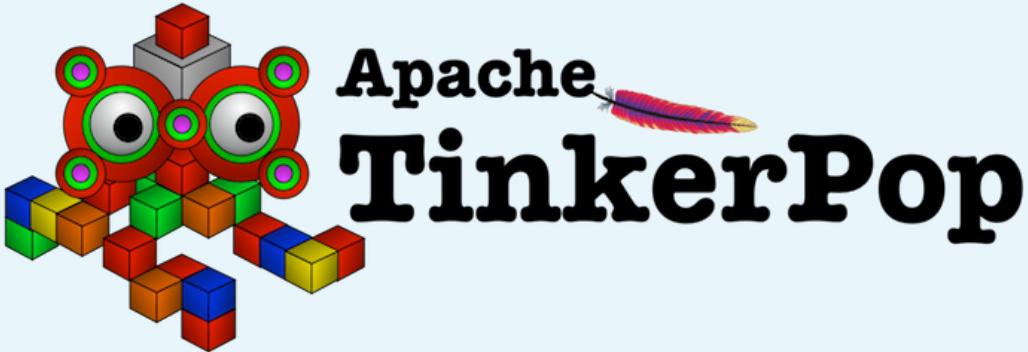
"What is The TinkerPop?"



# Apache TinkerPop

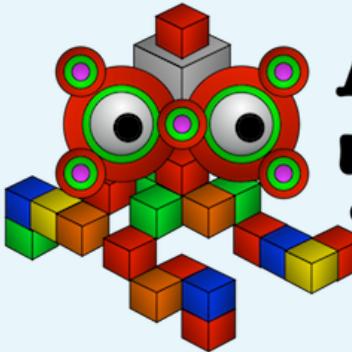


"What is The TinkerPop?"



- TinkerPop provides graph computing capabilities to any data processing system.
- TinkerPop supports both OLTP (database) and OLAP (batch processor) systems.
- TinkerPop is used to process the order fulfillment network of Amazon.com ("100 billion vertices" -- at least factor 10x on edges?).  
<https://videos.dabcc.com/aws-reinvent-2015-dat203-building-graph-databases-on-aws/>
- TinkerPop is the sole interface for >50% of all graph systems in the market/wild.
- TinkerPop currently processes a near trillion edge graph represented over 1000 machines.  
(can't disclose company name)
- TinkerPop works with Apache Cassandra, Apache HBase, Apache Hadoop, Apache Spark, Apache Giraph, Apache Atlas, Apache Falcon, ...
- TinkerPop is always looking for more contributors and perhaps TinkerPop could be your future software/language design and development home.

"Why should you care about The TinkerPop?"

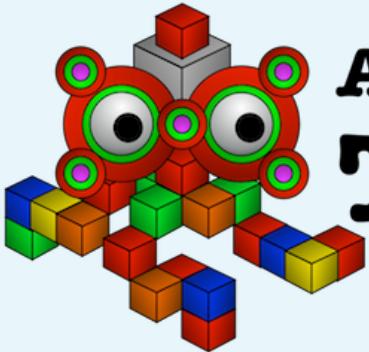


# Apache TinkerPop

TinkerPop has language bindings in every major programming language.

- TinkerPop provides graph computing capabilities to any data processing system.
- TinkerPop supports both OLTP (database) and OLAP (batch processor) systems.
- TinkerPop is used to process the order fulfillment network of Amazon.com ("100 billion vertices" -- at least factor 10x on edges?).  
<https://videos.dabcc.com/aws-reinvent-2015-dat203-building-graph-databases-on-aws/>
- TinkerPop is the sole interface for >50% of all graph systems in the market/wild.
- TinkerPop currently processes a near trillion edge graph represented over 1000 machines.
- TinkerPop works with Apache Cassandra, Apache HBase, Apache Hadoop, Apache Spark, Apache Giraph, Apache Atlas, Apache Falcon, ...
- TinkerPop is always looking for more contributors and perhaps TinkerPop could be your future software/language design and development home.

"Why should you care about The TinkerPop?"

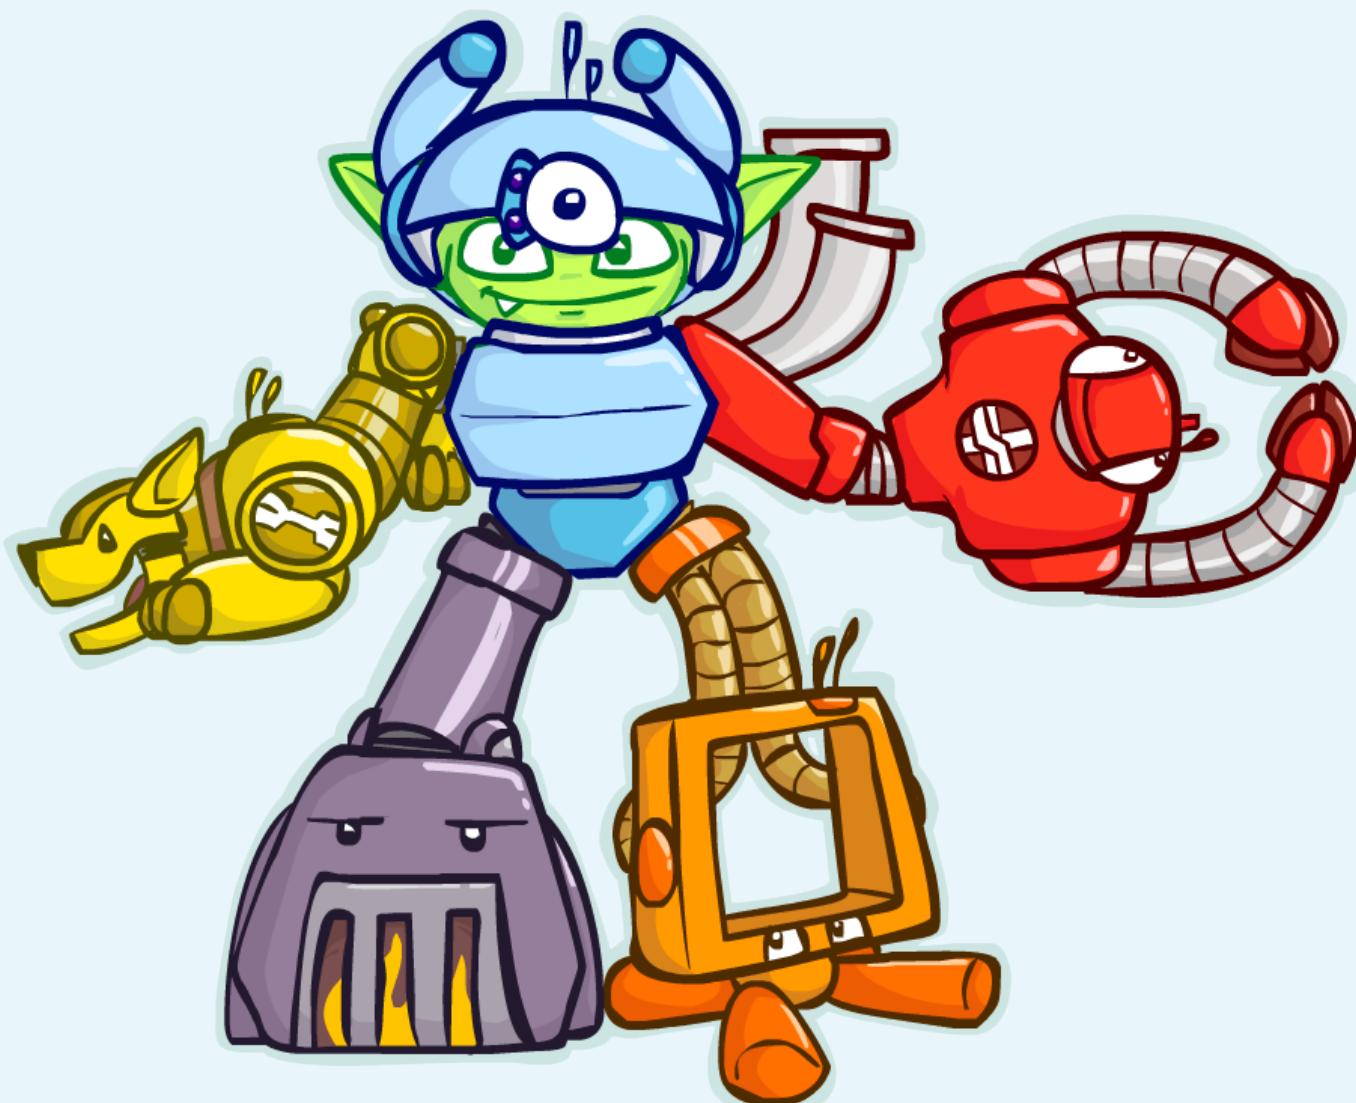


# Apache TinkerPop

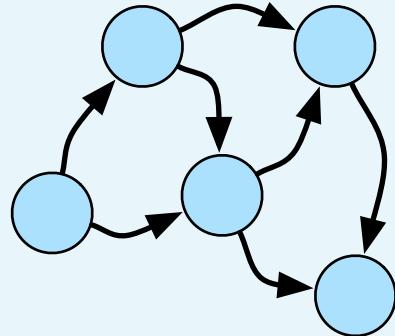
- TinkerPop provides graph computing capabilities to any data processing system.
- TinkerPop supports both OLTP (database) and OLAP (batch processor) systems.
- TinkerPop is used to process the order fulfillment network of Amazon.com ("100 billion vertices" -- at least factor 10x on edges).  
<https://videos.dabcc.com/aws-reinvent-2015-dat203-building-graph-databases-on-aws/>
- TinkerPop is the sole interface for >50% of all graph systems in the market/wild.
- TinkerPop currently processes a near trillion edge graph represented over 1000 machines.
- TinkerPop works with Apache Cassandra, Apache HBase, Apache Hadoop, Apache Spark, Apache Giraph, Apache Atlas, Apache Falcon, ...
- TinkerPop is always looking for more contributors and perhaps TinkerPop could be your future software/language design and development home.

"Why should you care about The TinkerPop?"

# Part 1: The Gremlin Traversal Machine



# The Machine Components



The Graph



The Traverser



The Traversal

The Data

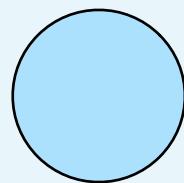
The CPU

The Program



# The Graph

**vertex**

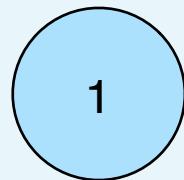


"The graph is the 'RAM'."



# The Graph

**vertex id**



"Vertices and edges in the graph have unique ids (addresses)."



# The Graph

**vertex label**

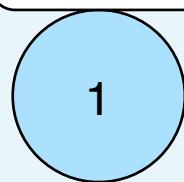




# The Graph

## vertex properties

name:marko  
age:36

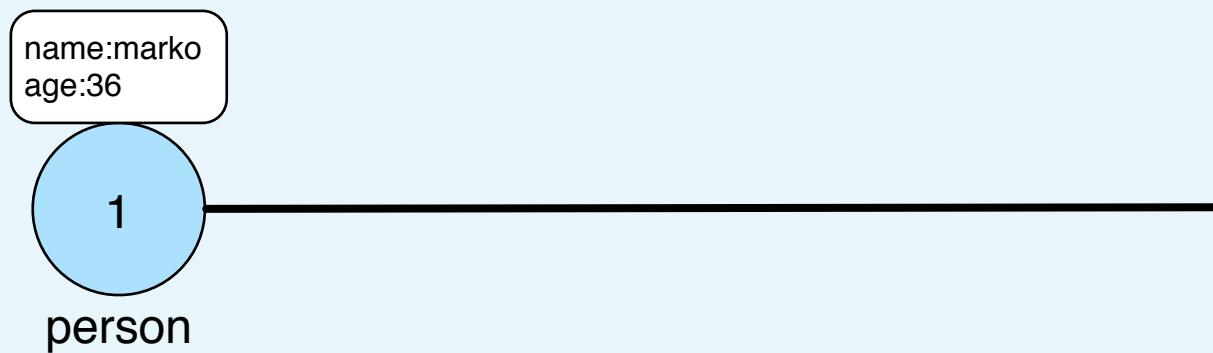


person



# The Graph

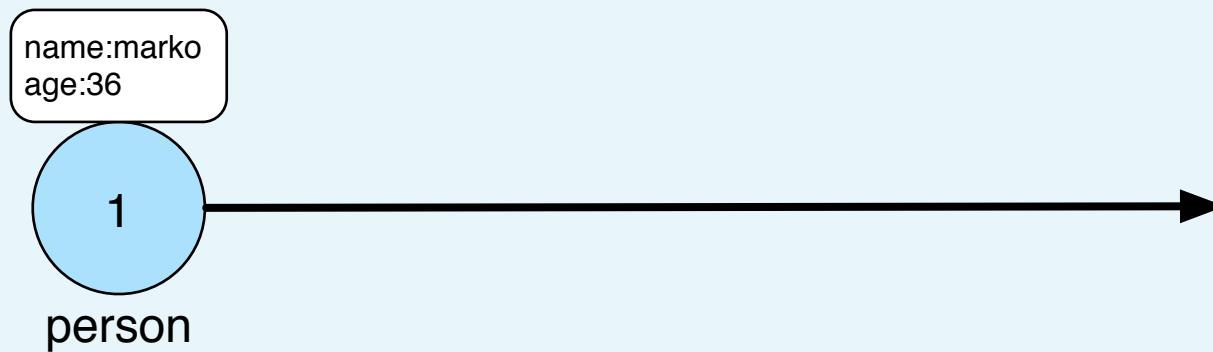
edge





# The Graph

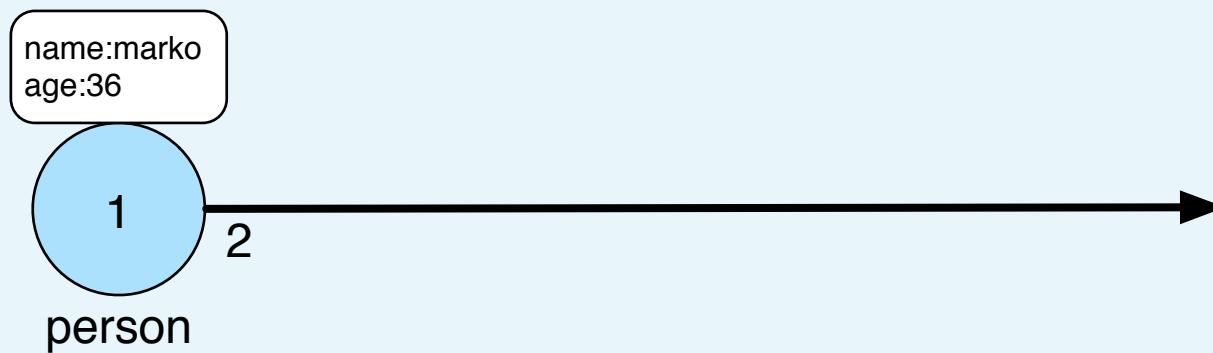
**edge direction**





# The Graph

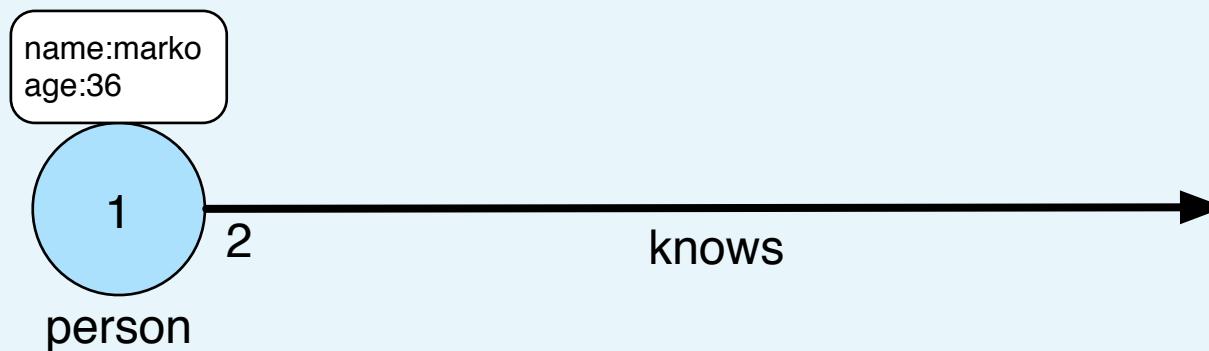
**edge id**





# The Graph

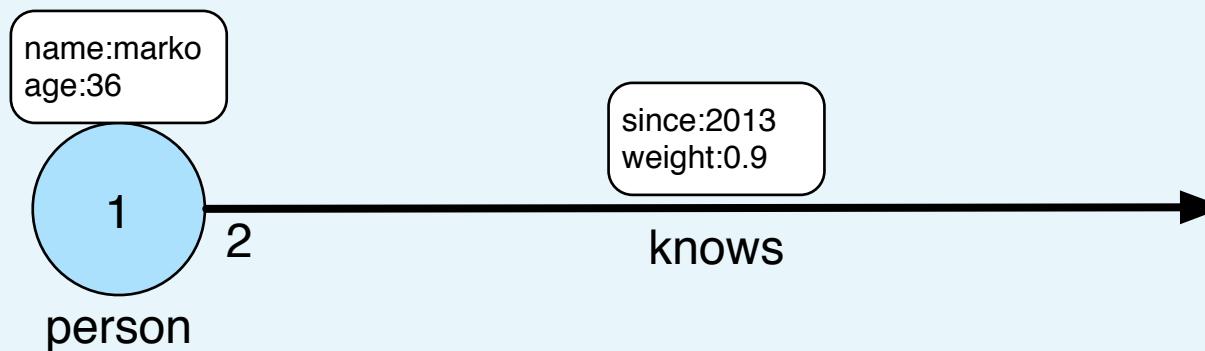
**edge label**





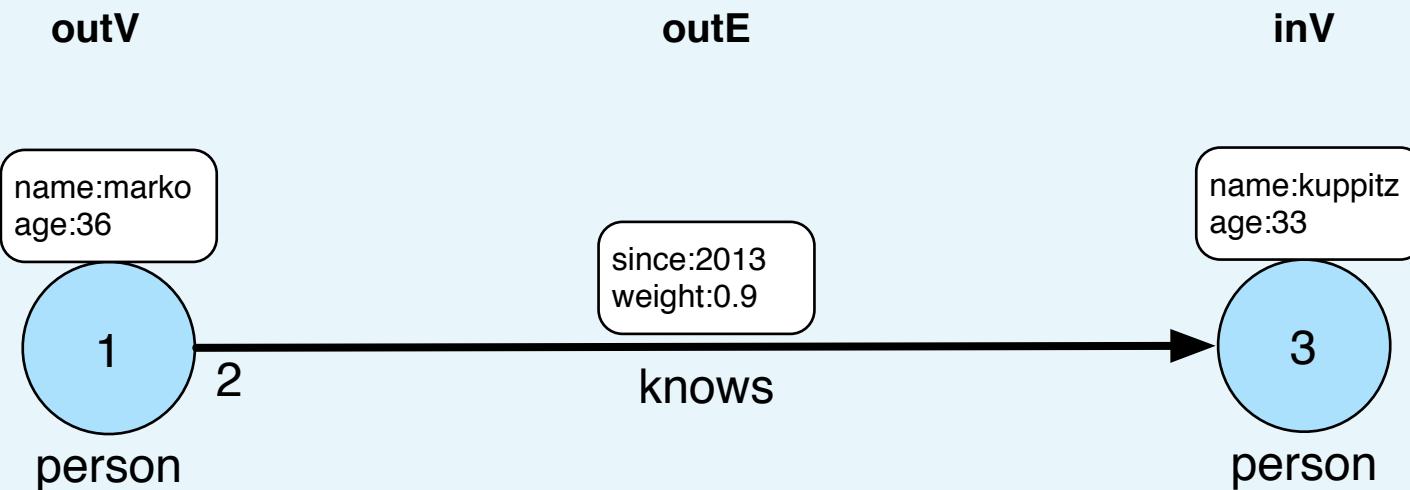
# The Graph

## edge properties





# The Graph

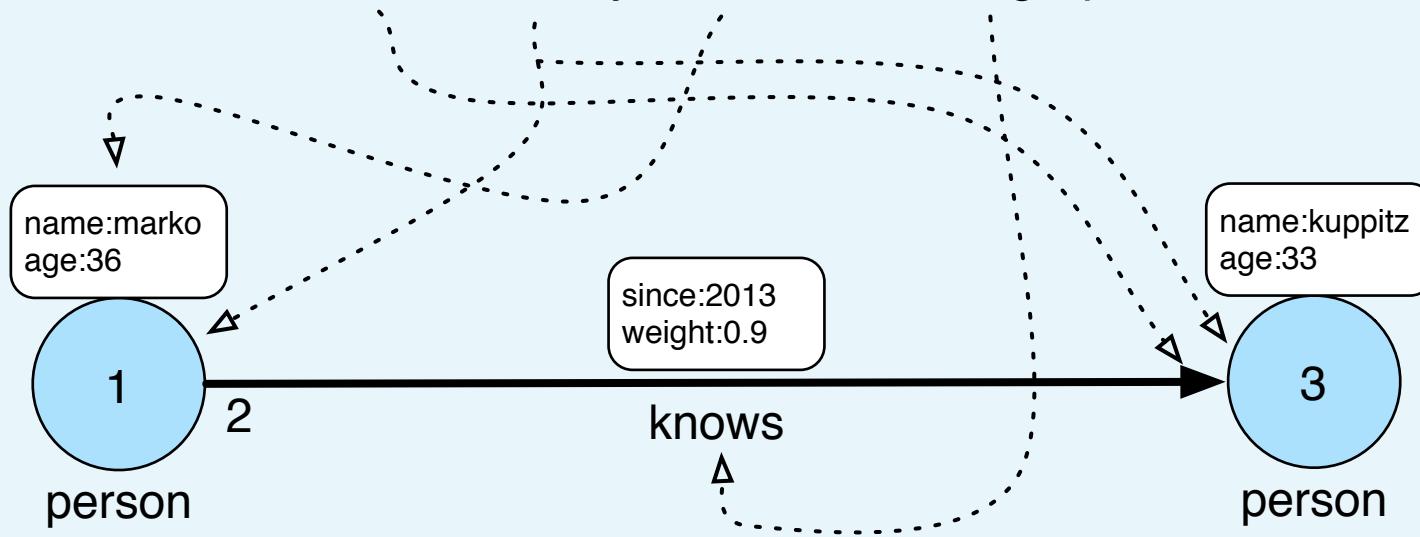


"Vertices are related by edges (addresses have pointers to each other)."



# The Graph

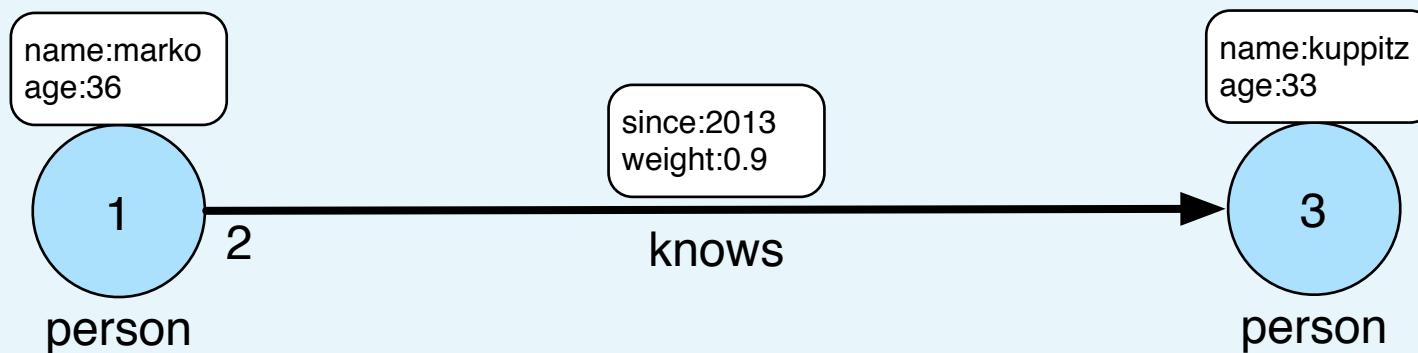
Directed, binary, attributed multi-graph.





# The Graph

Directed, binary, attributed multi-graph.



edges

directed, binary

vertices or edges

property key  
(string)

properties can't reference  
vertices or edges

$$G = (V, E \subseteq (V \times V), \lambda : (V \cup E) \times \Sigma^* \rightarrow U \setminus (V \cup E))$$

vertices

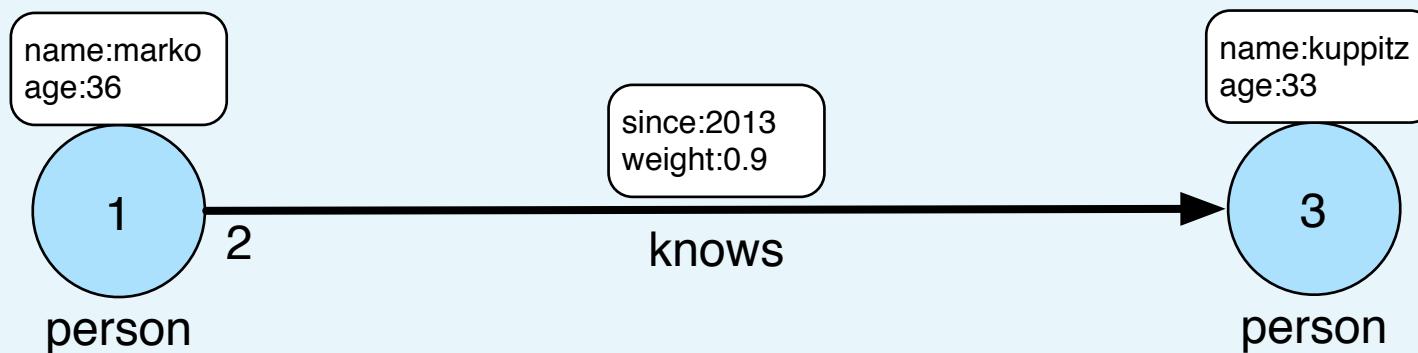
labels+properties  
function

U = anything



# The Graph

Directed, binary, attributed multi-graph.



$$G = (V, E \subseteq (V \times V), \lambda : (V \cup E) \times \Sigma^* \rightarrow U \setminus (V \cup E))$$

$$\lambda(v_1, \text{name}) \mapsto \text{marko}$$

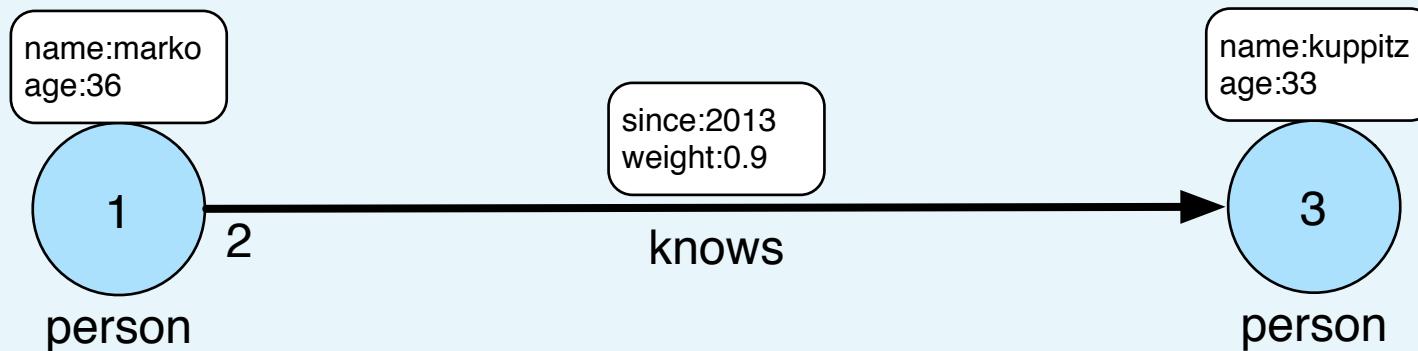
$$\lambda(e_2, \text{label}) \mapsto \text{knows}$$

$$((e_2)_1, (e_2)_2) = (v_1, v_3)$$



# The Graph

~~Directed, binary, attributed multi graph.~~



## Property graph.

\* Multi- and meta-properties are not discussed in this presentation nor in the associated conference article.  
<http://tinkerpop.incubator.apache.org/docs/3.0.2-incubating/#vertex-properties>

```
~/tinkerpop3$ bin/gremlin.sh
```



```
gremlin>
```

```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/  
(o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin>
```

```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/
(o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin>
```

"TinkerGraph is an in-memory graph system provided by TinkerPop."

```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/  
(o o)  
----o00o-(3)-o00o----
```

```
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin> graph = TinkerGraph.open()  
==>tinkergraph[vertices:0 edges:0]  
gremlin>
```



TitanGraph.open(...)



Neo4jGraph.open(...)



OrientGraph.open(...)



StardogGraph.open(...)



DSEGraph.open(...)



HadoopGraph.open(...)



HadoopGraph.  
compute(GiraphGraphComputer)



HadoopGraph.  
compute(SparkGraphComputer)

etc...

"Gremlin is agnostic to the underlying graph system."

```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/  
(o o)  
----o00o-(3)-o00o----  
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin> graph = TinkerGraph.open()  
==>tinkergraph[vertices:0 edges:0]  
gremlin> v1 = graph.addVertex(id,1,label,'person')  
==>v[1]  
gremlin>
```



"A graph system provider is Gremlin-enabled if they implement the `gremlin.structure` API suite."

```
~/tinkerpop3$ bin/gremlin.sh
```

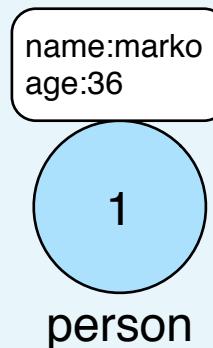
```
\,,,/  
(o o)  
----o00o-(3)-o00o----  
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin> graph = TinkerGraph.open()  
==>tinkergraph[vertices:0 edges:0]  
gremlin> v1 = graph.addVertex(id,1,label,'person')  
==>v[1]  
gremlin> v1.property('name','marko')  
==>vp[name->marko]  
gremlin>
```



"The APIs have to do with CRUD operations on a graph structure."

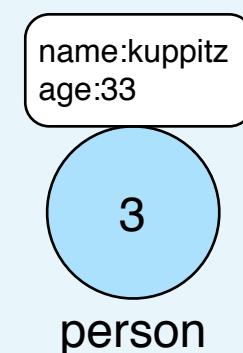
```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/  
(o o)  
----o00o-(3)-o00o----  
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin> graph = TinkerGraph.open()  
==>tinkergraph[vertices:0 edges:0]  
gremlin> v1 = graph.addVertex(id,1,label,'person')  
==>v[1]  
gremlin> v1.property('name','marko')  
==>vp[name->marko]  
gremlin> v1.property('age',36)  
==>vp[age->36]  
gremlin>
```



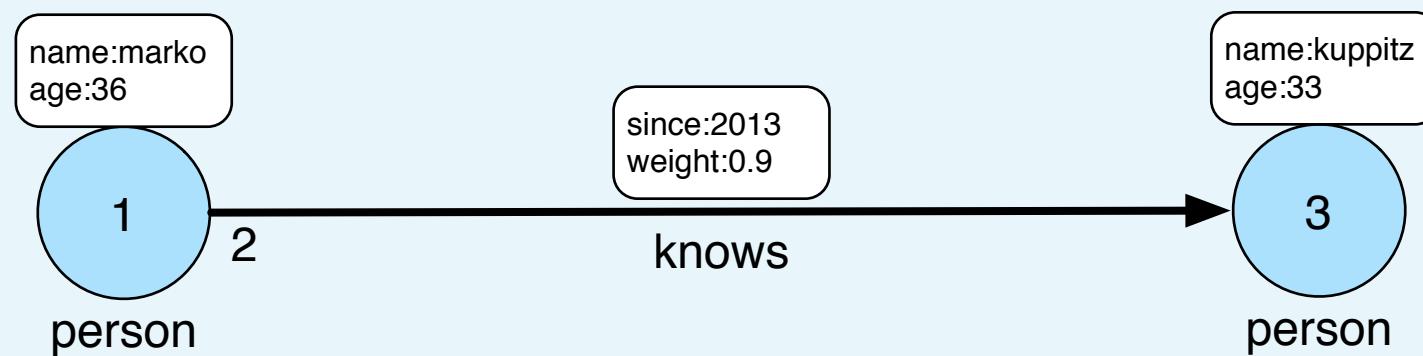
```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/
(o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> v1 = graph.addVertex(id,1,label,'person')
==>v[1]
gremlin> v1.property('name','marko')
==>vp[name->marko]
gremlin> v1.property('age',36)
==>vp[age->36]
gremlin> v3 = graph.addVertex(id,3,label,'person','name','kuppitz','age',33)
==>v[3]
gremlin>
```

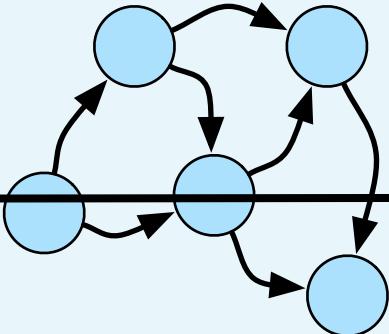


```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/ 
(o o)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> v1 = graph.addVertex(id,1,label,'person')
==>v[1]
gremlin> v1.property('name','marko')
==>vp[name->marko]
gremlin> v1.property('age',36)
==>vp[age->36]
gremlin> v3 = graph.addVertex(id,3,label,'person','name','kuppitz','age',33)
==>v[3]
gremlin> v1.addEdge('knows',v3,id,2,'since',2013,'weight',0.9)
==>e[2][1-knows->3]
gremlin>
```



# The Machine Components



The Graph



The Traverser



The Traversal

The Data

The CPU

The Program

# The Traverser



"The traverser is a 'CPU core'."

# The Traverser

$t \in T$



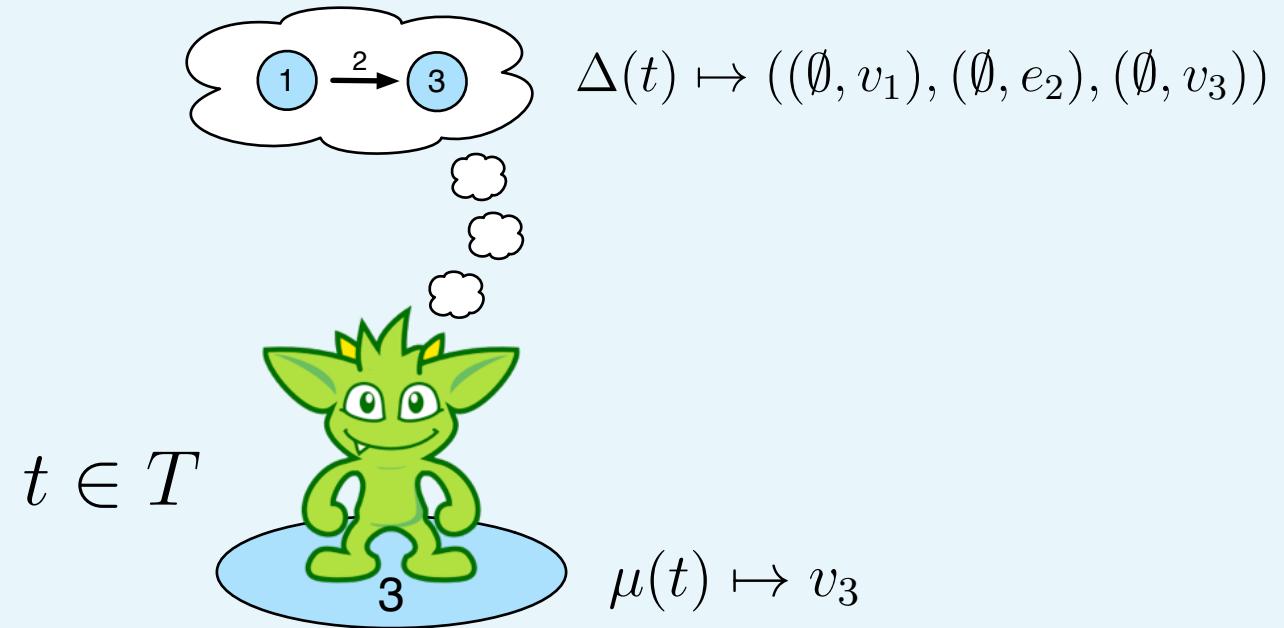
"The traverser  $t$  is in a set of traversers  $T$ ."

# The Traverser



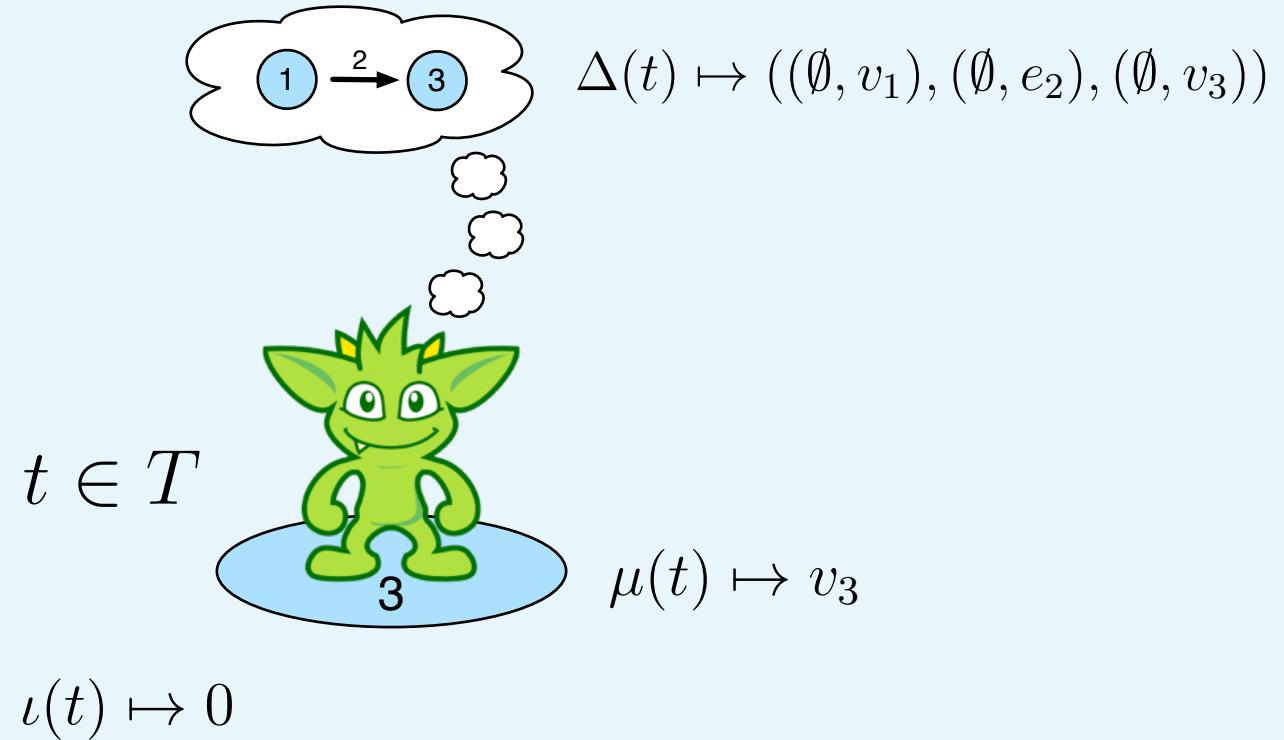
"The traverser's current graph location is vertex 3."

# The Traverser



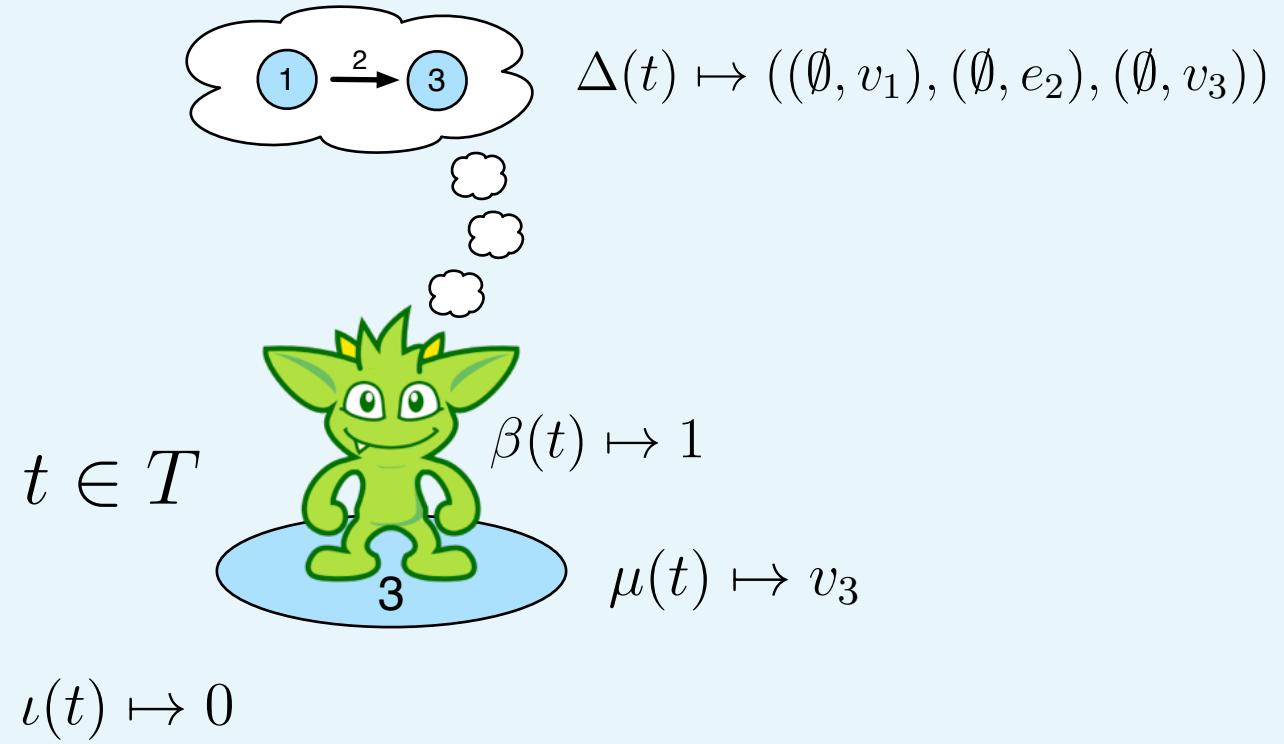
"The traverser may remember his history in the graph."

# The Traverser



"The traverser records how many times he has been through a loop sequence."

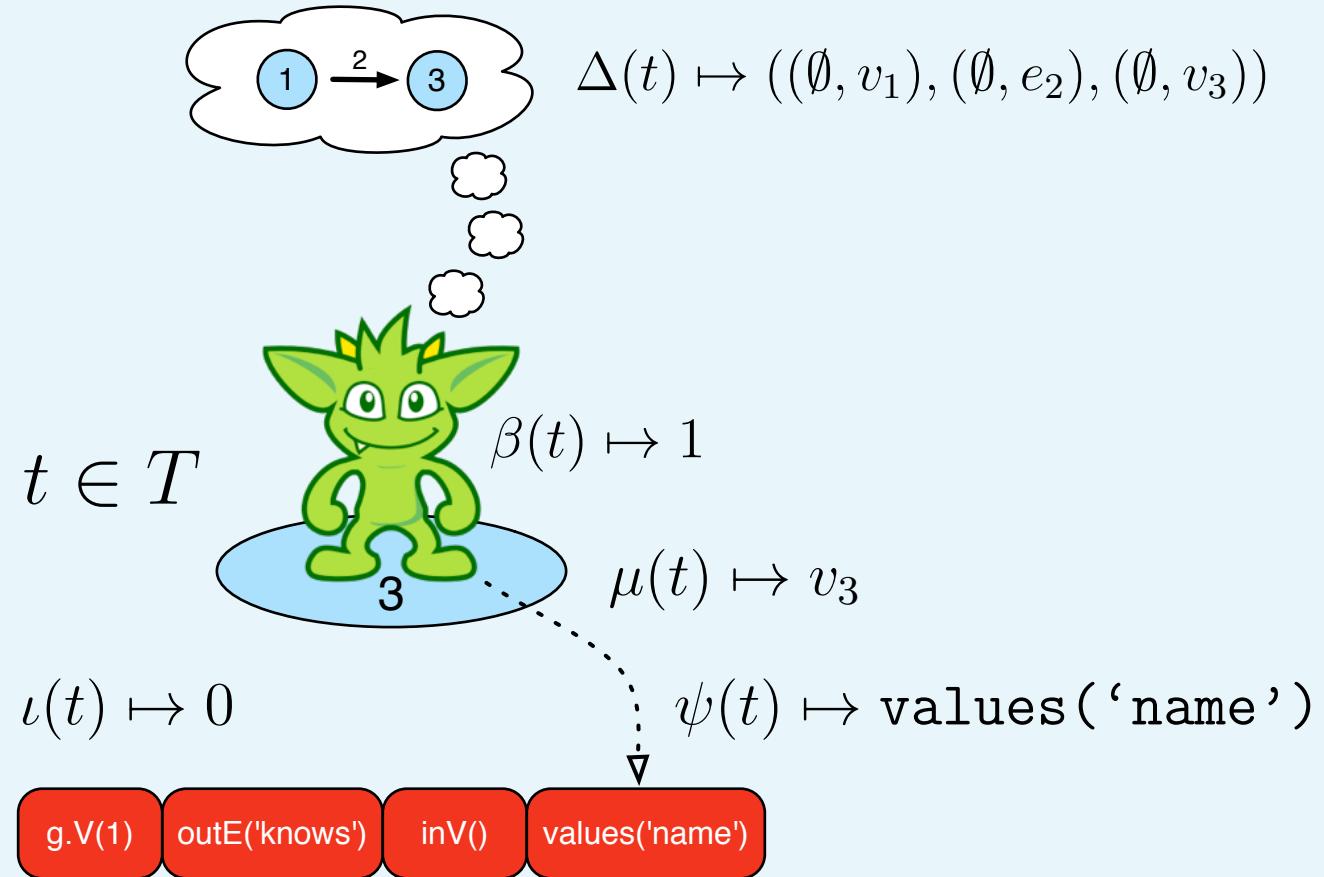
# The Traverser



*"It's what you would do if you thought in terms of "energy flows." "*

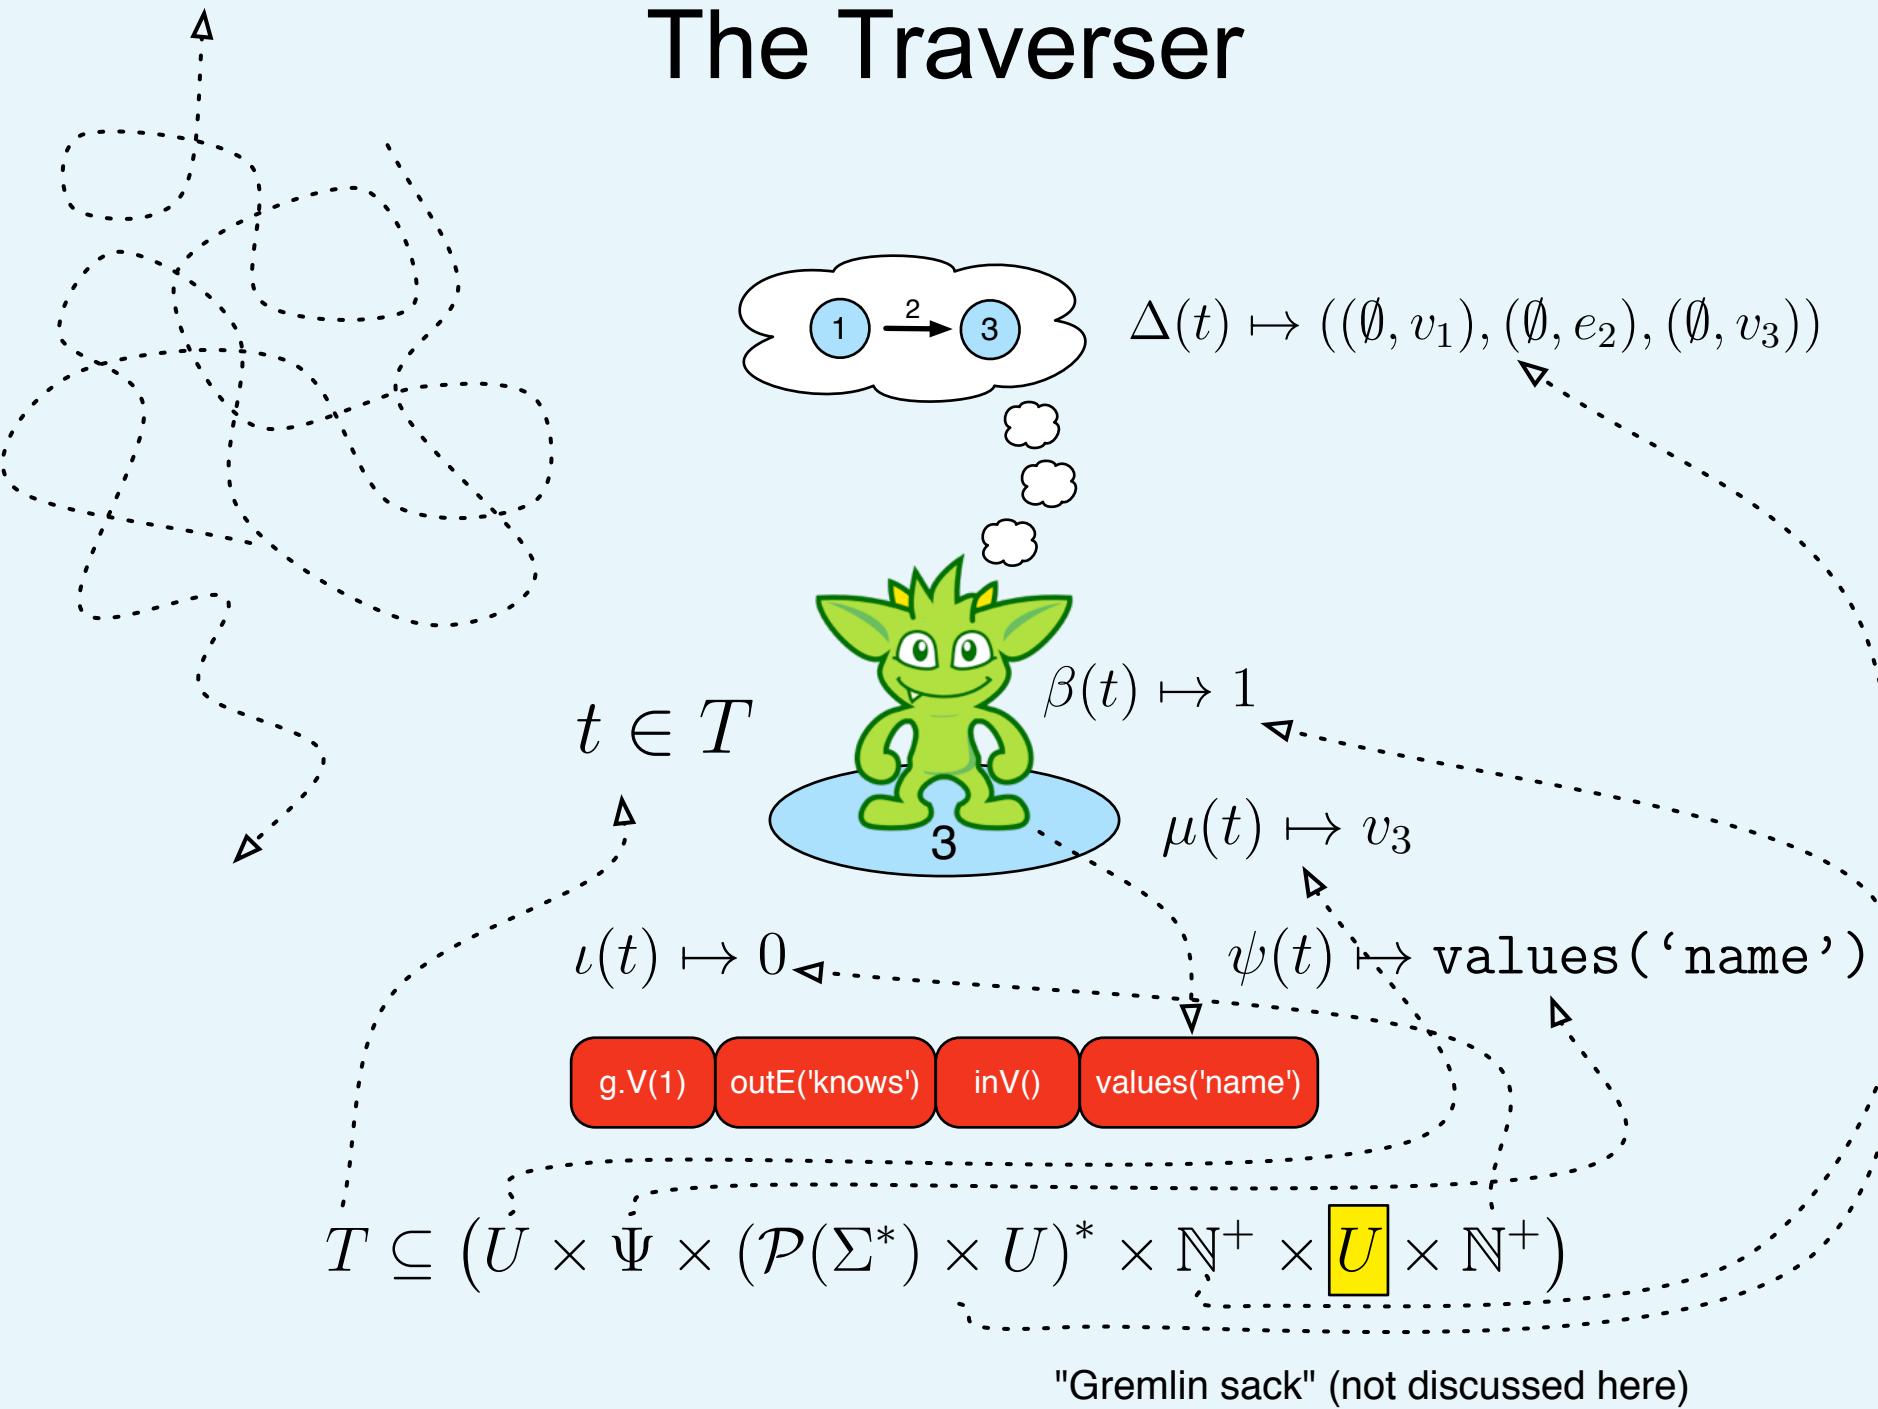
"The traverser can represent more than one traverser."

# The Traverser

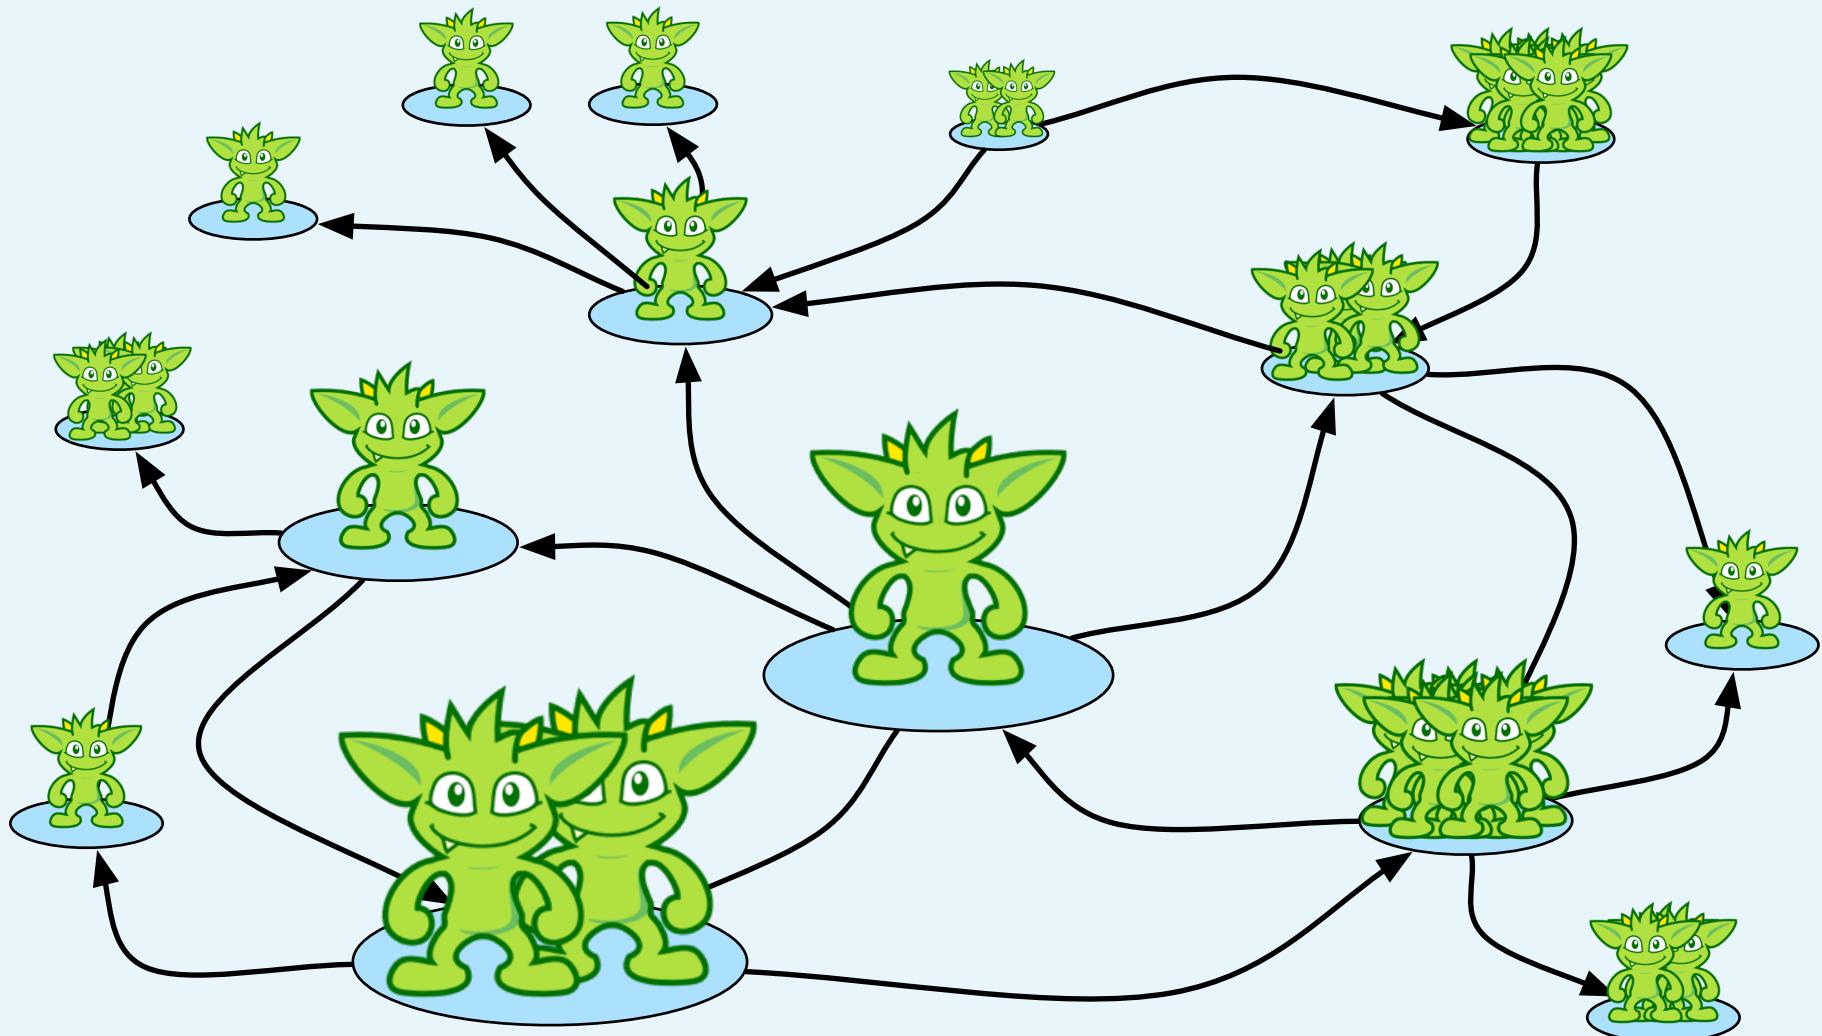


"The traverser has a reference to his location in the traversal (program counter)."

# The Traverser



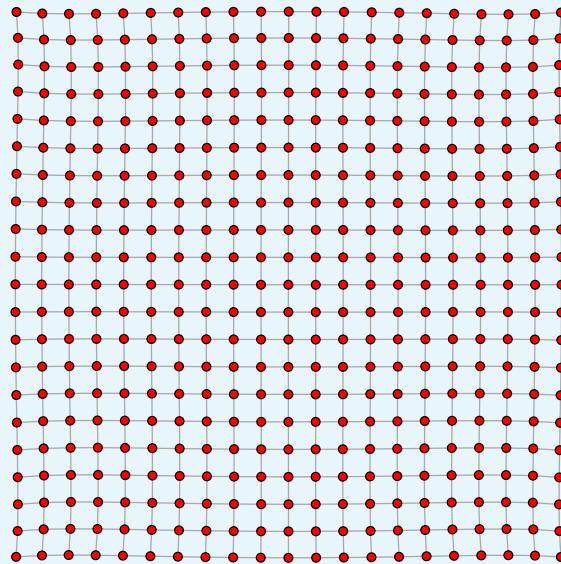
# The Traverser



Combinatoric  
explosion

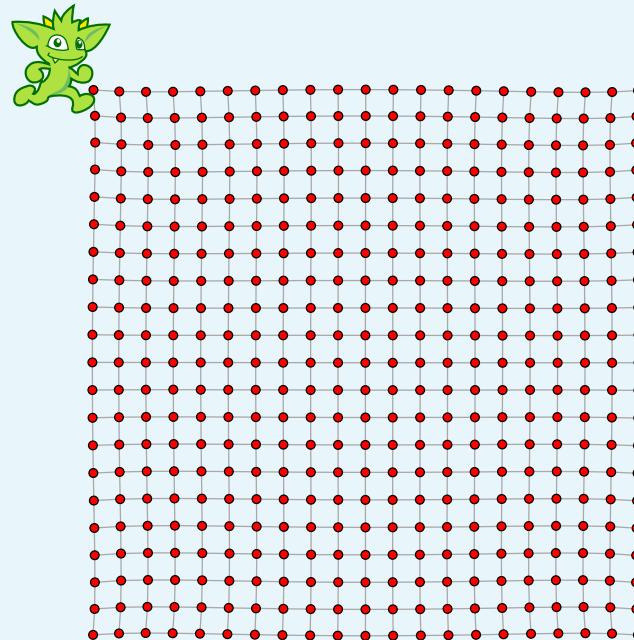
"Numerous (many many many) traversers are spawned during a traversal (computation)."  
many  
many

# The Traverser



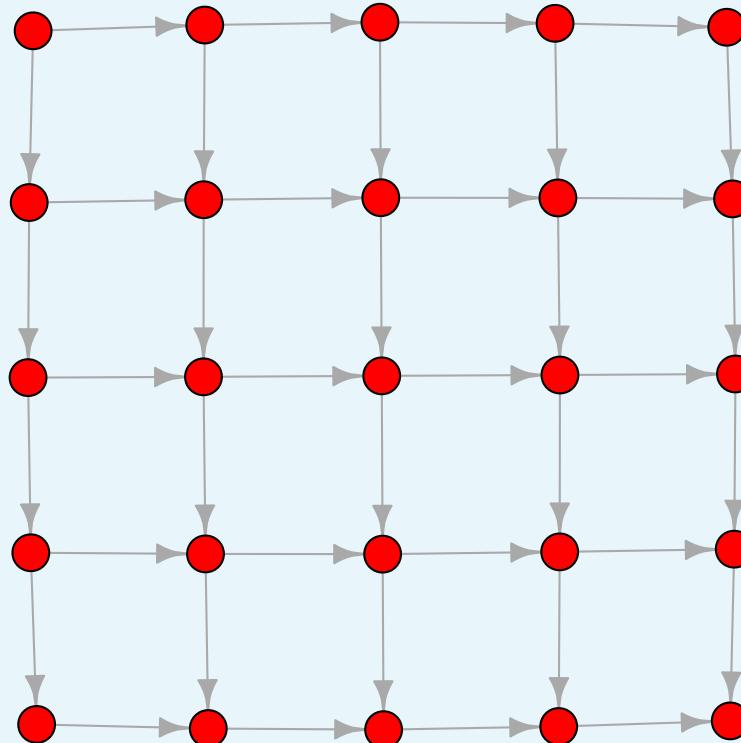
"A 21x21 vertex lattice -- 20x20 edges."

# The Traverser



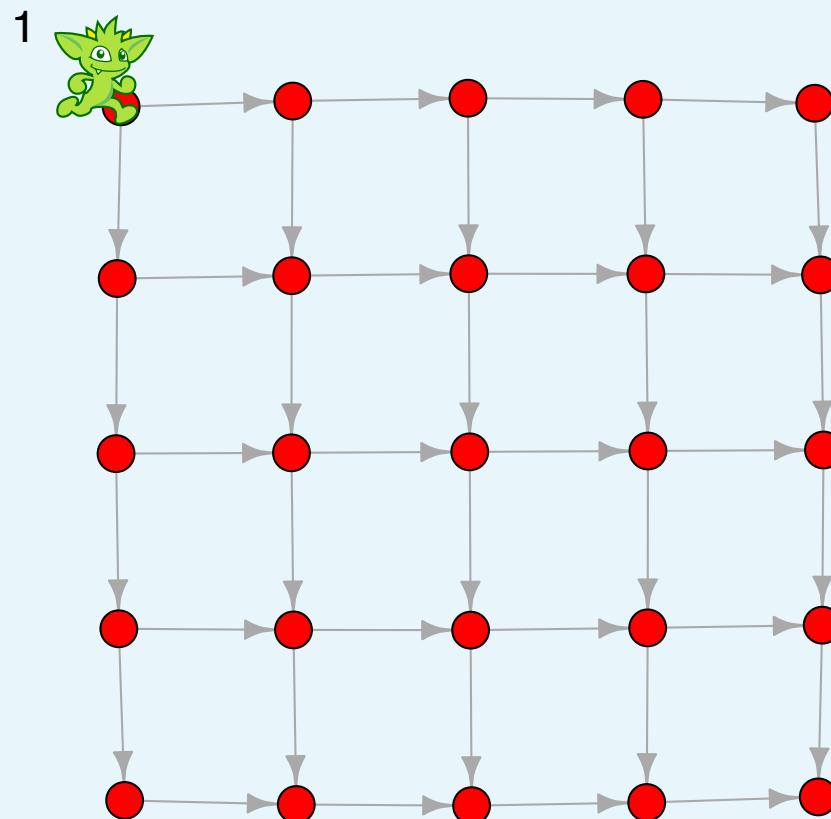
"A Gremlin starts at (0,0) and can only go right and down to reach (20,20).  
How many traversers (paths) ultimately reach (20,20)?"

# The Traverser



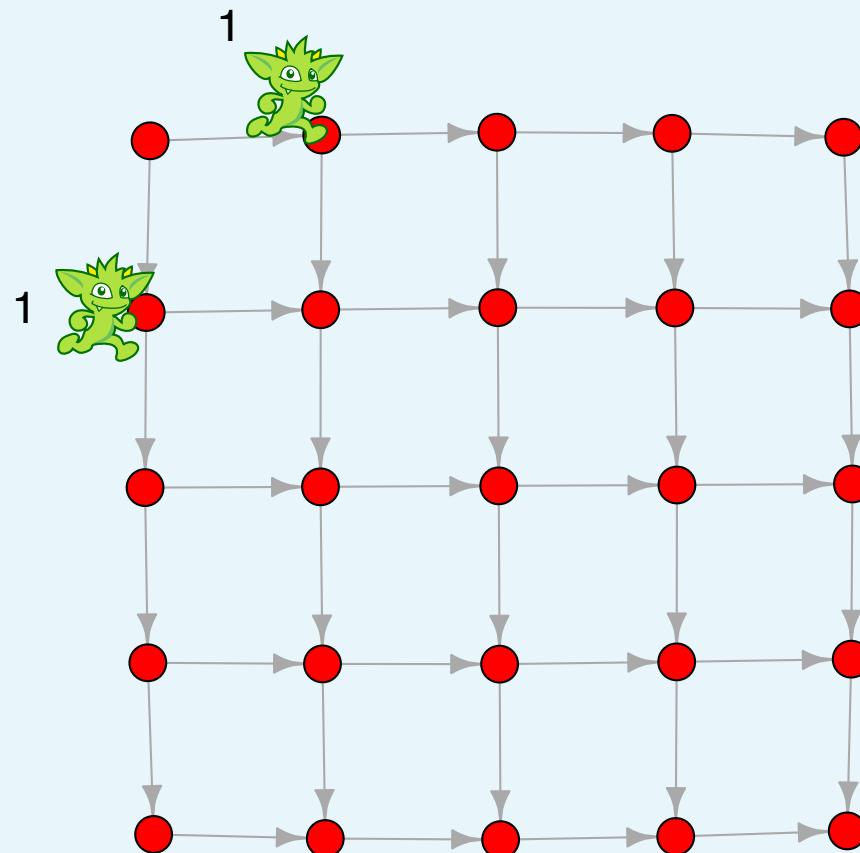
"For the sake of demonstration, lets look at a 5x5 lattice (4x4 edges)."

# The Traverser



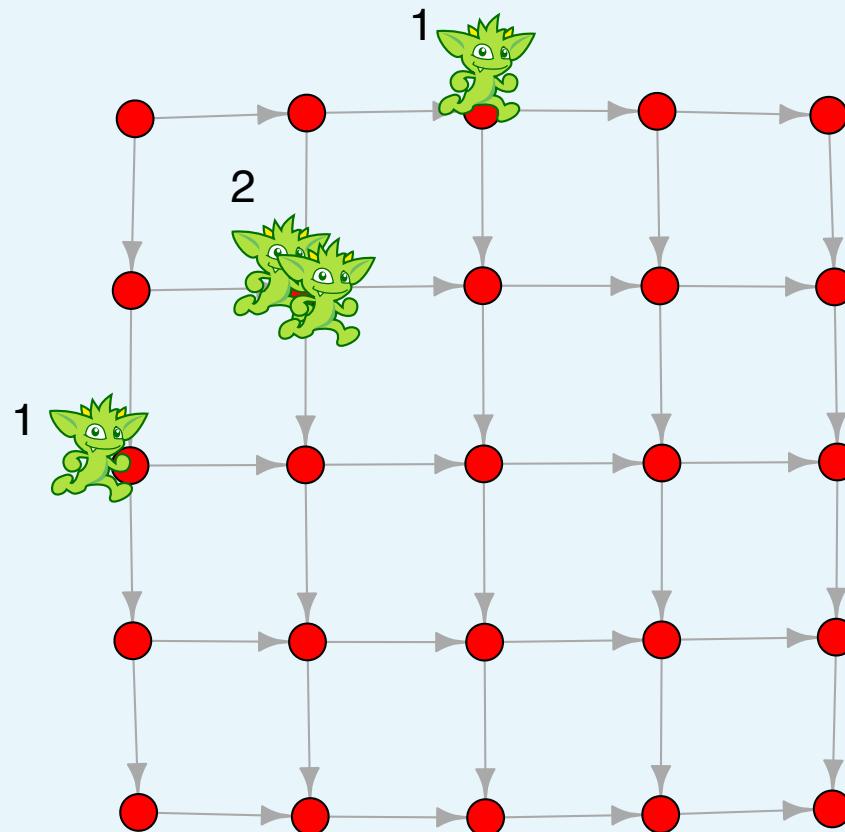
"Total = 1"

# The Traverser



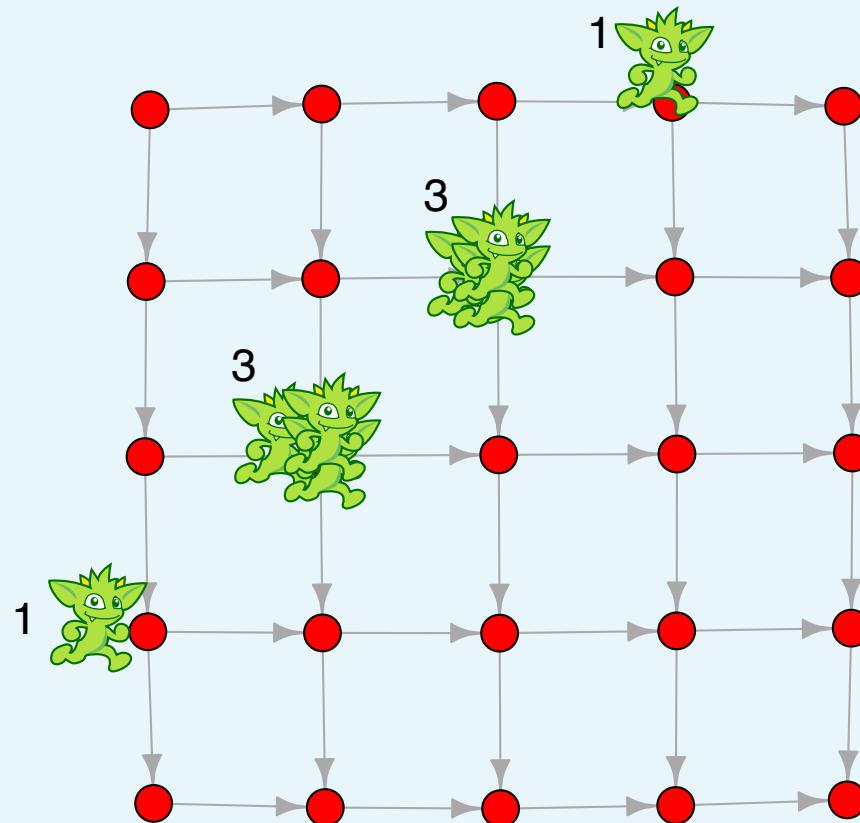
"Total = 2"

# The Traverser



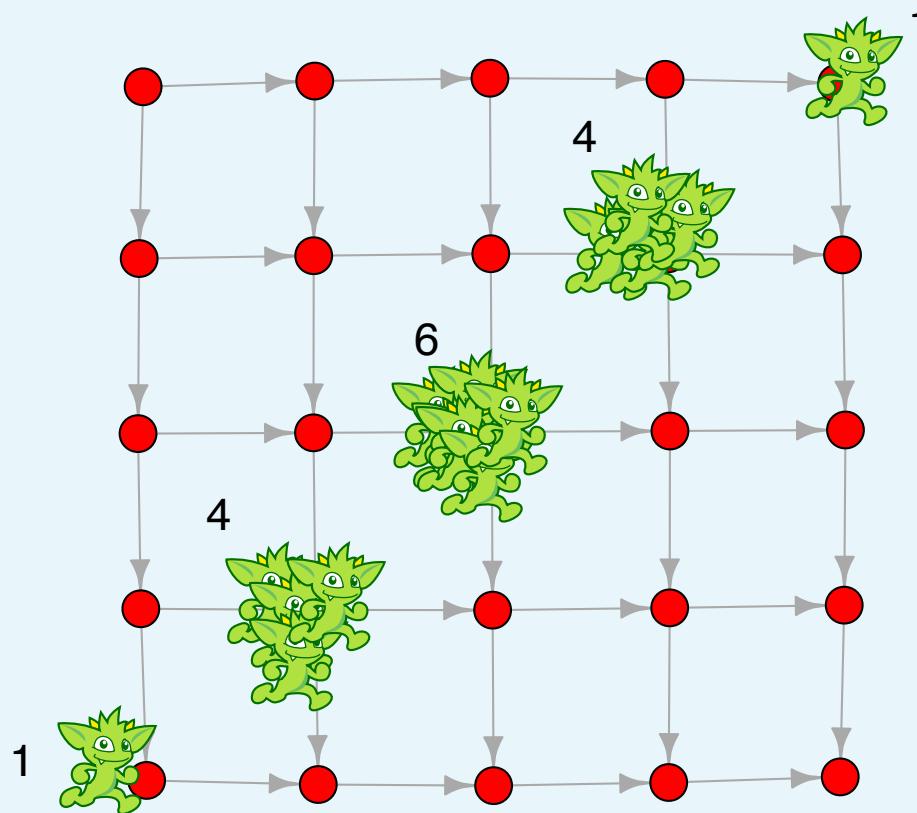
"Total = 4"

# The Traverser

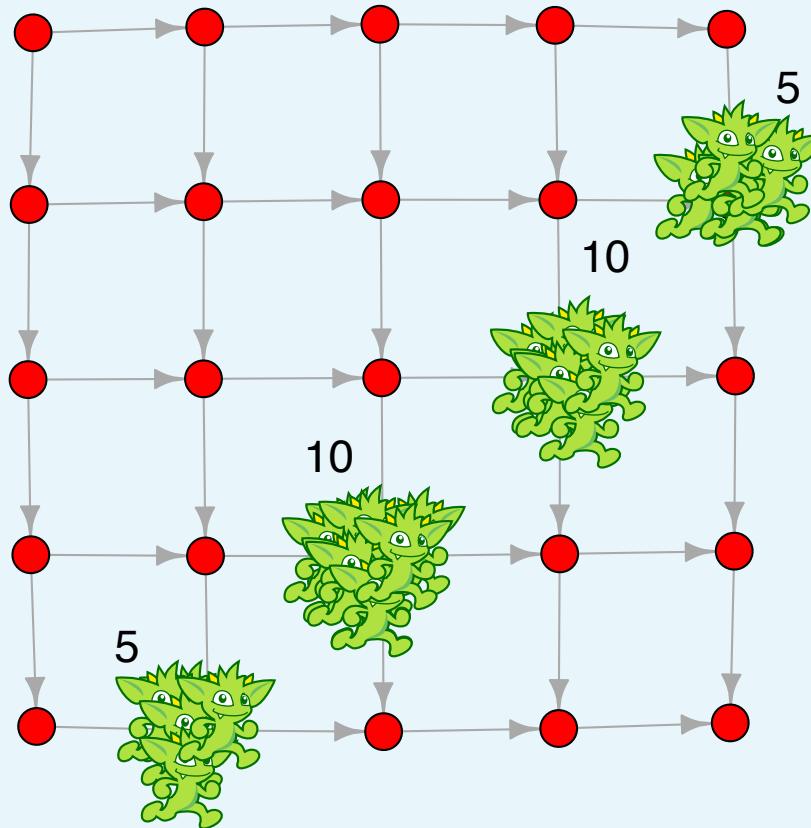


"Total = 8"

# The Traverser

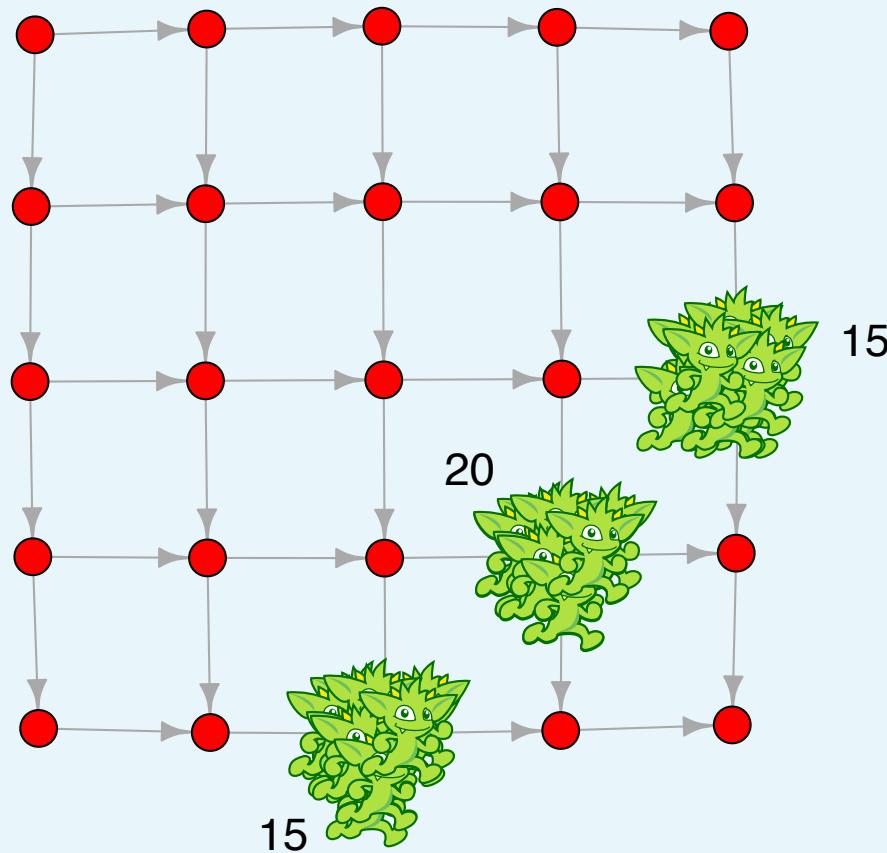


# The Traverser

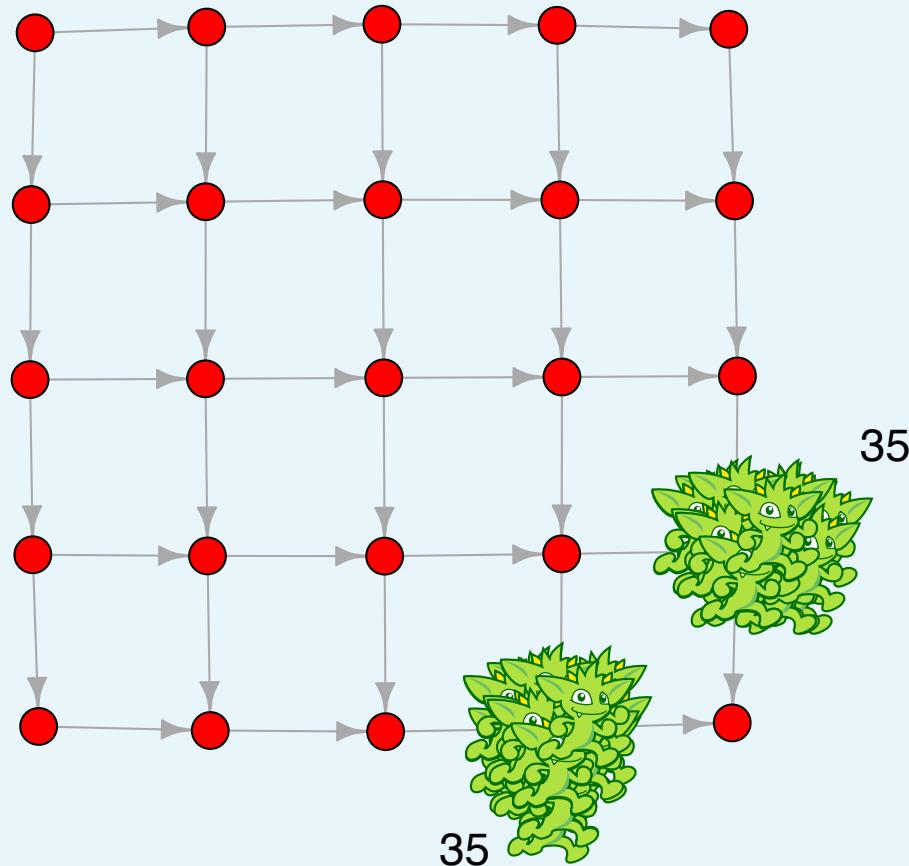


"Total = 20"

# The Traverser

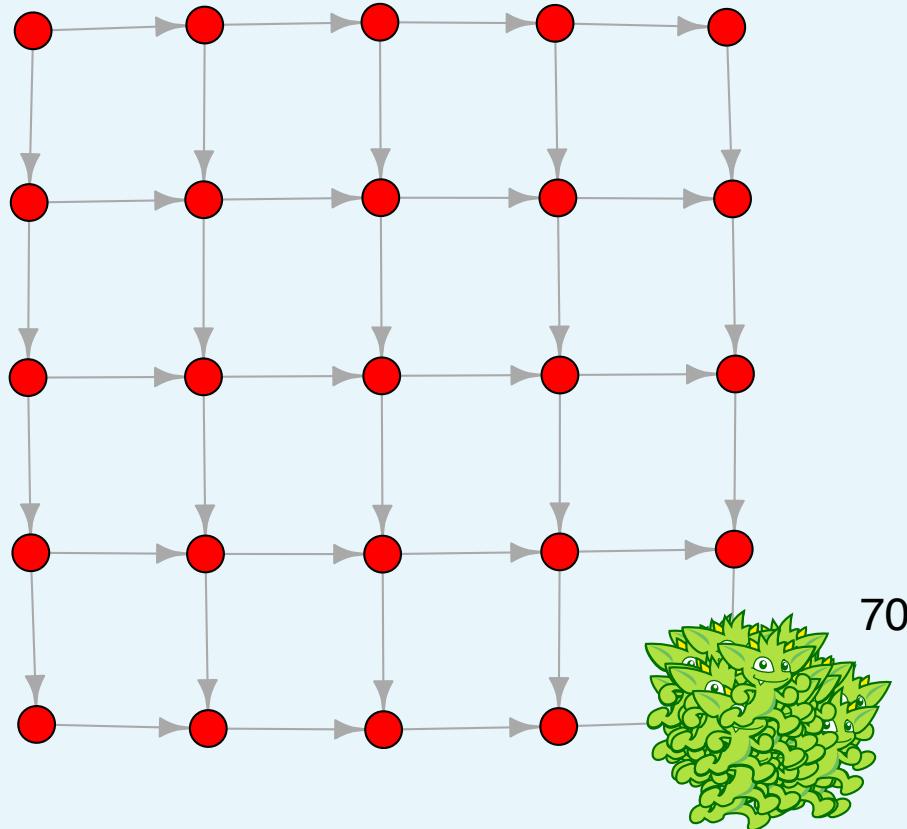


# The Traverser



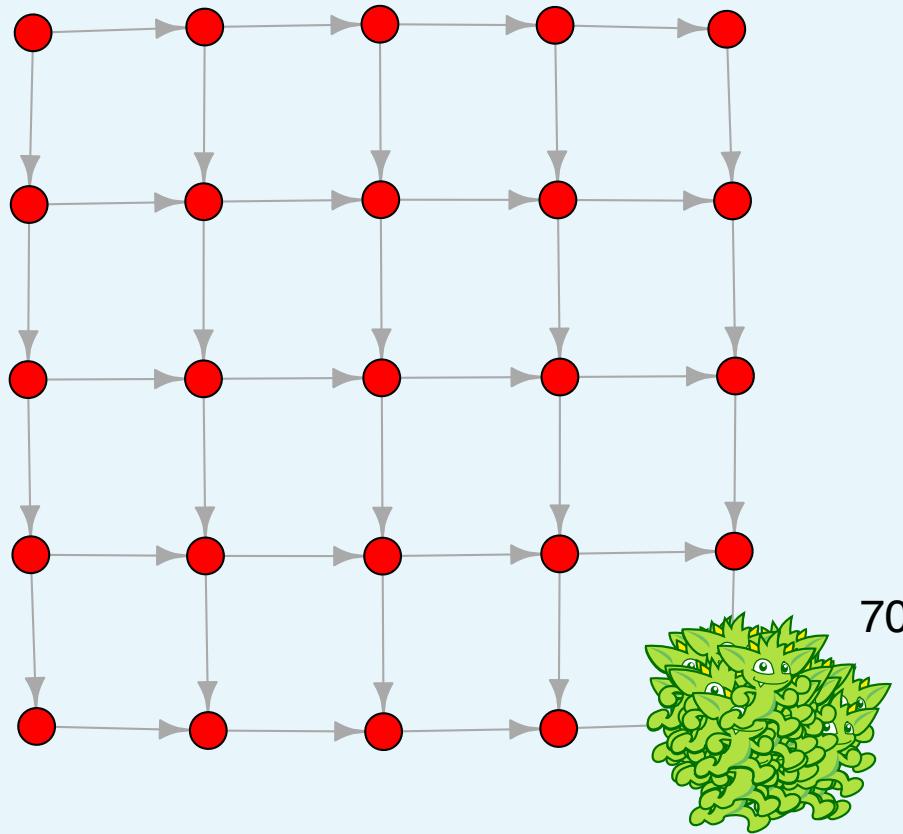
"Total = 70"

# The Traverser



"Total = 70"

# The Traverser

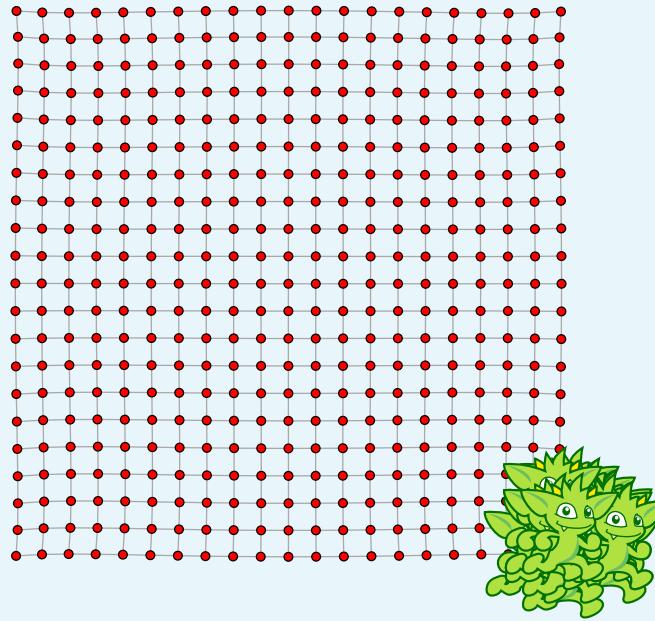


8 choose 4 = 70

$$\binom{2n}{n} = \frac{(2n)!}{(n!)^2} = \frac{(2n)(2n-1)(2n-2)\dots 1}{(n(n-1)(n-2)\dots 1)^2}$$

*n = number of edges on a side*

# The Traverser

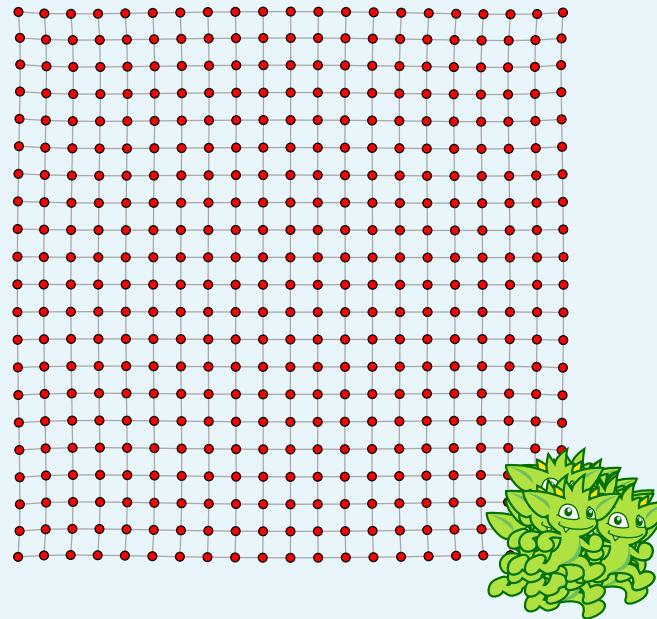


$$\binom{2n}{n} = \frac{(2n)!}{(n!)^2} = \frac{(2n)(2n-1)(2n-2)\dots 1}{(n(n-1)(n-2)\dots 1)^2}$$

40 choose 20

"Analytically, ~137 billion traversers will exist at (20,20)."

# The Traverser

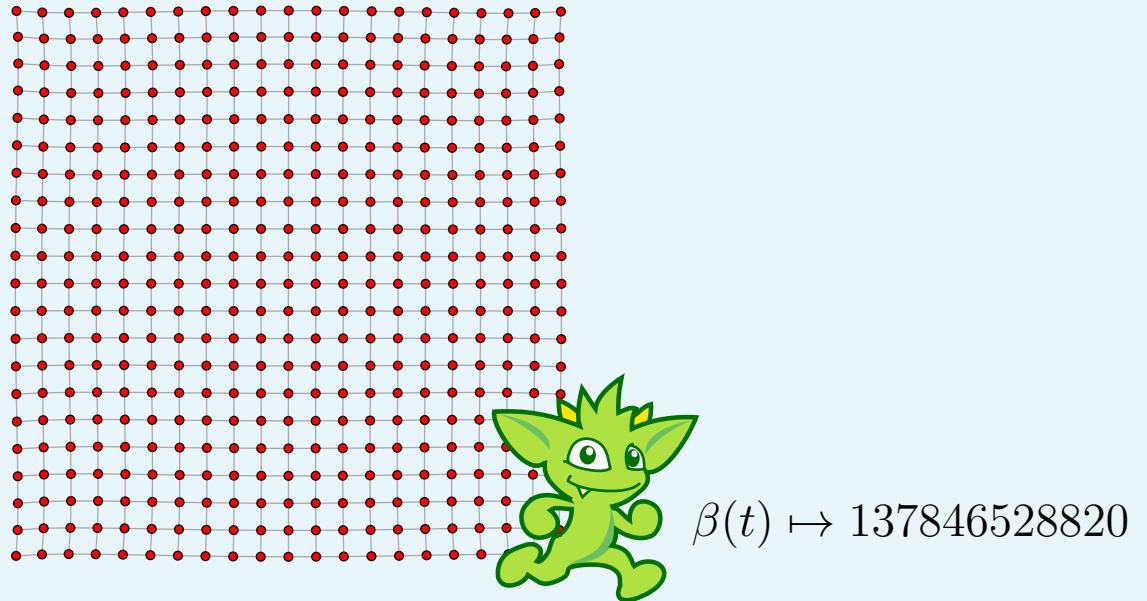


```
gremlin> g.V(0).repeat(out('down','right')).times(40).count()  
==>137846528820
```

```
gremlin>
```

"That was fast... The Gremlin machine spawned 137 billion traversers?!"

# The Traverser



```
gremlin> g.V(0).repeat(out('down','right')).times(40).count()  
==>137846528820  
gremlin> clock(100){g.V(0).repeat(out('down','right')).times(40).count().next()}  
==>0.8890942  
less than a millisecond
```

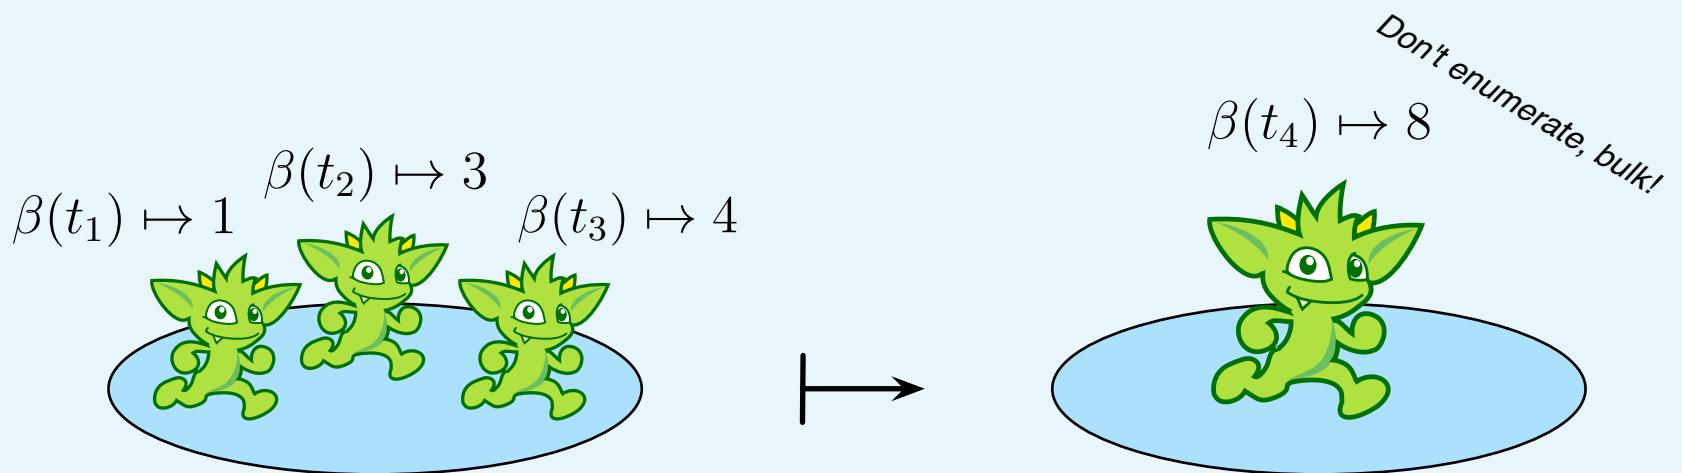
"No. At vertex 400 (20,20), there exists 1 traverser with a bulk of ~137 billion."

Bulking is a crucial component of OLAP when near trillion edge graphs are moving traversers between machines.

# The Traverser

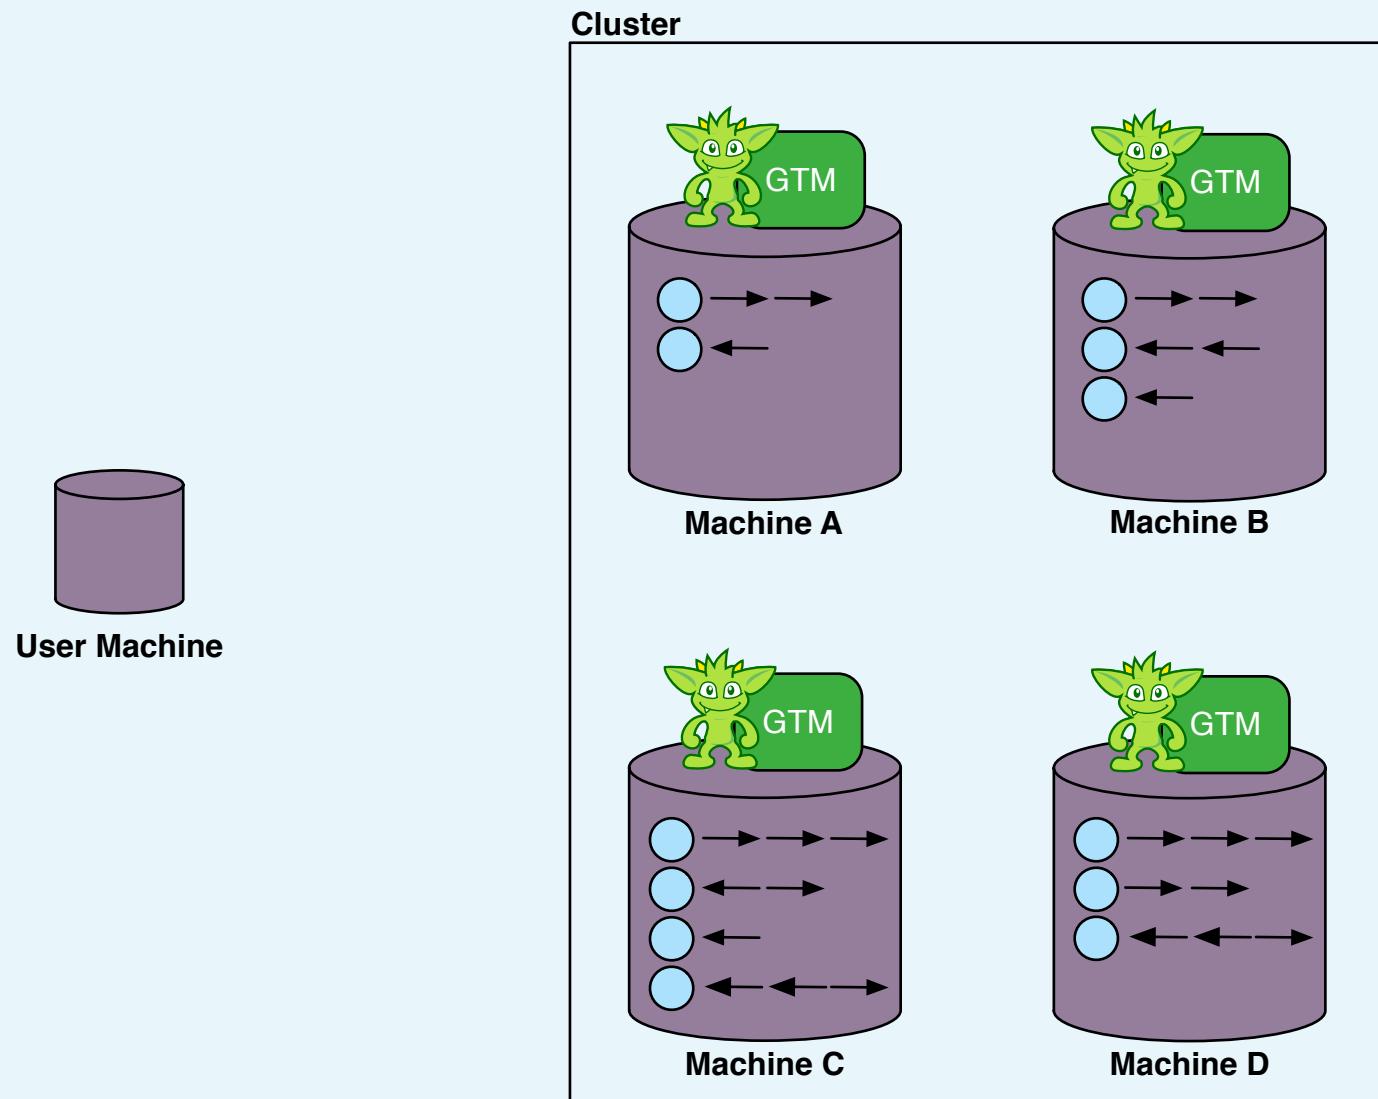
## Traverser Equivalence Class

$[t] = \{t' \in T \mid \mu(t) = \mu(t') \wedge \psi(t) = \psi(t') \wedge \Delta(t) = \Delta(t') \wedge \iota(t) = \iota(t')\}.$	<i>graph location</i>	<i>data location</i>
$\psi(t) = \psi(t')$	<i>traversal location</i>	<i>program counter</i>
$\Delta(t) = \Delta(t')$	<i>path history</i>	<i>MAY NOT BE REQUIRED</i>
	<i>same loop counter</i>	<i>recursion of program counter</i>



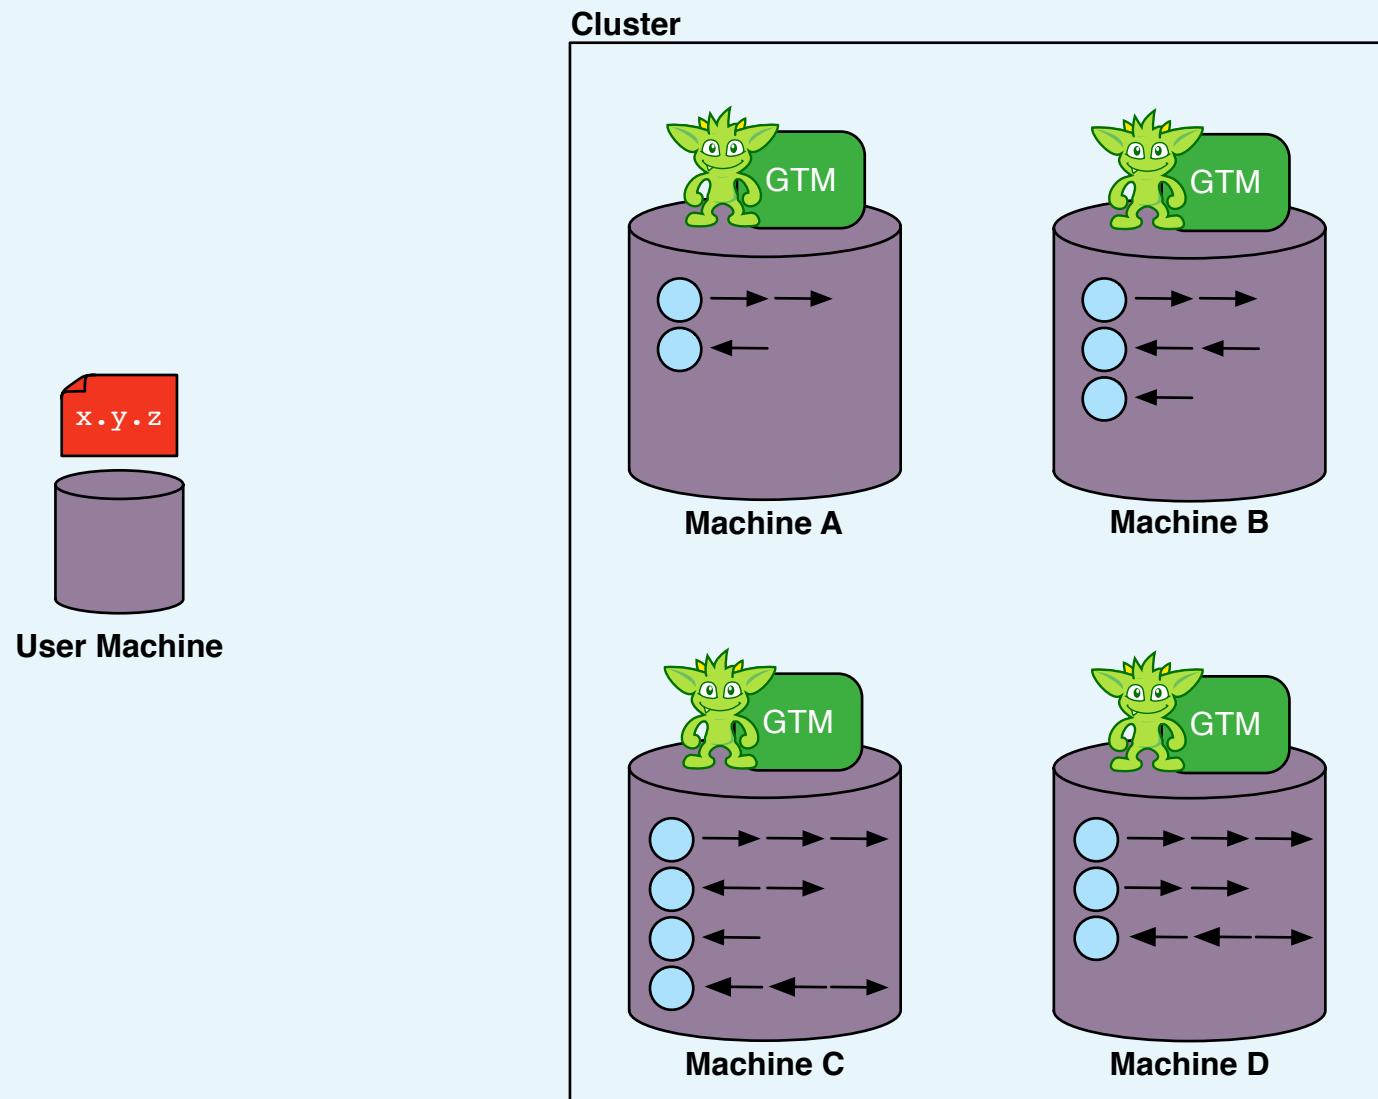
"If 2 traversers (each with bulk 1) are at the same location in the graph and have similar other properties, make 1 traverser with a bulk of 2."

# The Traverser



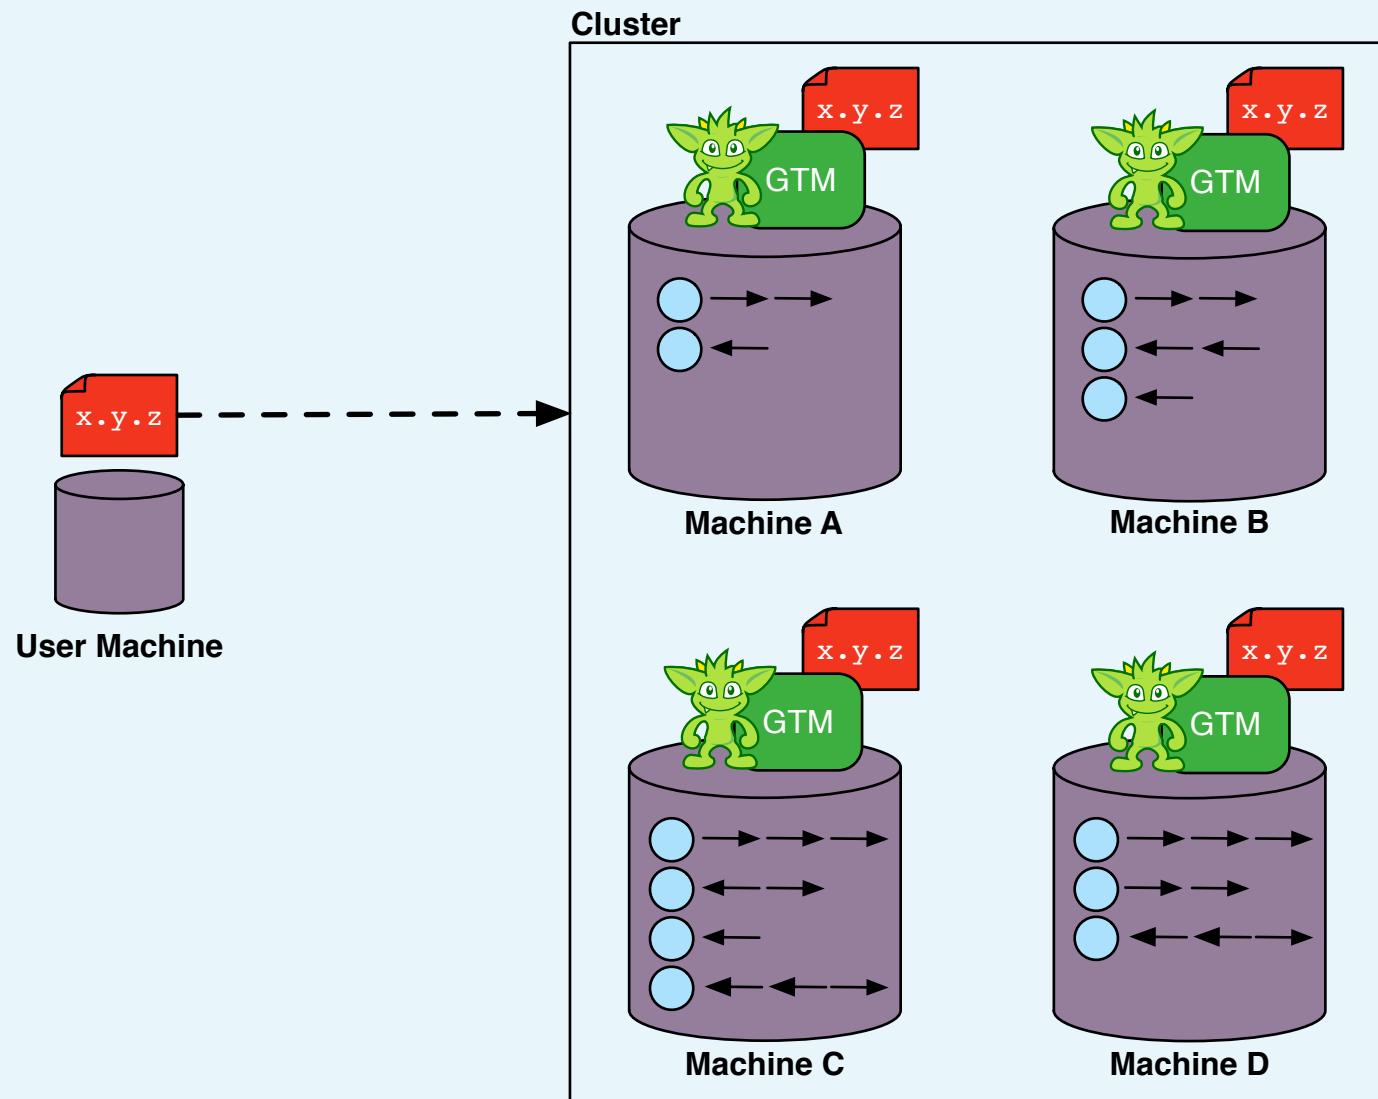
"In Gremlin OLAP, the graph is represented across a distributed system (e.g. Hadoop) as a partitioned adjacency list."

# The Traverser



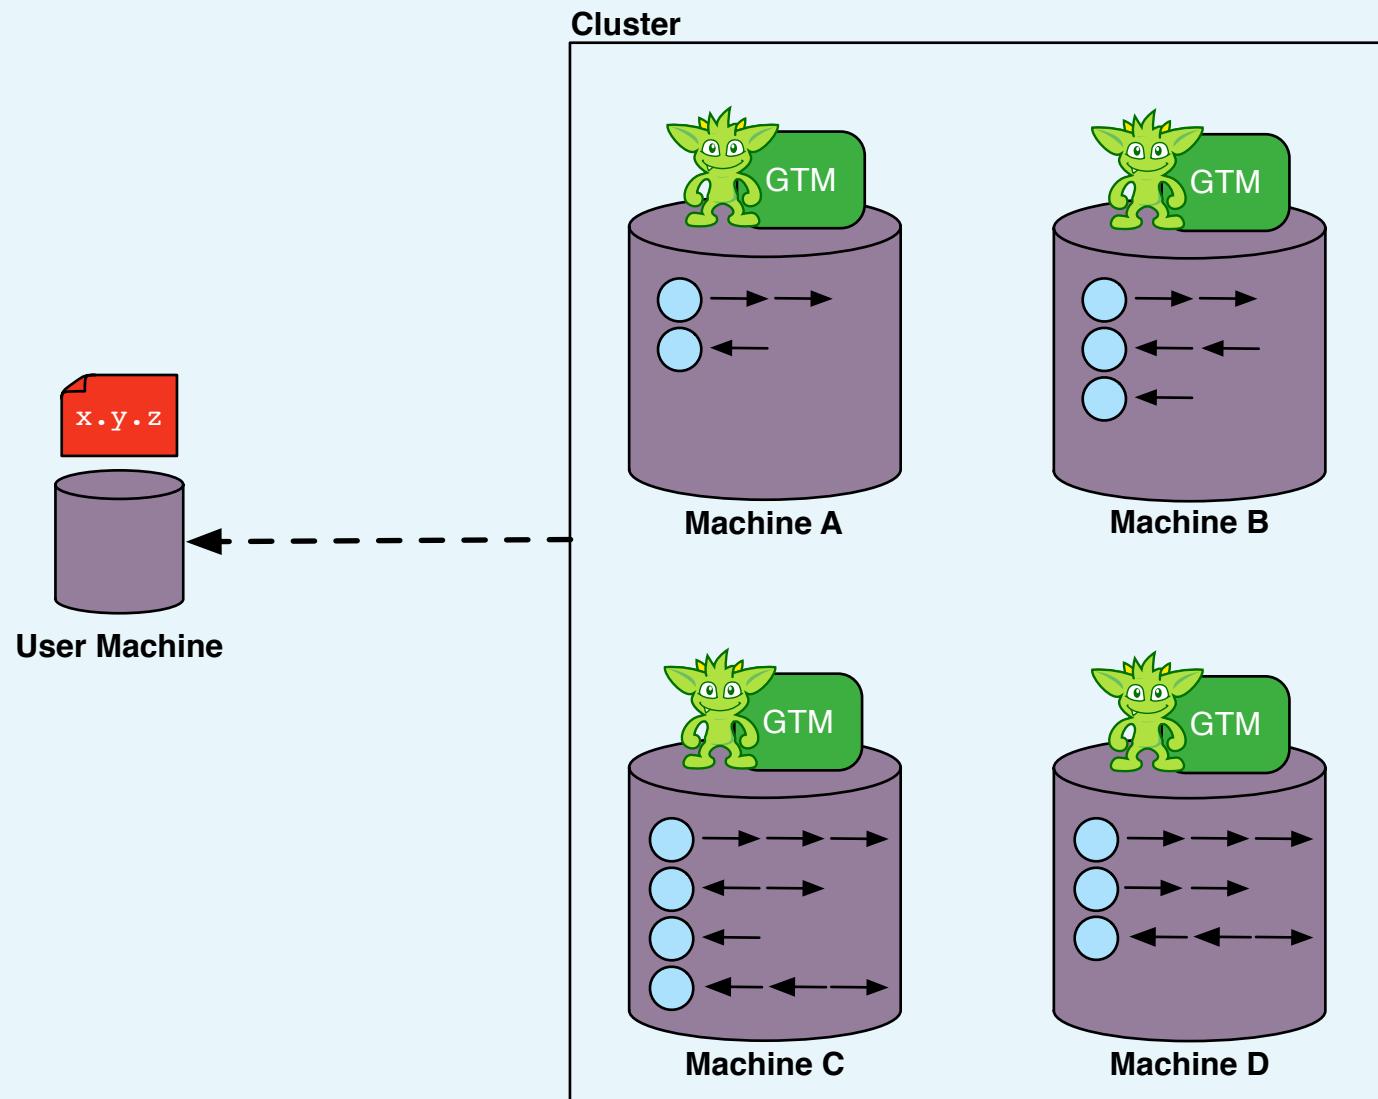
"The user creates a traversal (compiled from any language)."

# The Traverser



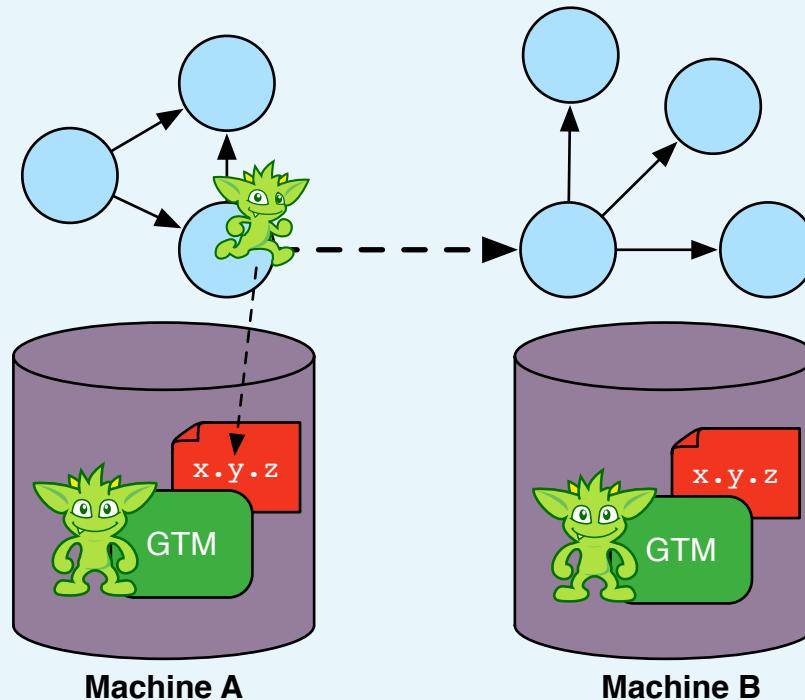
"The compiled traversal is sent to every machine in the cluster which has a piece of the graph and a Gremlin traversal machine."

# The Traverser



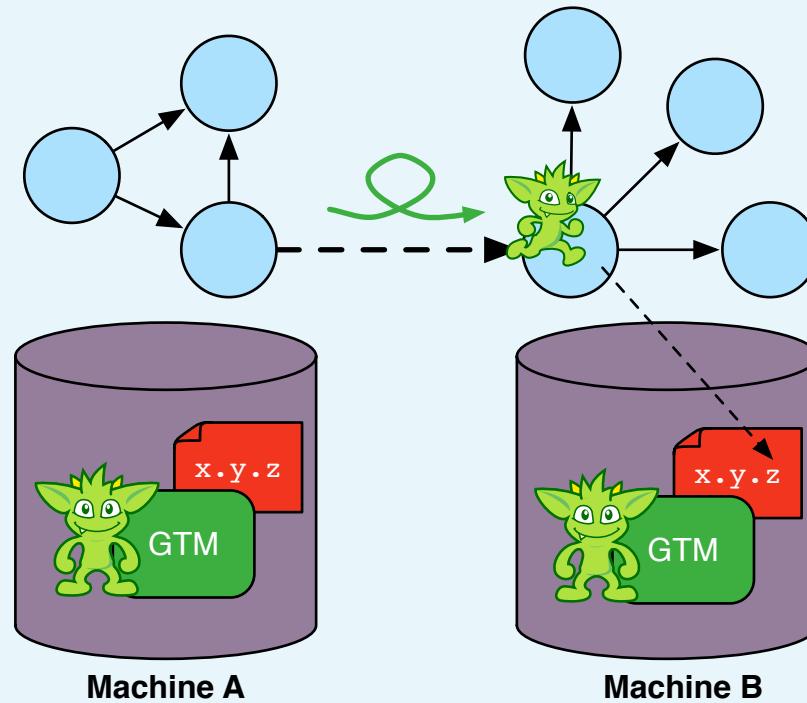
"When the graph computation is complete, a reference to the result is returned. For instance, in HadoopGremlin, HDFS stores both the graph and sideEffect data."

# The Traverser



"The OLAP algorithm is the classic Bulk Synchronous Parallel algorithm popularized by Google Pregel and Apache Giraph."

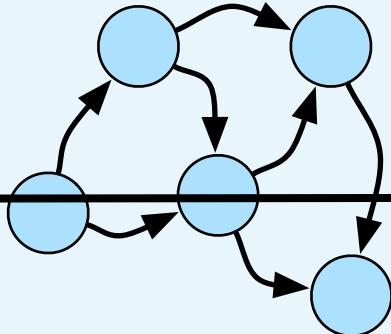
# The Traverser



"The messages are traversers. Given the traverser equivalence class [t], they can be bulked."

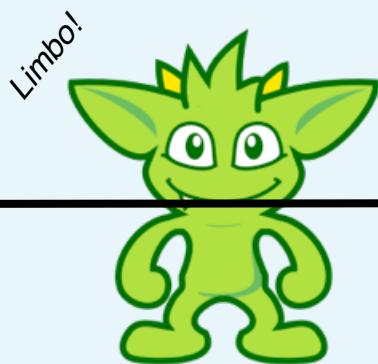
"Its just "complex energy."

# The Machine Components



The Graph

The Data



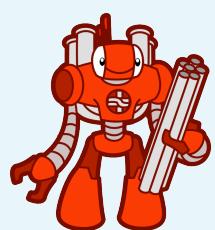
The Traverser

The CPU

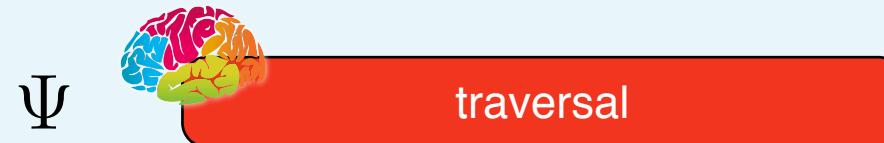


The Traversal

The Program



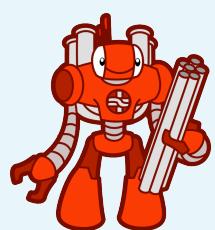
# The Traversal



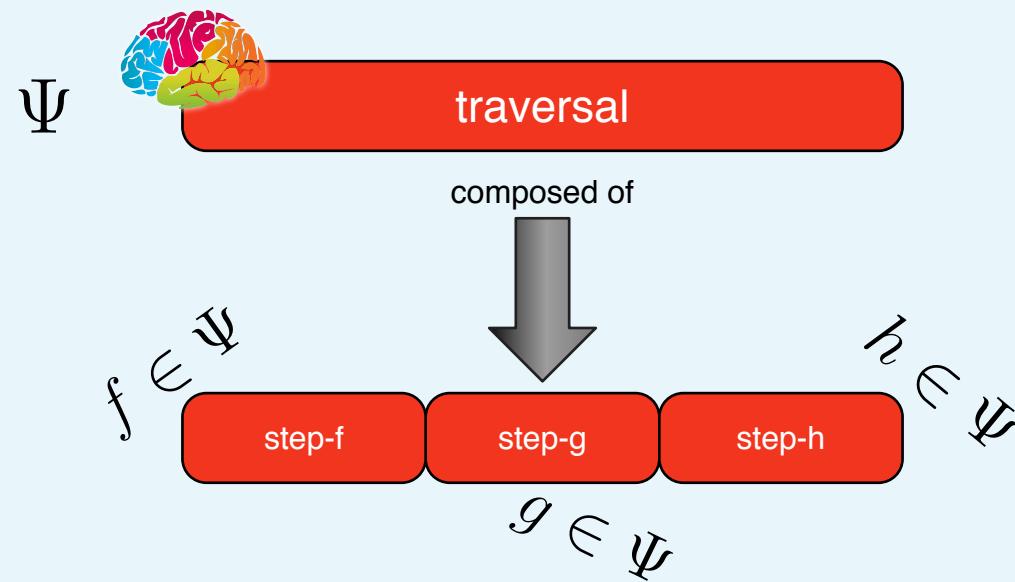
$\Psi$

traversal

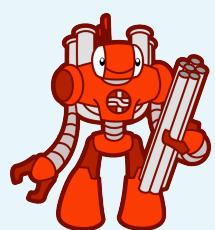
"The traversal is the 'software program'."



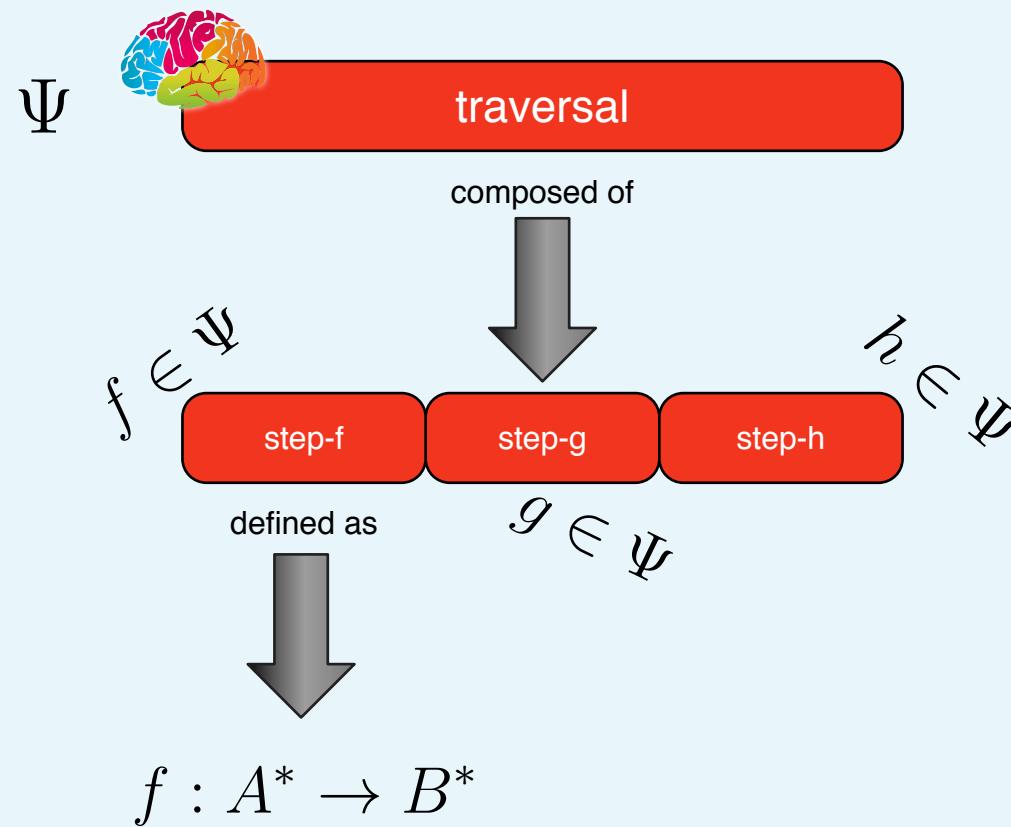
# The Traversal



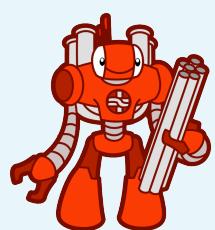
"The steps are the 'instructions'."



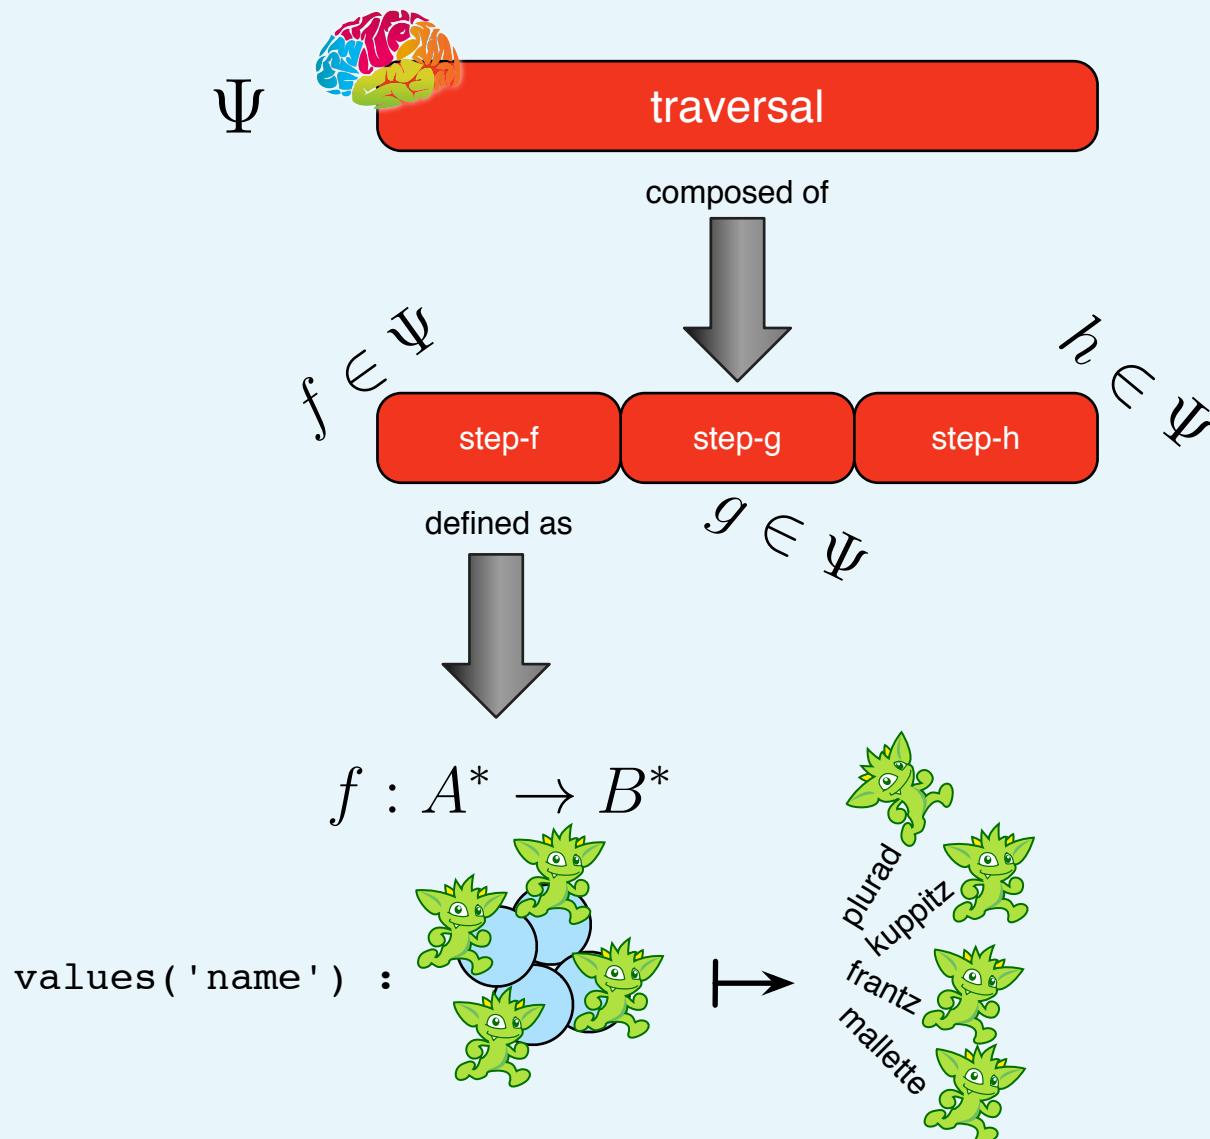
# The Traversal



"Step  $f$  maps a stream (multi-set) of traversers located at objects of type  $A$  to a stream (multi-set) of traversers located at objects of type  $B$ ."



# The Traversal



"Step `values('name')` maps a stream (multi-set) of traversers located at vertices to a stream (multi-set) of traversers located at strings."

# The Traversal

$\text{map} : A^* \rightarrow B^*$



"For a traverser at  $a$ , move the traverser to an object at  $b$ ."

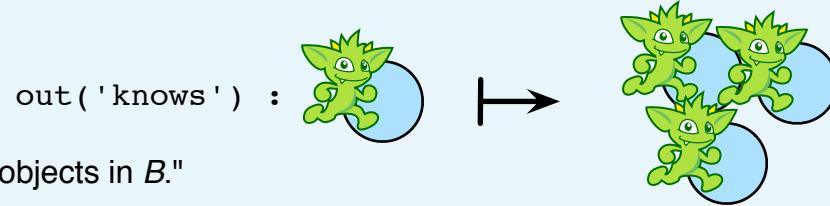
# The Traversal

$\text{map} : A^* \rightarrow B^*$



"For a traverser at  $a$ , move the traverser to an object at  $b$ ."

$\text{flatMap} : A^* \rightarrow B^*$



"For a traverser at  $a$ , clone the traverser across multiple objects in  $B$ ."

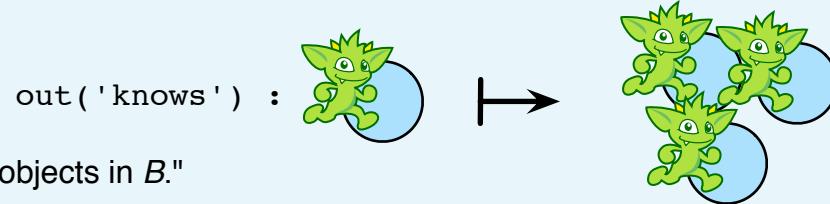
# The Traversal

$\text{map} : A^* \rightarrow B^*$



"For a traverser at  $a$ , move the traverser to an object at  $b$ ."

$\text{flatMap} : A^* \rightarrow B^*$



"For a traverser at  $a$ , clone the traverser across multiple objects in  $B$ ."

$\text{filter} : A^* \rightarrow A^*$



"For a traverser at  $a$ , either kill the traverser or leave him alone."

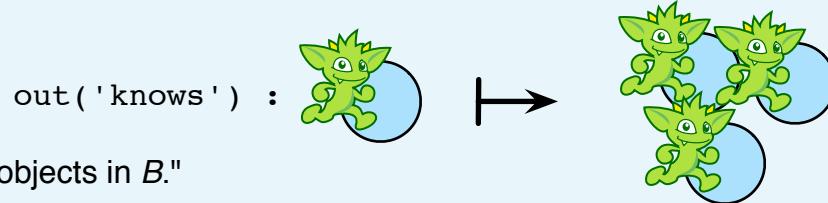
# The Traversal

$\text{map} : A^* \rightarrow B^*$



"For a traverser at  $a$ , move the traverser to an object at  $b$ ."

$\text{flatMap} : A^* \rightarrow B^*$



"For a traverser at  $a$ , clone the traverser across multiple objects in  $B$ ."

$\text{filter} : A^* \rightarrow A^*$



"For a traverser at  $a$ , either kill the traverser or leave him alone."

$\text{sideEffect} : A^* \rightarrow_x A^*$



"For a traverser at  $a$ , leave him alone though manipulate some data structure  $x$ ."

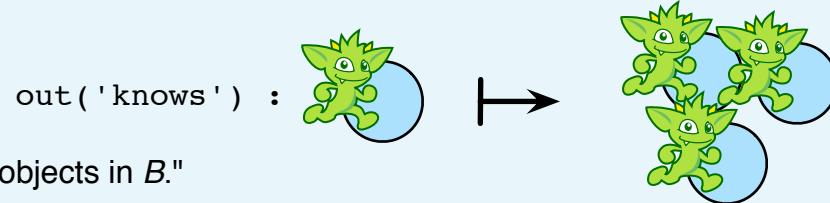
# The Traversal

$\text{map} : A^* \rightarrow B^*$



"For a traverser at  $a$ , move the traverser to an object at  $b$ ."

$\text{flatMap} : A^* \rightarrow B^*$



"For a traverser at  $a$ , clone the traverser across multiple objects in  $B$ ."

$\text{filter} : A^* \rightarrow A^*$



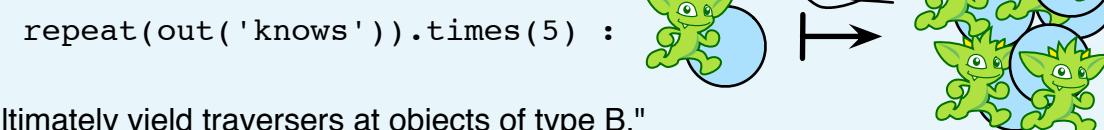
"For a traverser at  $a$ , either kill the traverser or leave him alone."

$\text{sideEffect} : A^* \rightarrow_x A^*$



"For a traverser at  $a$ , leave him alone though manipulate some data structure  $x$ ."

$\text{branch} : A^* \rightarrow^b B^*$



"For a traverser at  $a$ , choose some internal branch  $b$  to ultimately yield traversers at objects of type  $B$ ."

# The Traversal

map	flatMap	filter	sideEffect	branch
LabelStep	VertexStep	AndStep	GroupCountStep	BranchStep
IdStep	EdgeVertexStep	CoinStep	GroupStep	ChooseStep
PathStep	CoalesceStep	CyclicPathStep	StoreStep	LocalStep
SackStep	...	DedupStep	AggregateStep	RepeatStep
SelectStep		DropStep	InjectStep	UnionStep
MatchStep		FilterStep	SubgraphStep	...
ConstantStep		HasStep	TreeStep	
...		IsStep	IdentityStep	
		NotStep	...	
		OrStep		
		RangeStep		
		...		

"This is the step library (instruction set) of the Gremlin graph traversal machine (virtual machine)."

# The Traversal

Linear Motif

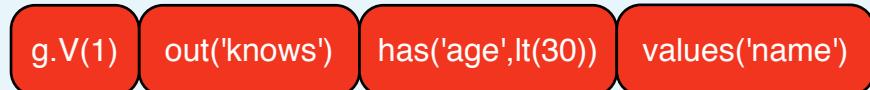
$f \circ g \circ h$

$f$

$g$

$h$

`g.V(1).out('knows').has('age',lt(30)).values('name')`



"The 'byte code' of Gremlin is a sequence of steps ...."

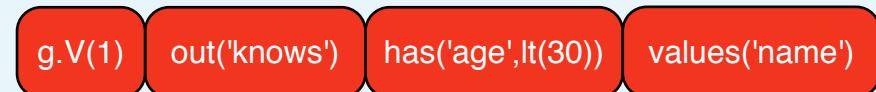
"

# The Traversal

## Linear Motif

$$f \circ g \circ h$$

```
g.V(1).out('knows').has('age',lt(30)).values('name')
```



## Nested Motif

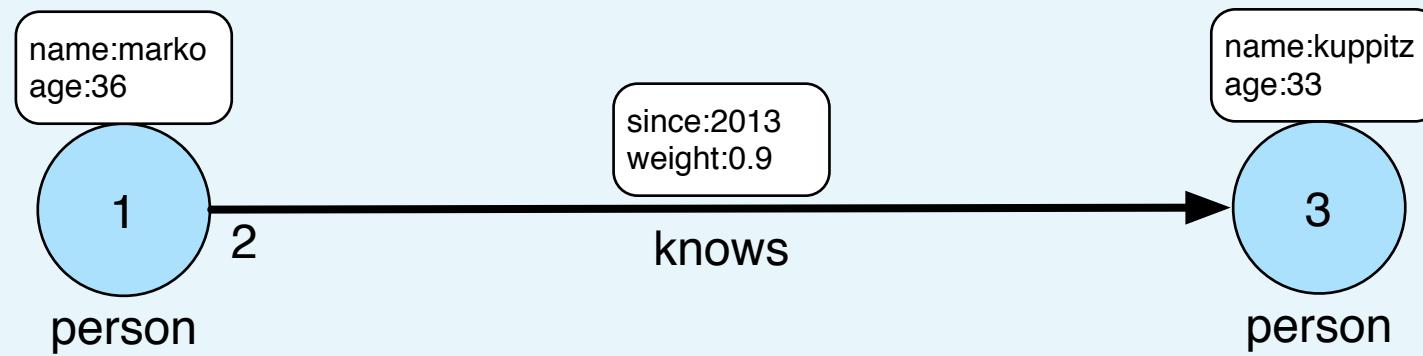
$$f(g \circ h) \circ k$$

```
g.V(1).repeat(out('knows').has('age',lt(30))).values('name')
```



"The 'byte code' of Gremlin is a sequence of steps with some steps nested within each other."

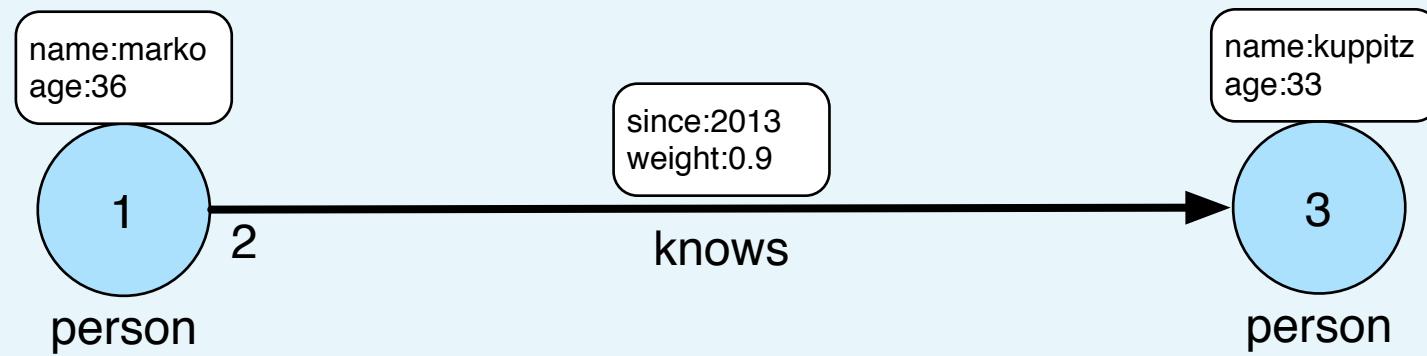
```
gremlin> graph  
==>tinkergraph[vertices:2 edges:1]  
gremlin>
```



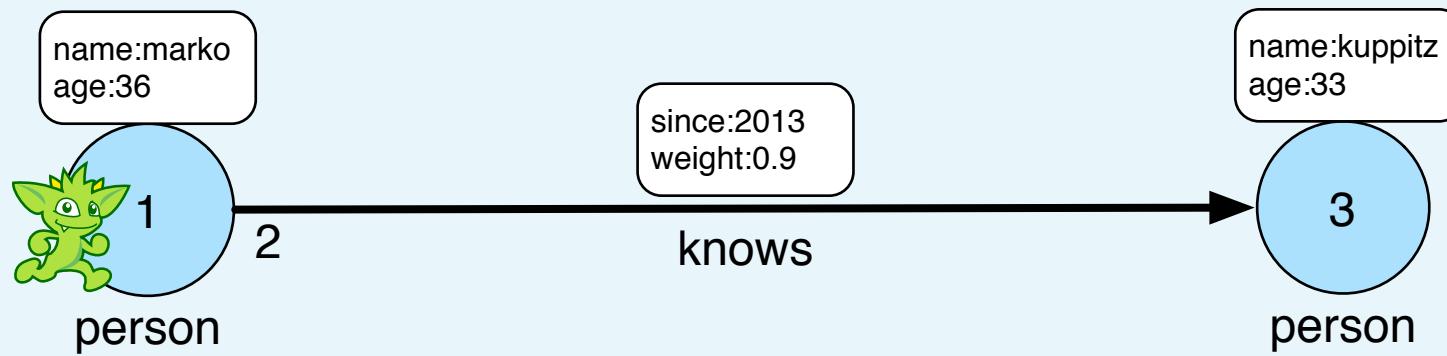
"Back to our super baby toy graph."

```
gremlin> graph  
==>tinkergraph[vertices:2 edges:1]  
gremlin> g = graph.traversal()  
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]  
gremlin>
```

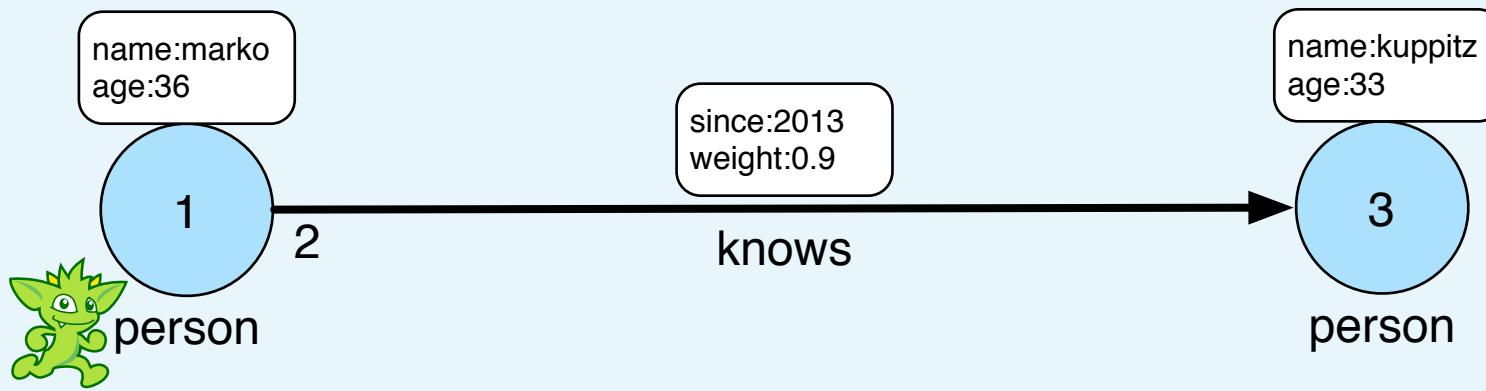
OLTP vs. OLAP  
graph database vs. processor  
Neo4j vs. Giraph



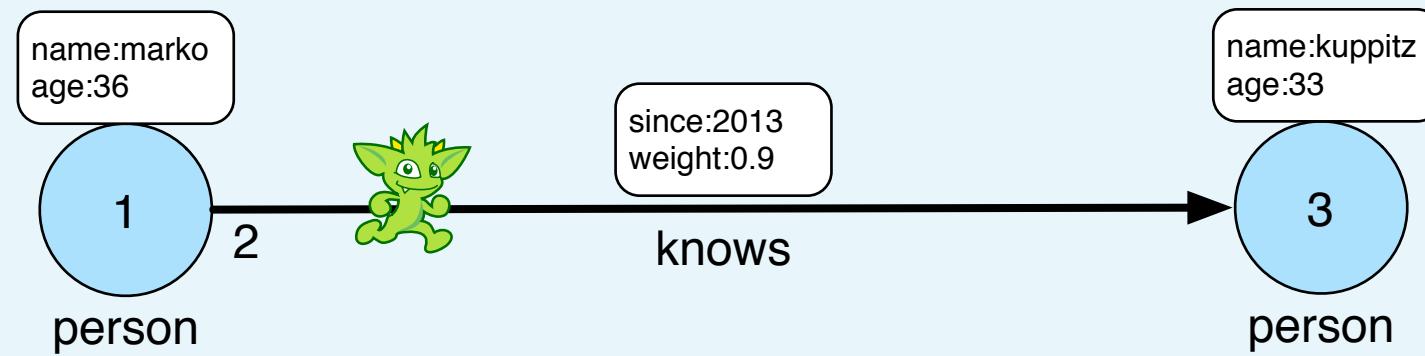
```
gremlin> graph
==>tinkergraph[vertices:2 edges:1]
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]
gremlin> g.V(1)
==>v[1]
gremlin>
```



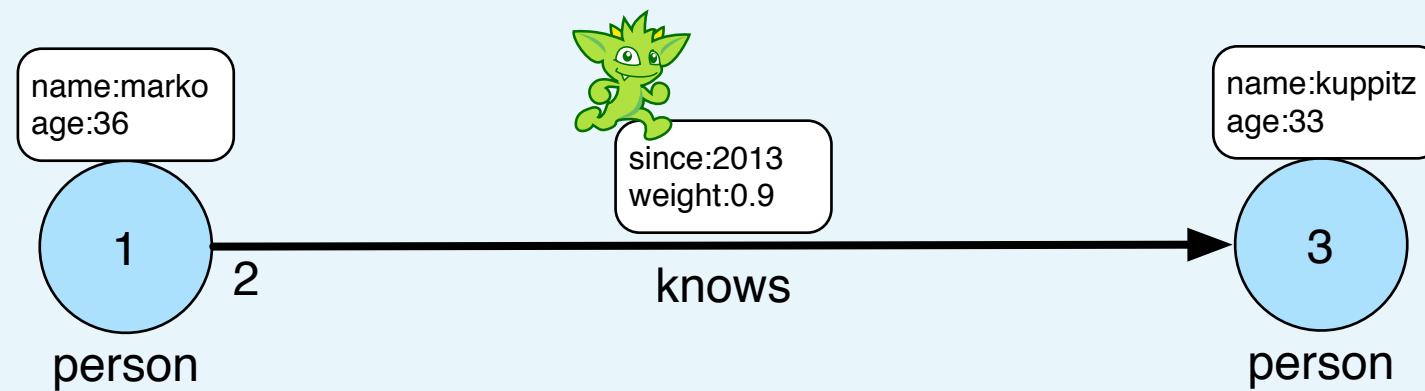
```
gremlin> graph
==>tinkergraph[vertices:2 edges:1]
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]
gremlin> g.V(1)
==>v[1]
gremlin> g.V(1).label()
==>person
gremlin>
```



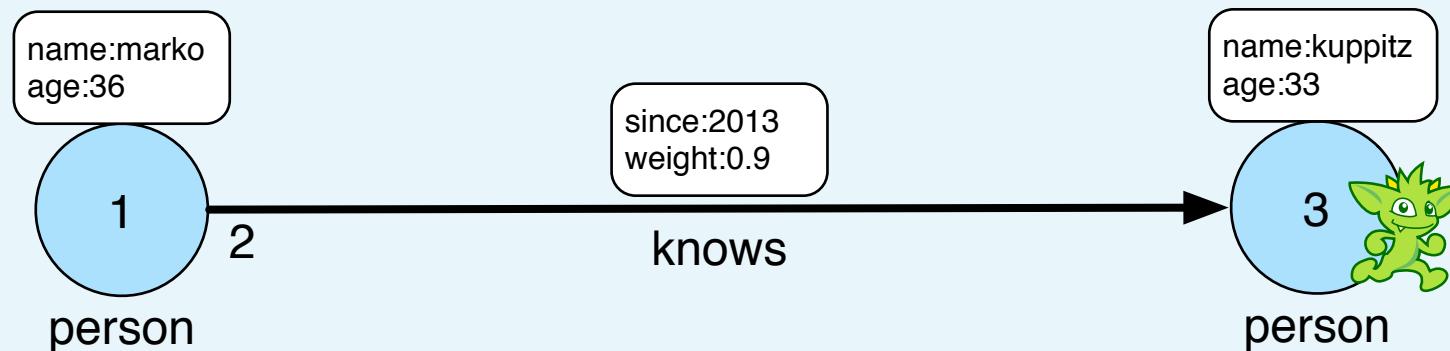
```
gremlin> graph
==>tinkergraph[vertices:2 edges:1]
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]
gremlin> g.V(1)
==>v[1]
gremlin> g.V(1).label()
==>person
gremlin> g.V(1).outE()
==>e[2][1-knows->3]
gremlin>
```



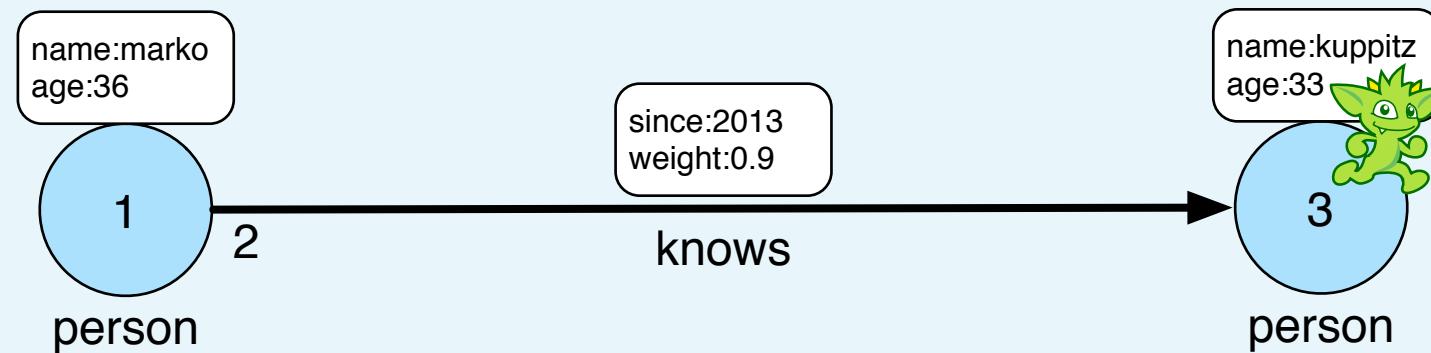
```
gremlin> graph
==>tinkergraph[vertices:2 edges:1]
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]
gremlin> g.V(1)
==>v[1]
gremlin> g.V(1).label()
==>person
gremlin> g.V(1).outE()
==>e[2][1-knows->3]
gremlin> g.V(1).outE().valueMap()
==>[weight:0.9, since:2013]
gremlin>
```



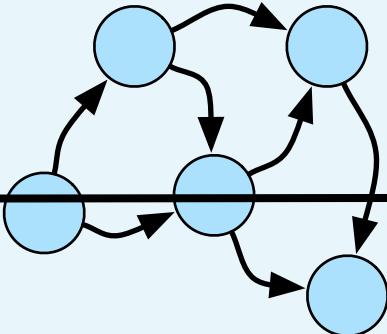
```
gremlin> graph
==>tinkergraph[vertices:2 edges:1]
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]
gremlin> g.V(1)
==>v[1]
gremlin> g.V(1).label()
==>person
gremlin> g.V(1).outE()
==>e[2][1-knows->3]
gremlin> g.V(1).outE().valueMap()
==>[weight:0.9, since:2013]
gremlin> g.V(1).out()
==>v[3]
gremlin>
```



```
gremlin> graph
==>tinkergraph[vertices:2 edges:1]
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:2 edges:1], standard]
gremlin> g.V(1)
==>v[1]
gremlin> g.V(1).label()
==>person
gremlin> g.V(1).outE()
==>e[2][1-knows->3]
gremlin> g.V(1).outE().valueMap()
==>[weight:0.9, since:2013]
gremlin> g.V(1).out()
==>v[3]
gremlin> g.V(1).out().values('age')
==>33
gremlin>
```

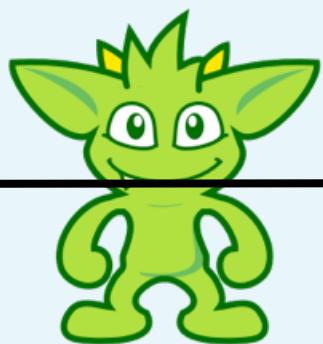


# The Machine Components



The Graph

The Data



The Traverser

The CPU



The Traversal

The Program

# The Machine Components

$$G \xleftarrow[\psi]{\mu} \frac{t \in T}{\{\Delta, \beta, \varsigma, \iota\}} \xrightarrow{\psi} \Psi$$

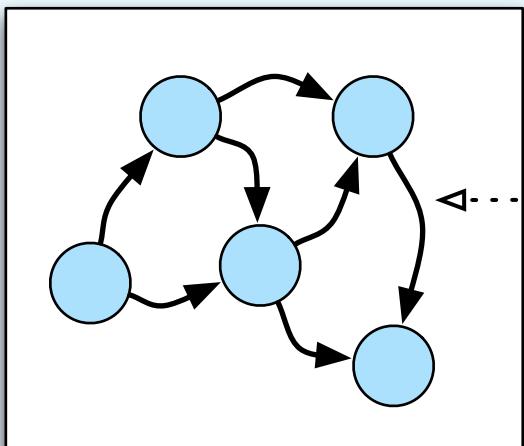
# The Machine Components

$$G \xleftarrow{\mu} \frac{t \in T}{\{\Delta, \beta, \varsigma, \iota\}} \xrightarrow{\psi} \Psi$$

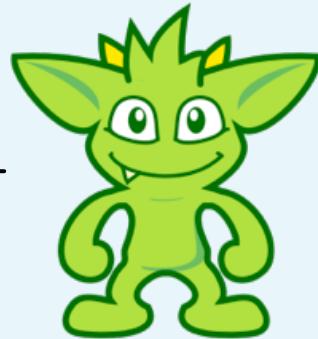


# The Machine Components

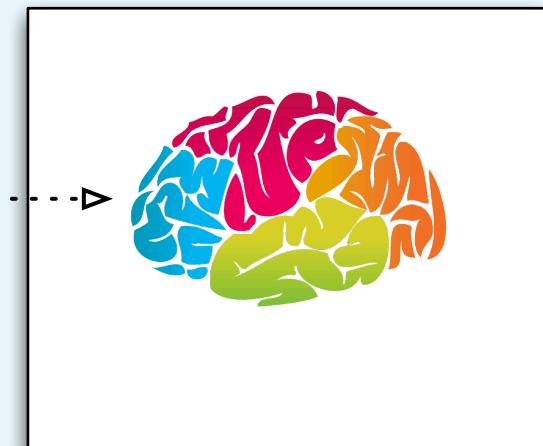
$$G \leftarrow \mu \frac{t \in T}{\{\Delta, \beta, \varsigma, \iota\}} \psi \rightarrow \Psi$$



Graph Structure  
(Data)



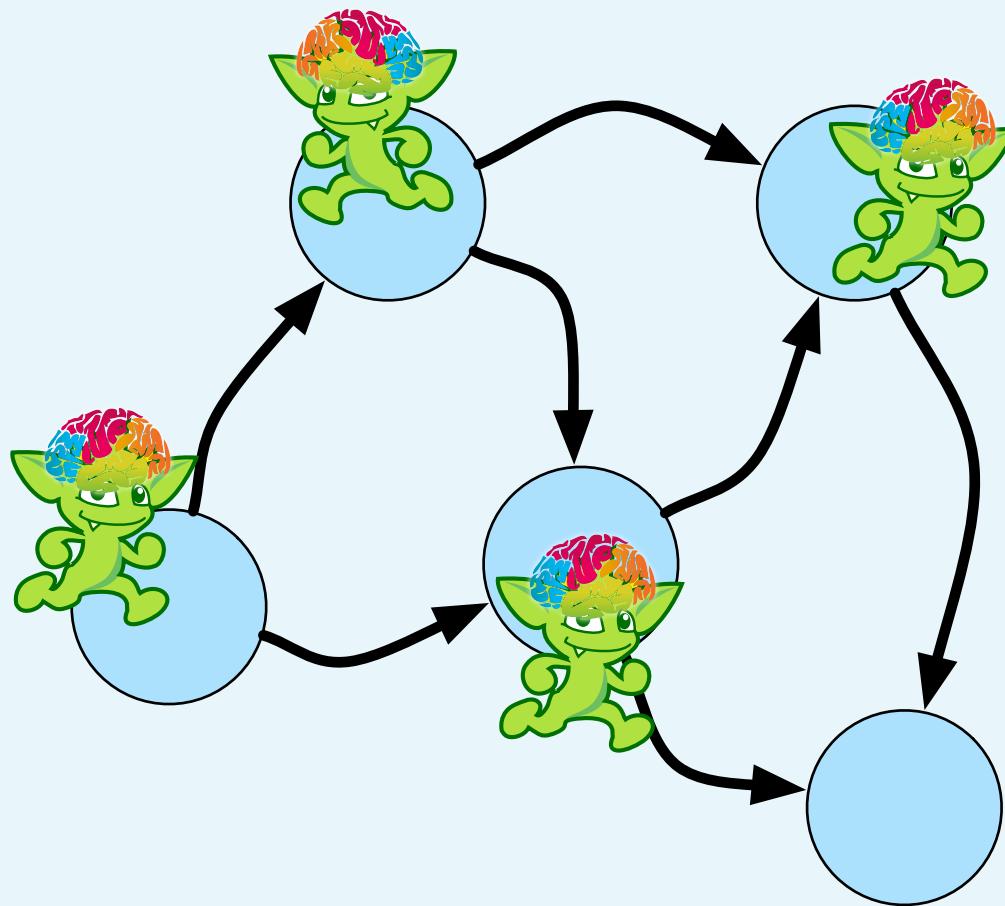
+



Graph Process  
(Algorithm)

= Graph Computing

# The Machine Components



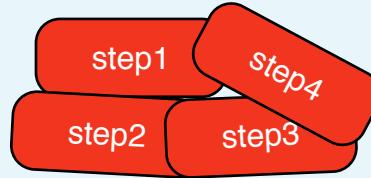
"Traversers move about a graph as instructed by their traversal. The result of the computation is

- 1.) the final location of all halted traversers and
- 2.) the state of any side-effect data structures."

# Part 2: The Gremlin Traversal Language



Gremlin's step library is the instruction set of the Gremlin traversal machine.



A linear/nested composition of steps forms a traversal which is executed by the Gremlin traversal machine.



Gremlin-Java8 is the language provided by TinkerPop, though any language (with respective compiler) can be used to write Gremlin traversals.

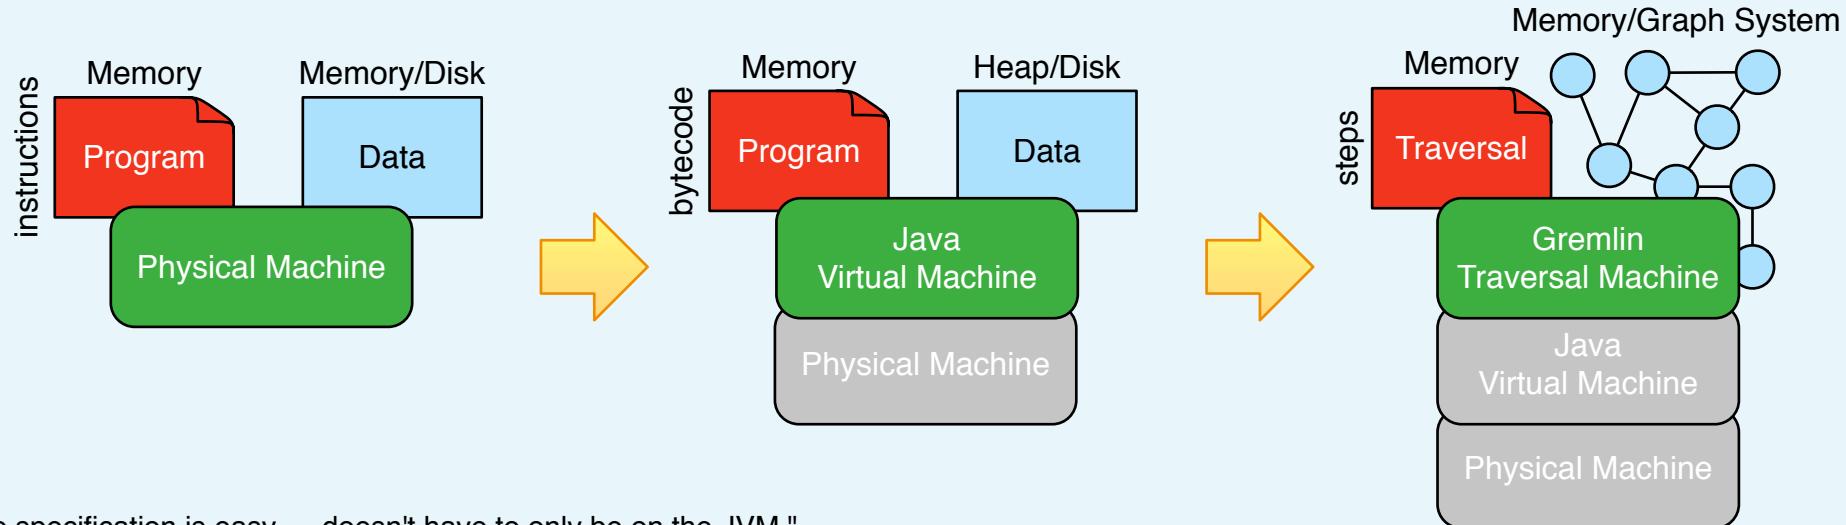
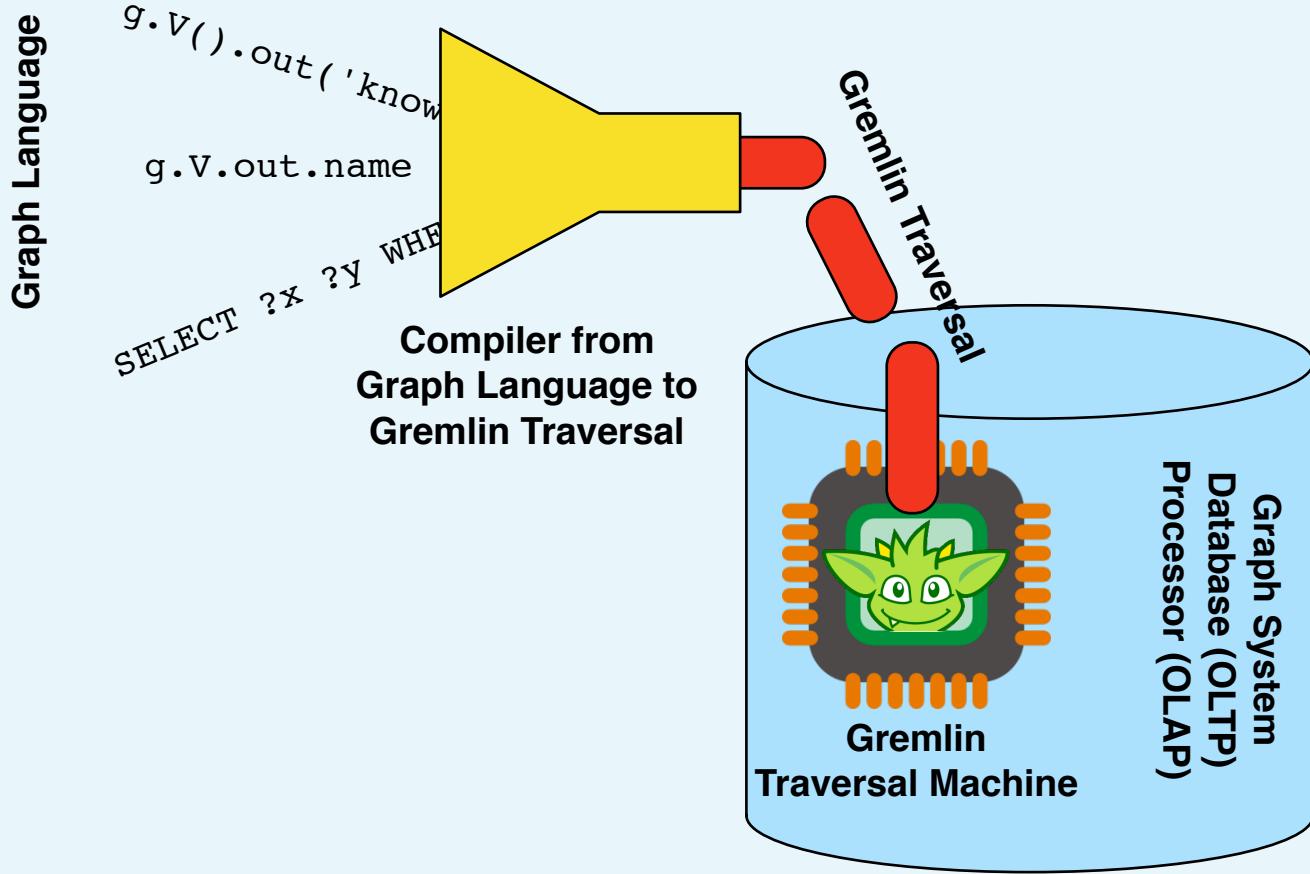


```
g.V(1).as("a").  
out("knows").as("b").  
select("a", "b")
```

Gremlin-Java8



```
SELECT ?a ?b WHERE {  
?a e:knows ?b  
?a v:id ?c  
FILTER (?c == 1)  
}
```



"The specification is easy ... doesn't have to only be on the JVM."

```
~/tinkerpop3$ bin/gremlin.sh
```



```
gremlin>
```

"Gremlin-Groovy can be executed in the Gremlin Console REPL and thus, is good for demonstrations."

```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/  
(o o)
```

```
-----o00o-(3)-o00o-----
```

```
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin>
```

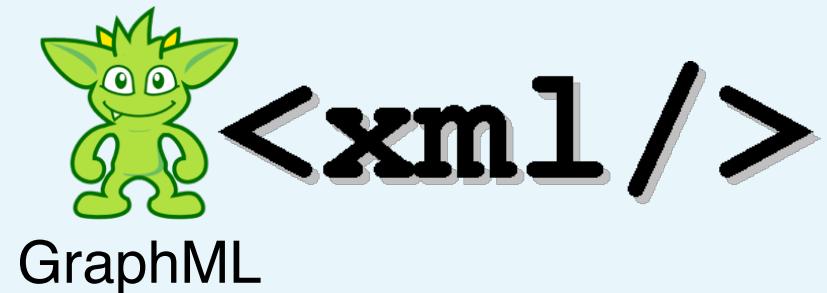
```
~/tinkerpop3$ bin/gremlin.sh
```

```
\,,,/
(○ ○)
-----o00o-(3)-o00o-----
plugin activated: tinkerpop.server
plugin activated: tinkerpop.utilities
plugin activated: tinkerpop.tinkergraph
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin>
```

```
~/tinkerpop3$ bin/gremlin.sh\n\n  \\",,/\\n  (o o)\\n-----o00o-(3)-o00o-----\\nplugin activated: tinkerpop.server\\nplugin activated: tinkerpop.utilities\\nplugin activated: tinkerpop.tinkergraph\\ngremlin> graph = TinkerGraph.open()\\n==>tinkergraph[vertices:0 edges:0]\\ngremlin> graph.io(graphml()).readGraph('data/grateful-dead.xml')\\n==>null\\ngremlin>
```

```
~/tinkerpop3$ bin/gremlin.sh
```

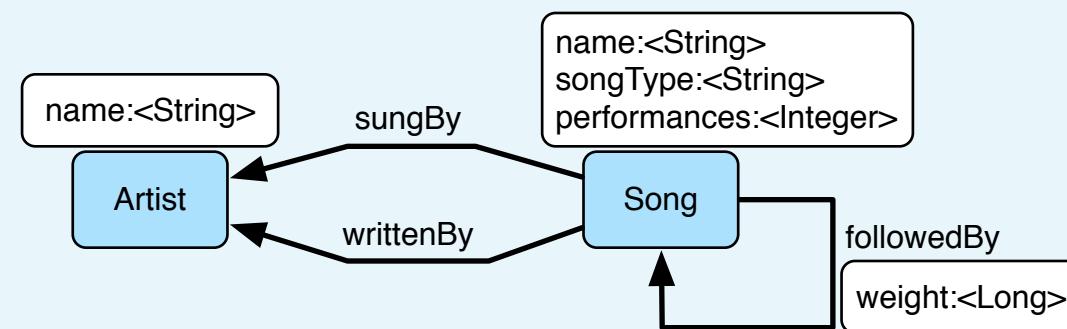
```
\,,,/  
(o o)  
-----o00o-(3)-o00o-----  
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin> graph = TinkerGraph.open()  
==>tinkergraph[vertices:0 edges:0]  
gremlin> graph.io(graphml()).readGraph('data/grateful-dead.xml')  
==>null  
gremlin>
```



"TinkerPop supports three I/O graph formats. Though the user can add their own format/parser."

```
~/tinkerpop3$ bin/gremlin.sh
```

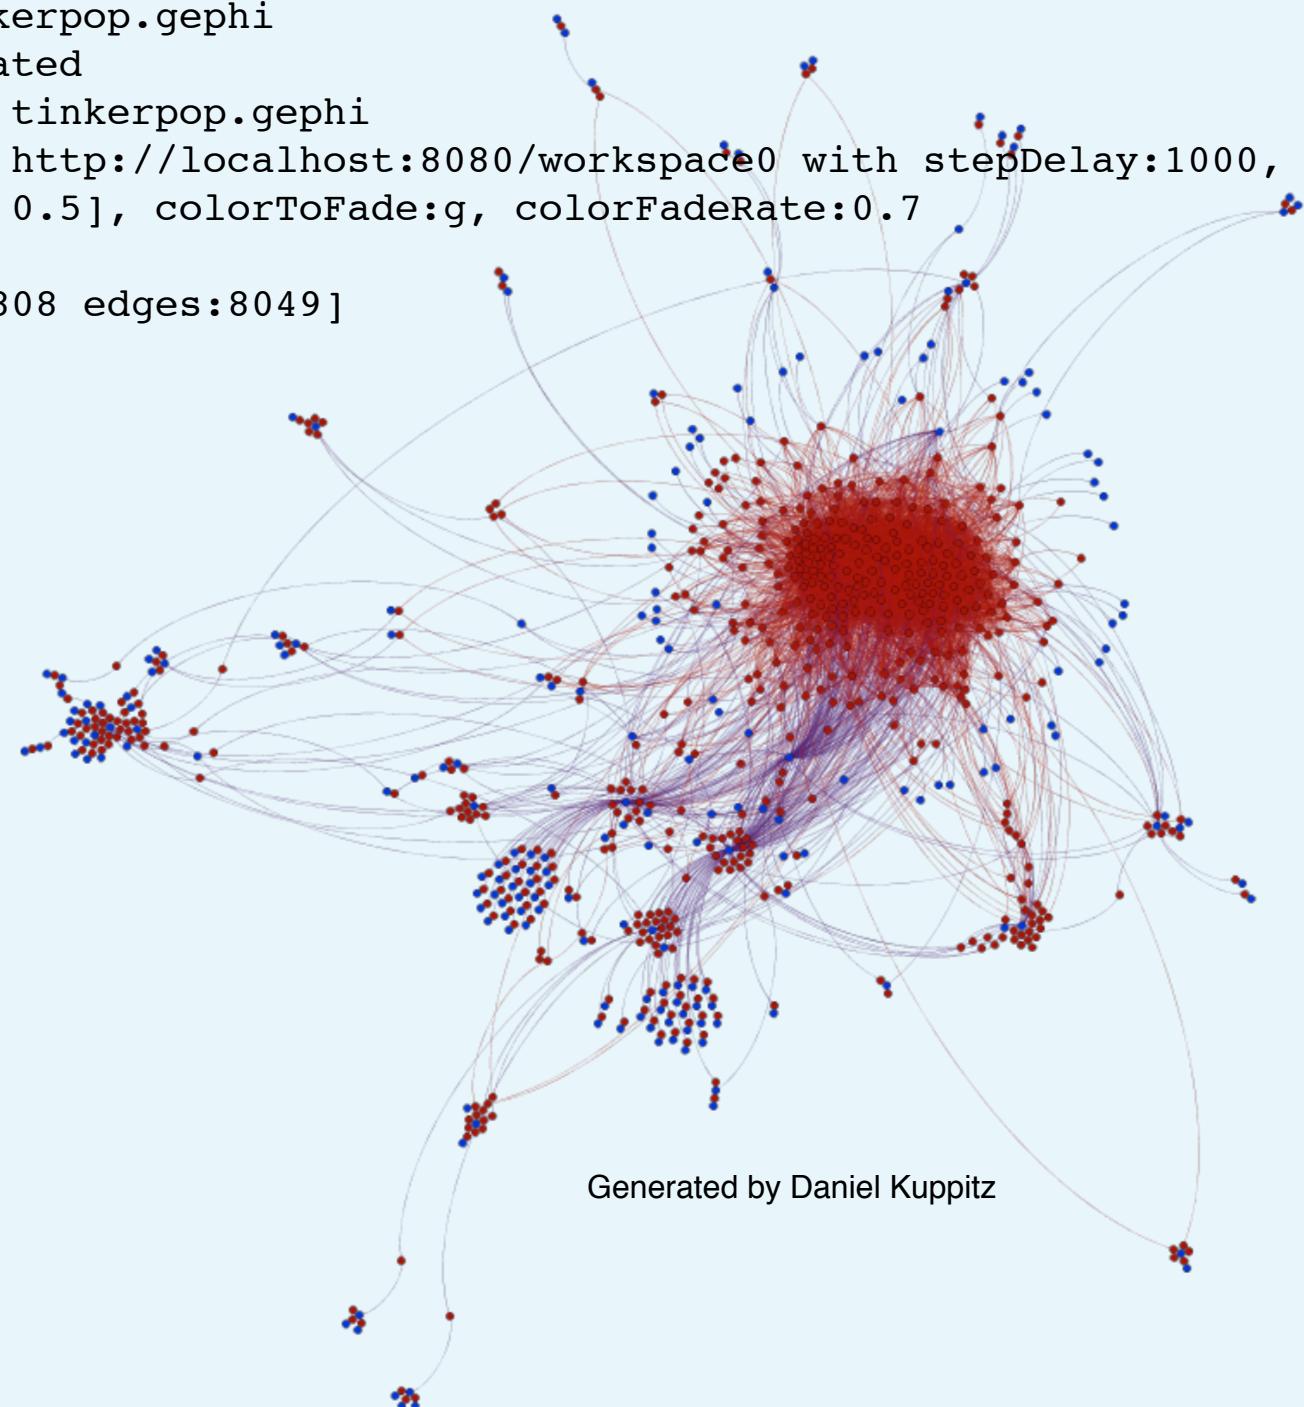
```
\,,,/  
(o o)  
----o00o-(3)-o00o----  
plugin activated: tinkerpop.server  
plugin activated: tinkerpop.utilities  
plugin activated: tinkerpop.tinkergraph  
gremlin> graph = TinkerGraph.open()  
==>tinkergraph[vertices:0 edges:0]  
gremlin> graph.io(graphml()).readGraph('data/grateful-dead.xml')  
==>null  
gremlin>
```



```
gremlin> :plugin use tinkerpop.gephi
==>tinkerpop.gephi activated
gremlin> :remote connect tinkerpop.gephi
==>Connection to Gephi - http://localhost:8080/workspace0 with stepDelay:1000,
startRGBColor:[0.0, 1.0, 0.5], colorToFade:g, colorFadeRate:0.7
gremlin> :> graph
==>tinkergraph[vertices:808 edges:8049]
gremlin>
```



<http://gephi.org>



Rodriguez, M.A., Kuppitz, D., Yim, K., "Tales From the TinkerPop," DataStax Engineering Blog, 2015.  
<http://www.datastax.com/dev/blog/tales-from-the-tinkerpop>

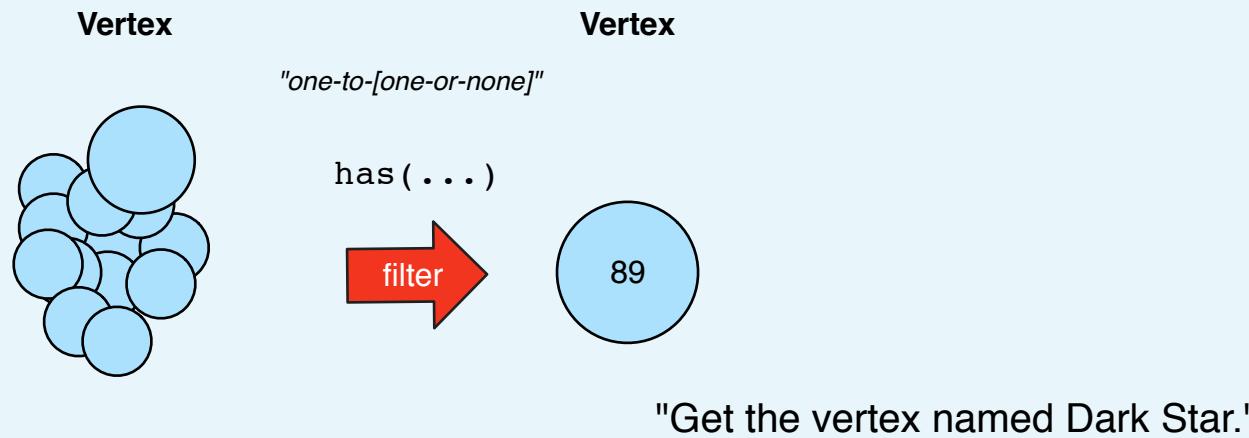
This is an actual Gremlin/R session I performed for this presentation.  
I was interested in understanding why Grateful Dead concerts continue to fascinate me.



**SPOILER ALERT:** No local correlations -- correlations exist in the stationary distribution. EigenDead..the long run."

```
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(graphml()).readGraph('data/grateful-dead.xml')
==>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:808 edges:8049], standard]
gremlin>
```

```
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(graphml()).readGraph('data/grateful-dead.xml')
==>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:808 edges:8049], standard]
gremlin> g.V().has('name', 'DARK STAR')
==>v[89]
gremlin>
```



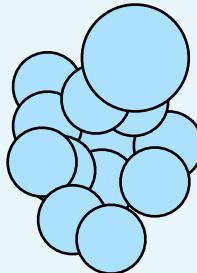
```
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(graphml()).readGraph('data/grateful-dead.xml')
==>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:808 edges:8049], standard]
gremlin> g.V().has('name', 'DARK STAR')
==>v[89]
```



### ProviderOptimizationStrategy

\* OLTP graph system providers turn this into an index-lookup.

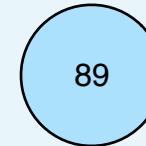
**Vertex**



**Vertex**

"one-to-[one-or-none]"

has( . . . )

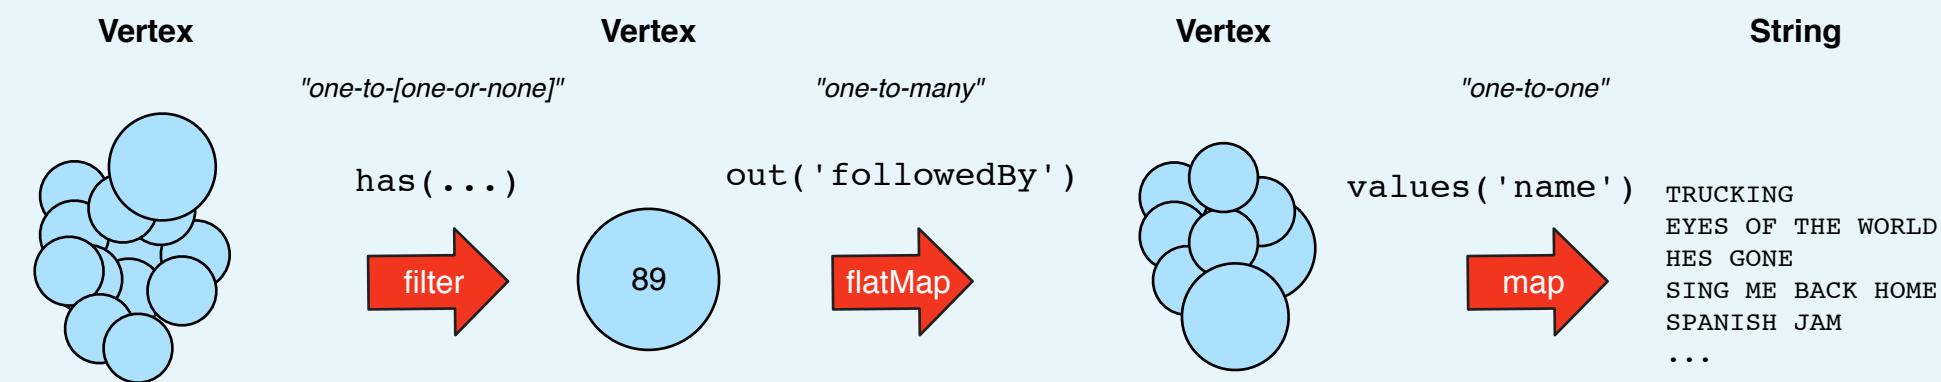


"TraversalStrategies are an important part of Gremlin (discussed at length in the conference article)."

```

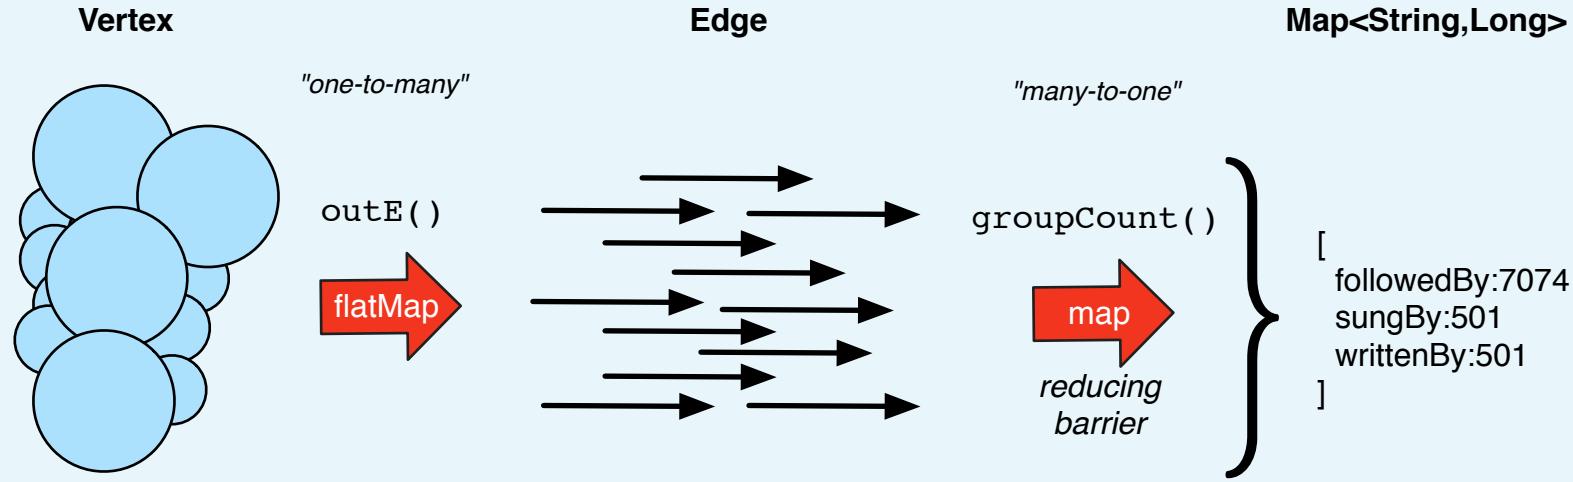
gremlin> graph = TinkerGraph.open()
==>tinkergraph[vertices:0 edges:0]
gremlin> graph.io(graphml()).readGraph( 'data/grateful-dead.xml' )
==>null
gremlin> g = graph.traversal()
==>graphtraversalsource[tinkergraph[vertices:808 edges:8049], standard]
gremlin> g.V().has('name', 'DARK STAR')
==>v[89]
gremlin> g.V().has('name', 'DARK STAR').out('followedBy').values('name')
==>TRUCKING
==>EYES OF THE WORLD
==>HES GONE
==>SING ME BACK HOME
==>SPANISH JAM
==>CHINA DOLL
==>MORNING DEW
==>WHARF RAT
==>THE OTHER ONE
==>MIND LEFT BODY JAM
...
gremlin>

```



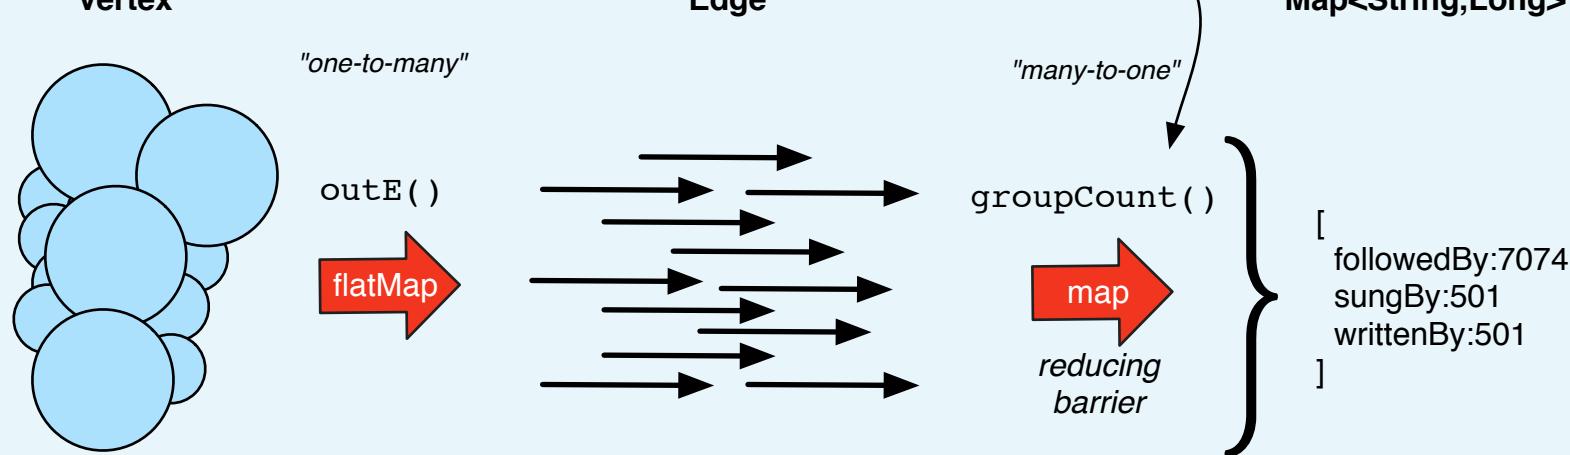
"What are the names of the songs that have followed Dark Star in concert?"

```
gremlin> g.V().outE().groupCount().by(label)
==>[followedBy:7047, sungBy:501, writtenBy:501]
gremlin>
```



"What is the distribution of edge labels in the graph?"

```
gremlin> g.V().outE().groupCount().by(label)  
==>[followedBy:7047, sungBy:501, writtenBy:501]  
gremlin>
```



"What is the distribution of edge labels in the graph?"

```
gremlin> g.v().groupCount().by(outE().count())
==>0=224
==>1=26
==>2=254
==>3=45
==>4=24
==>5=20
==>6=10
==>7=10
==>8=6
==>9=8
...
gremlin>
```

"What is the out-degree distribution of the graph?"

```
gremlin> g.V().groupCount().by(outE().count())
==>0=224
==>1=26
==>2=254
==>3=45
==>4=24
==>5=20
==>6=10
==>7=10
==>8=6
==>9=8
...
gremlin> f = new File('grateful-dead-distribution.txt')
==>grateful-dead-distribution.txt
gremlin> g.V().groupCount().by(outE().count()).
    next().each{degree,freq -> f.append(degree + '\t' + freq + '\n')}
gremlin> :quit
```

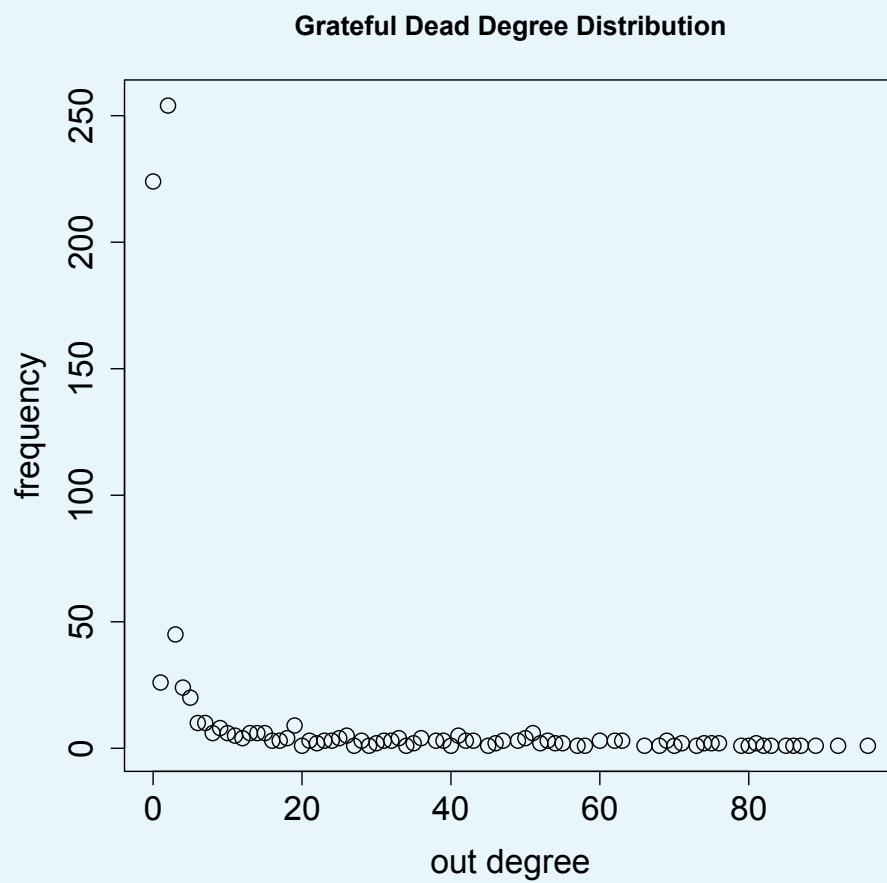


"Save result to a tab delimited file for further analysis."

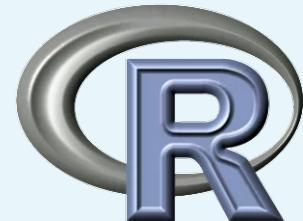
```
$ r  
> t <- read.table("grateful-dead-distribution.txt")  
>
```



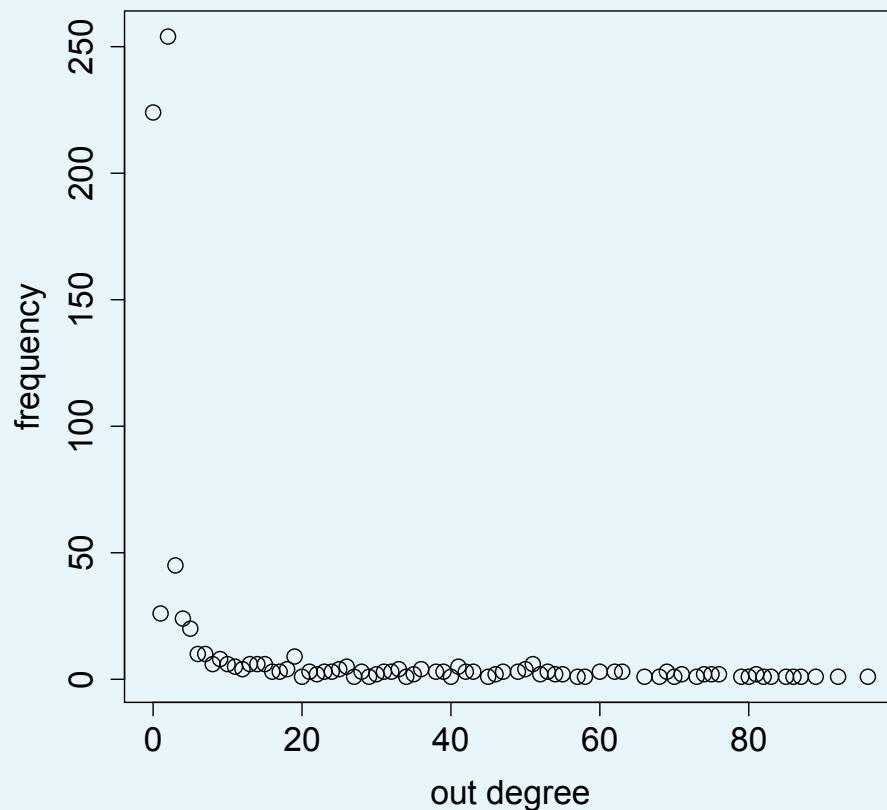
```
$ r  
> t <- read.table("grateful-dead-distribution.txt")  
> plot(t,xlab="out degree", ylab="frequency",cex=1.5,cex.lab=1.5,cex.axis=1.5,  
      main="Grateful Dead Degree Distribution")  
>
```



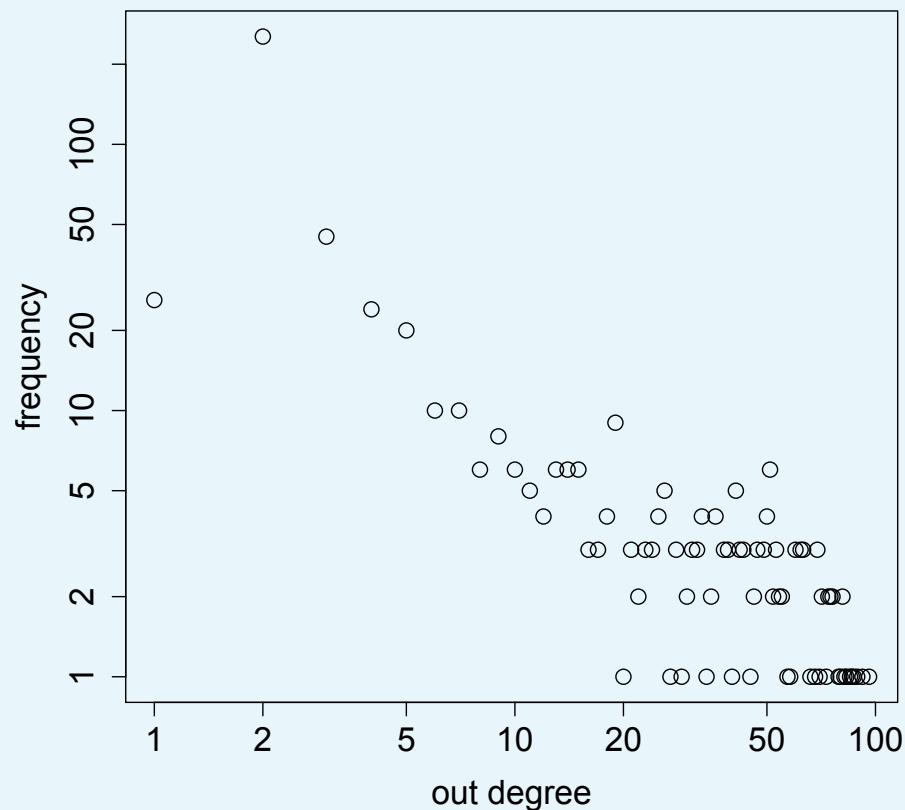
```
$ r  
> t <- read.table("grateful-dead-distribution.txt")  
> plot(t,xlab="out degree", ylab="frequency",cex=1.5,cex.lab=1.5,cex.axis=1.5,  
       main="Grateful Dead Degree Distribution")  
> plot(t,xlab="out degree", ylab="frequency",cex=1.5,cex.lab=1.5,cex.axis=1.5,  
       main="Grateful Dead Degree Distribution",log="xy")  
>
```



Grateful Dead Degree Distribution



Grateful Dead Degree Distribution



"Yea, yet another power law. Its 2005 and I get a free publication!"

```
gremlin> g.V().as('a').out('followedBy').as('b').  
           select('a','b').by('performances')  
==>[a:5, b:1]  
==>[a:5, b:531]  
==>[a:5, b:394]  
==>[a:5, b:293]  
==>[a:5, b:1]  
==>[a:1, b:473]  
==>[a:1, b:24]  
==>[a:531, b:87]  
==>[a:531, b:41]  
==>[a:531, b:254]  
...  
gremlin> f = new File('performance-assortativity.txt')  
==>performance-assortativity.txt  
gremlin> g.V().as('a').out('followedBy').as('b').  
           select('a','b').by('performances').  
           each{f.append(it.a + '\t' + it.b + '\n')}
```

*"Gremlins of a green, bulk together."*



"Are songs that follow each other assortative by the number of times they are played in concert?"

<https://en.wikipedia.org/wiki/Assortativity>

```

gremlin> g.V().as('a').out('followedBy').as('b').
           select('a','b').by('performances')
==>[a:5, b:1]
==>[a:5, b:531]
==>[a:5, b:394]
==>[a:5, b:293]
==>[a:5, b:1]
==>[a:1, b:473]
==>[a:1, b:24]
==>[a:531, b:87]
==>[a:531, b:41]
==>[a:531, b:254]
...
gremlin> f = new File('performance-assortativity.txt')
==>performance-assortativity.txt
gremlin> g.V().as('a').out('followedBy').as('b').      "...you know: birds and feathers and flocking together."
           select('a','b').by('performances').
           each{f.append(it.a + '\t' + it.b + '\n')}

```



```

> cor.test(t[,1],t[,2], test='pearson')
Pearson's product-moment correlation

data: t[, 1] and t[, 2]
t = -3.9525, df = 7045, p-value = 7.809e-05
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
-0.07030966 -0.02371587
sample estimates:
cor
-0.04703835

```

No correlation  
between songs  
by performance.



```
gremlin> g.V().as('a').out('followedBy').as('b') .  
      select('a','b').by('songType')  
==>[a:cover, b:cover]  
==>[a:cover, b:cover]  
==>[a:cover, b:original]  
==>[a:cover, b:cover]  
==>[a:cover, b:cover]  
==>[a:cover, b:original]  
==>[a:cover, b:original]  
==>[a:cover, b:original]  
==>[a:cover, b:cover]  
==>[a:cover, b:cover]  
...  
gremlin> f = new File('songType-assortativity.txt')  
==>songType-assortativity.txt  
gremlin> g.V().as('a').out('followedBy').as('b') .  
      select('a','b').by('songType') .  
      each{f.append((it.a == 'cover' ? 0 : 1) + '\t' +  
                (it.b == 'cover' ? 0 : 1) + '\n')}  
    "Hmmm.. need numbers -- ah! only two song types."
```



"Are songs that follow each other assortative by either being a cover or an original song?"

```

gremlin> g.V().as('a').out('followedBy').as('b').
    select('a','b').by('songType')
==>[a:cover, b:cover]
==>[a:cover, b:cover]
==>[a:cover, b:original]
==>[a:cover, b:cover]
==>[a:cover, b:cover]
==>[a:cover, b:original]
==>[a:cover, b:original]
==>[a:cover, b:original]
==>[a:cover, b:cover]
==>[a:cover, b:cover]
...
gremlin> f = new File('songType-assortativity.txt')
==>songType-assortativity.txt
gremlin> g.V().as('a').out('followedBy').as('b').
    select('a','b').by('songType').
    each{f.append((it.a == 'cover' ? 0 : 1) + '\t' +
        (it.b == 'cover' ? 0 : 1) + '\n')}

```




---

```

> cramersV(chisq.test(t[,1],t[,2])$observed)
[1] 0.0093161

```

No correlation  
between songs  
by song type (original or cover).



"Cramer's V is for nominal data correlations."

```
gremlin> g.V().as('a').out('followedBy').as('b').  
           select('a','b').by(out('sungBy').values('name'))  
==>[a:Garcia, b:Spencer_Davis]  
==>[a:Garcia, b:Weir]  
==>[a:Garcia, b:Garcia]  
==>[a:Garcia, b:Garcia]  
==>[a:Garcia, b:Weir]  
==>[a:Spencer_Davis, b:Weir]  
==>[a:Spencer_Davis, b:Grateful_Dead]  
==>[a:Weir, b:Garcia]  
==>[a:Weir, b:All]  
==>[a:Weir, b:Garcia]  
...  
gremlin> f = new File('singer-assortativity.txt')  
==>singer-assortativity.txt  
gremlin> g.V().as('a').out('followedBy').as('b').  
           select('a','b').by(out('sungBy').values('name')).  
           each{f.append(it.a.hashCode() + '\t' + it.b.hashCode() + '\n')}
```

*"Duh! Now I need an algorithm to turn a String to a unique integer.  
Duh -- hashCode()."*



"Are songs that follow each other assortative by their singer?"

```

gremlin> g.V().as('a').out('followedBy').as('b').
           select('a','b').by(out('sungBy').values('name'))
==>[a:Garcia, b:Spencer_Davis]
==>[a:Garcia, b:Weir]
==>[a:Garcia, b:Garcia]
==>[a:Garcia, b:Garcia]
==>[a:Garcia, b:Weir]
==>[a:Spencer_Davis, b:Weir]
==>[a:Spencer_Davis, b:Grateful_Dead]
==>[a:Weir, b:Garcia]
==>[a:Weir, b:All]
==>[a:Weir, b:Garcia]
...
gremlin> f = new File('singer-assortativity.txt')
==>singer-assortativity.txt
gremlin> g.V().as('a').out('followedBy').as('b').
           select('a','b').by(out('sungBy').values('name')).  

           each{f.append(it.a.hashCode() + '\t' + it.b.hashCode() + '\n')}

```




---

```

> cramersV(chisq.test(t[,1],t[,2])$observed)
[1] 0.2354349

```

*Songs are weakly correlated by singers.*



"However, its typically just Jerry and Bobby trading off in ones or twos..."

```
gremlin> g.V().hasLabel('song').  
    repeat(out('followedBy').groupCount('m').by('name'))).times(8).  
        cap('m').  
    order(local).by(valueDecr).  
    limit(local,10)
```

"Sorta getting bored working on the slides. My improv-vibe died...  
oh yea, but I got a nasty classic riff I remember."

"What songs are most central in the concert network?"

"Not PaaaaageRank again."

```
gremlin> g.V().hasLabel('song').  
    repeat(out('followedBy').groupCount('m').by('name'))).times(8).  
        cap('m').  
    order(local).by(valueDecr).  
    limit(local,10)  
==>PLAYING IN THE BAND=34142246667508  
==>ME AND MY UNCLE=32094411419320  
==>JACK STRAW=31867238591868  
==>EL PASO=29973481580211  
==>TRUCKING=29819272116849  
==>PROMISED LAND=28663488257022  
==>CHINA CAT SUNFLOWER=28569992924918  
==>CUMBERLAND BLUES=26320323048221  
==>LOOKS LIKE RAIN=26138795229794      "Huh, those are the tracks from the 'greatest' Grateful Dead's greatest hits."  
==>RAMBLE ON ROSE=26059394903880  
gremlin>
```



[https://en.wikipedia.org/wiki/What\\_a\\_Long\\_Strange\\_Trip\\_It%27s\\_Been](https://en.wikipedia.org/wiki/What_a_Long_Strange_Trip_It%27s_Been)

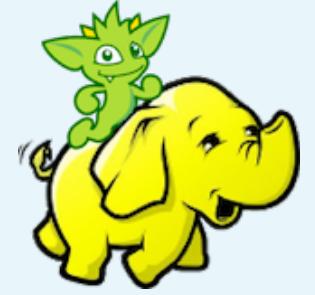
"Rodriguez, M.A., Gintautas, V., Pepe, A., "A Grateful Dead Analysis: The Relationship Between Concert and Listening Behavior," First Monday 14:1, ISSN:1396-0466, <http://arxiv.org/abs/0807.2466>, January 2009."

```
gremlin> g.V().hasLabel('song').  
    repeat(out('followedBy').groupCount('m').by('name'))).times(8).  
        cap('m').  
        order(local).by(valueDecr).  
        limit(local,10)  
==>PLAYING IN THE BAND=34142246667508  
==>ME AND MY UNCLE=32094411419320  
==>JACK STRAW=31867238591868  
==>EL PASO=29973481580211  
==>TRUCKING=29819272116849  
==>PROMISED LAND=28663488257022  
==>CHINA CAT SUNFLOWER=28569992924918  
==>CUMBERLAND BLUES=26320323048221  
==>LOOKS LIKE RAIN=26138795229794  
==>RAMBLE ON ROSE=26059394903880  
gremlin> clock(10){g.V().hasLabel('song').  
    repeat(out('followedBy').groupCount('m').by('name')).times(8).  
        cap('m').  
        order(local).by(valueDecr).  
        limit(local,10)}  
==>4040.761404  
gremlin>
```

"34 trillion traversers  
passed through this vertex."

"4 seconds to analyze that many paths -- that is the power of bulking."

```
gremlin> graph = HadoopGraph.open( 'hadoop-grateful-gryo.properties' )
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin>
```



"And now note how the same Gremlin queries you've seen thus far can execute across a cluster."

```
gremlin> graph = HadoopGraph.open( 'hadoop-grateful-gryo.properties' )
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin> g = graph.traversal(computer(SparkGraphComputer))
==>graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],
sparkgraphcomputer]
gremlin>
```

GraphComputer's  
are OLAP systems



"Lets use Apache Spark via TinkerPop's SparkGraphComputer."

```
gremlin> graph = HadoopGraph.open('hadoop-grateful-gryo.properties')
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin> g = graph.traversal(computer(SparkGraphComputer))
==>graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],
sparkgraphcomputer]
gremlin> g.V().match(
    __.as('a').out('sungBy').as('b'),
    __.as('a').out('writtenBy').as('b'),
    __.as('a').values('name').as('c'),
    __.as('b').values('name').as('d')).  
select('c','d')
```

" / " is required by Gremlin-Groovy  
as "as" is a keyword in Groovy.



"Which songs were written and sung by the same person?"

```

gremlin> graph = HadoopGraph.open('hadoop-grateful-gryo.properties')
==>hadoopgraph[gryoinputformat->gryooutputformat]
gremlin> g = graph.traversal(computer(SparkGraphComputer))
==>graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],
sparkgraphcomputer]
gremlin> g.V().match(
    __.as('a').out('sungBy').as('b'),
    __.as('a').out('writtenBy').as('b'),
    __.as('a').values('name').as('c'),
    __.as('b').values('name').as('d'))..
    select('c','d')
==>[c:ANY WONDER, d:Unknown]
==>[c:WALK DOWN THE STREET, d:Unknown]
==>[c:LEAVE YOUR LOVE AT HOME, d:Unknown]
==>[c:COWBOY SONG, d:Unknown]
==>[c:NEIGHBORHOOD GIRLS, d:Suzanne_Vega]
==>[c:MINDBENDER, d:Garcia_Lesh]
==>[c:EQUINOX, d:Lesh]
==>[c:NO LEFT TURN UNSTONED (CARDBOARD COWBOY), d:Lesh]
==>[c:CHILDHOODS END, d:Lesh]
==>[c:NEVER TRUST A WOMAN, d:Mydland]
...
gremlin>

```



"That was a distributed, declarative graph pattern match query.  
This could have been over a 1000 node cluster."



Gremlin-Groovy

```
g.V().match(  
  __.as('a').out('sungBy').as('b'),  
  __.as('a').out('writtenBy').as('b'),  
  __.as('a').values('name').as('c'),  
  __.as('b').values('name').as('d')).  
    select('c','d')
```



```
SELECT ?c ?d WHERE {  
  ?a e:writtenBy ?b .  
  ?a e:sungBy ?b .  
  ?a v:name ?c .  
  ?b v:name ?d  
}
```

"Gremlin supports the 'match'-style found in most pattern match languages such as SPARQL."



Gremlin-Groovy

```
g.V().match(  
  __.as('a').out('sungBy').as('b'),  
  __.as('a').out('writtenBy').as('b'),  
  __.as('a').values('name').as('c'),  
  __.as('b').values('name').as('d')).  
    select('c','d')
```



```
SELECT ?c ?d WHERE {  
  ?a e:writtenBy ?b .  
  ?a e:sungBy ?b .  
  ?a v:name ?c .  
  ?b v:name ?d  
}
```



**HOMEWORK:** Use Apache Calcite and convert SQL to a Gremlin traversal.

Ted Wilmes was here.

"Two different graph languages can be compiled to the Gremlin traversal machine."

```
gremlin> :install com.datastax sparql-gremlin 0.1
==>Loaded: [com.datastax, sparql-gremlin, 0.1]
gremlin> :plugin use datastax.sparql
==>datastax.sparql activated
gremlin>
```



"Install the SPARQL-Gremlin compiler developed by Daniel Kuppitz of DataStax."

```
gremlin> :install com.datastax sparql-gremlin 0.1
==>Loaded: [com.datastax, sparql-gremlin, 0.1]
gremlin> :plugin use datastax.sparql
==>datastax.sparql activated
gremlin> :remote connect datastax.sparql g
==>SPARQL[graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],
sparkgraphcomputer]]
gremlin>
```

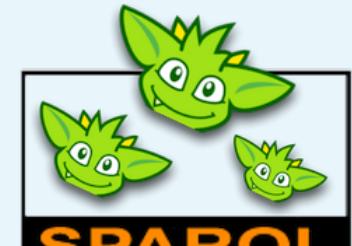


"Given that SPARQL is NOT Groovy, it is passed to the cluster as a remote String."

```

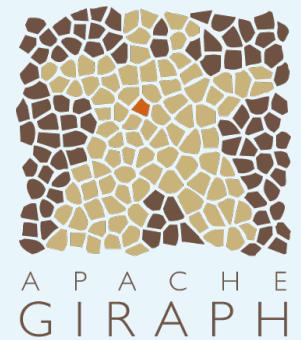
gremlin> :install com.datastax sparql-gremlin 0.1
==>Loaded: [com.datastax, sparql-gremlin, 0.1]
gremlin> :plugin use datastax.sparql
==>datastax.sparql activated
gremlin> :remote connect datastax.sparql g
==>SPARQL[graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],
sparkgraphcomputer]]
gremlin> :> SELECT ?c ?d WHERE {
    ?a e:writtenBy ?b .
    ?a e:sungBy ?b .
    ?a v:name ?c .
    ?b v:name ?d }
==>[c:ANY WONDER, d:Unknown]
==>[c:WALK DOWN THE STREET, d:Unknown]
==>[c:LEAVE YOUR LOVE AT HOME, d:Unknown]
==>[c:COWBOY SONG, d:Unknown]
==>[c:NEIGHBORHOOD GIRLS, d:Suzanne_Vega]
==>[c:MINDBENDER, d:Garcia_Lesh]
==>[c:EQUINOX, d:Lesh]
==>[c:NO LEFT TURN UNSTONED (CARDBOARD COWBOY), d:Lesh]
==>[c:CHILDHOODS END, d:Lesh]
==>[c:NEVER TRUST A WOMAN, d:Mydland]
...
gremlin>

```



"SPARQL was just executed over Apache Spark by way of the Gremlin traversal machine."

```
gremlin> g = graph.traversal(computer(GiraphGraphComputer))  
==>graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],  
giraphgraphcomputer]  
gremlin>
```



"Use a different GraphComputer."

```
gremlin> g = graph.traversal(computer(GiraphGraphComputer))
==>graphtraversalsource[hadoopgraph[gryoinputformat->gryooutputformat],
giraphgraphcomputer]
gremlin> g.V().match(
    __.as('a').out('sungBy').as('b'),
    __.as('a').out('writtenBy').as('b'),
    __.as('a').values('name').as('c'),
    __.as('b').values('name').as('d')).  
select('c','d')
```

```
INFO org.apache.hadoop.mapreduce.Job - Running job: job_1445539638801_0011
INFO org.apache.hadoop.mapreduce.Job - map 50% reduce 0%
INFO org.apache.hadoop.mapreduce.Job - map 100% reduce 0%
INFO org.apache.hadoop.mapreduce.Job - Counters: 55
```

...

### Giraph Stats

```
Superstep 1 GiraphComputation (ms)=2022
Superstep 2 GiraphComputation (ms)=1994
Superstep 3 GiraphComputation (ms)=999
Superstep 4 GiraphComputation (ms)=1001
Superstep 5 GiraphComputation (ms)=1254
Total (ms)=21425
```

...

```
==>[c:IT MUST HAVE BEEN THE ROSES, d:Hunter]
==>[c:EASY WIND, d:Hunter]
==>[c:WHATLL YOU RAISE, d:Hunter]
==>[c:CRYPTICAL ENVELOPMENT, d:Garcia]
==>[c:CREAM PUFF WAR, d:Garcia]
==>[c:DRUMS, d:Grateful_Dead]
...
gremlin>
```

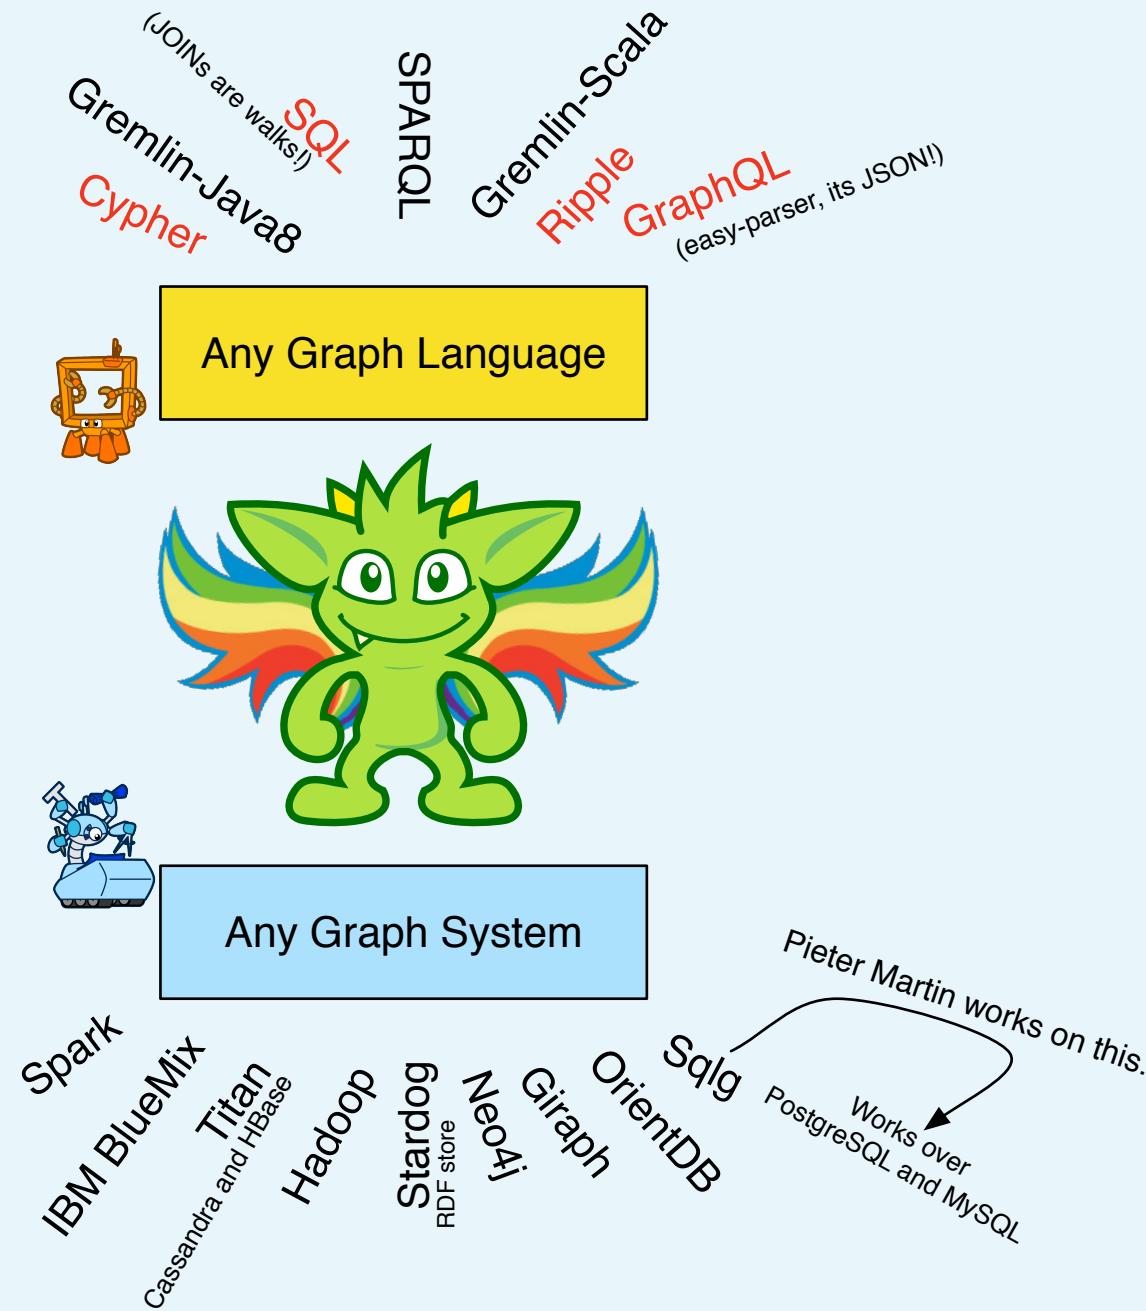
"I'm a Gingerbremlin."



A P A C H E  
G I R A P H

"Gremlin-Groovy just executed over Apache Giraph by way of the Gremlin traversal machine."

RED = "No known compiler."



\* Many system providers are still on TinkerPop2 and thus, haven't migrated to TinkerPop3.

# Keep on Traversin' ...



Thank you....



Association for  
Computing Machinery



Rodriguez, M.A., Kuppitz, D., Yim, K., "Tales From the TinkerPop," DataStax Engineering Blog, 2015.  
<http://www.datastax.com/dev/blog/tales-from-the-tinkerpop>